

What This Book Is For

What data exists about American democracy, and what it says about the country

Democracy's Data

Last updated 1 September 2026

American democracy writes itself down at a scale no other subject can match. Every ten years the government attempts to count every person in the country, at every address, and publishes the result. Every vote in every county is tallied and certified. Every dollar a federal campaign raises is disclosed with a name attached. Every roll call in Congress since 1789 is on the record. Nearly all of it is public, most of it is free, and anyone with a laptop can hold it.

This book is about that shelf. It asks two questions, in this order, of everything on it.

What data is available? Who counts the population, who publishes the returns, who keeps the voter rolls, what a campaign must disclose, what a survey can ask, what a traffic stop leaves behind. Most people have never opened these files and could not say what is in them, and that includes most people who argue about politics for a living. By the end of this book you will know what exists, where it lives, and how to get it yourself.

And what does it say about the country? The data is not a filing cabinet; it is the country's autobiography. Read it and you can watch the population shift south and west until House seats follow. You can watch the electorate age, the parties pull apart in Congress, contested House races become scarce, and the same counties flip that decided the last election. Those are not opinions. They are findings you can hold in your hand, because somebody kept the record.

That second question is what separates this book from its neighbors on the syllabus. A textbook on American politics will tell you, as settled fact, that the president's party loses seats at the midterm. A book on democratic theory will tell you what democracy requires: equality, participation, accountability. This book does neither.

It takes the textbook's sentence and turns it back into a question: *does* the president's party lose seats at the midterm? Then it opens the record, every House election since 1858, and answers with a chart you can read for yourself. It does not argue that democracy requires equal turnout across groups. It measures what turnout actually is, group by group, and hands you the measurement. Assertions become questions, and questions get answered from the file.

There is a third thing, quieter, that makes the first two work: **knowing what a file is before trusting what it says.** Somebody made every dataset, for a reason that was almost never your question. A file will answer some questions well, others badly, and some not at all. So each chapter describes its data before arguing from it, and you will learn to do the same. That habit, not statistics and not programming, is what this book adds to the books you could already read.

01 What the data can show you

A sample of what is waiting in these files, all of it developed in a chapter of this book.

The census shows a country in motion. For a hundred years the population has drained toward the South and West. Because the Constitution ties House seats to the count, political power has moved with it, and you can watch states gain and lose seats decade by decade, in one table. The ACS shows who moves between states and who stays, down to the flow between any pair of states. The returns show that most House races are no longer seriously contested, and the rosters show how long a congressional career really lasts and how it usually ends – retirement, far more often than defeat.

The surveys show what no record can: what the country believes. Party identification drifting over seventy years. The gender gap opening. Confidence in institutions falling, unevenly, for fifty years. Each party convinced the other is more extreme than it is. The roll-call record shows the parties in Congress further apart than at any point since Reconstruction, and it shows it with a method you will read about and be able to question.

People also leave records while doing something else: registering to vote, filing a campaign report, being stopped at a traffic light. Those records show who participates in American democracy and who it touches, in ways no one designed anyone to see.

None of this requires trusting anybody's summary, this book's included. The files are public, every chapter tells you where they live, and every number on these pages is computed from them in code you can inspect. That is the standing invitation of the book: check.

02 Four sections, and each one is a kind of data

The book is organised by **what kind of record the data came from**, not by what it is about. That is unusual, and it is the whole design. Turnout shows up in three different sections of this book, and it means something slightly different in each, because a different kind of record produced it for a different reason. Grouping by topic would hide that. Grouping by kind of data is what makes it visible.

There are countless datasets that bear on American democracy. Nearly all of them are one of three kinds, and each kind answers a different family of questions.

Section	Docs	What it covers
I. Data About the Population	19	Counting people and places: the enumeration, the rolling survey, the estimates between, and the geography they are published on
II. Survey Data	17	Asking a sample of people questions, and weighting the answers until they describe the country
III. Administrative Data	55	Records kept to run a system: returns, rolls, filings, votes, stops
IV. Putting Data Together	20	Combining the kinds, which is where the hard questions and the wrong answers are

Section I is data about the population. Nearly every political number is a fraction, and the bottom of the fraction comes from the Census Bureau. The Bureau produces three main instruments, and telling them apart is half of this section.

The **decennial census** is an enumeration: a once-a-decade attempt to count every person, at a single moment, and the count that allocates seats in Congress. The **American Community Survey** is a rolling sample survey: it asks far more questions, it can describe small places if you let it average over enough years, and it shows trends the once-a-decade snapshot cannot. The **Population Estimates Program** is arithmetic between the counts — the number you get for a year when nobody counted. Underneath all three sits the Bureau's **geography**, the system of nested boxes every American statistic is published in.

Section II is survey data — asking a sample of people questions and using statistics, rather than enumeration, to describe everyone. A survey is the only instrument in this book that can tell you what people *believe*: no record of any kind holds an opinion. But a sample only describes the country after it has been weighted to look like the country, and the benchmark it is weighted to is almost always the census. So this section leans on Section I, on purpose, and says exactly where the leaning happens.

Section III is administrative data — records compiled while running something, by people who were not trying to describe anyone. The certified election result. The voter roll. A campaign's finance filing, a member's roll-call vote, a lobbying registration. A first name on a Social Security card, a traffic stop, a grade stamped on a neighbourhood in 1937.

These records have a great advantage over both counting and asking: nobody is estimating, because the record *is* the event. And they share one great weakness: the record was kept to run a system, its columns do that system's job and no other, and nobody in it volunteered to be described.

Section IV puts the kinds together. Everything before it works with one source at a time. Here you combine them. You will guess a voter's race from a name and an address, and check a survey's claim about turnout against the counted ballots. You will estimate how a group voted from precinct totals, and grade a district map against a rule somebody wrote down.

That is where the genuinely interesting questions live. It is also where nearly all the wrong answers come from — because every join carries the assumptions of both its sources, plus one more about how they line up.

03 Chapters and briefs

Each section is built the same way, out of two kinds of document.

A **chapter** describes a kind of data. What the record is, who makes it and under what obligation, what it is published as, and what it cannot tell you however carefully you ask. The chapters are the readings: the decennial census gets one, the ACS gets one, voter files get one, campaign finance gets one. Read a chapter and you know what the file on your desk actually is.

A **brief** uses the data. One real dataset, one real question, a figure or two, and a conclusion you could defend — followed by what you should have learned and what you could try next. The briefs are the labs, they are short, and each one leans on exactly one chapter: the chapter for the kind of data it uses. 15 chapters and 92 briefs are in the book as this page is built, counted by reading the folder this file sits in.

04 What a chapter does, in the same order

The sections differ; the moves do not. Not every document makes every one, so the count beside each says how many of the 111 chapters do — measured by reading them, like everything else on this page.

The move	What that means	Chapters
Where it came from, and why	Who produced this file, under what obligation, and for whose purpose — which was almost never yours. The address is printed, so you can go and get it yourself	107
What it actually looks like	One real record, shown in full before anything is summarised — because a row is where the surprises are. The chapters that skip it have nothing small enough to print	42
What it says	Summary numbers and figures, built from the file in front of you, with the wrong reading shown as well as the right one. Every number is computed from the data as the page is built, never typed in	111
What it cannot say	The question this source will not answer however carefully you ask — a limit built into the file, not into the analysis	108
Extensions	The tables the figures rest on, linked so you can open them in a spreadsheet, and questions the document left for you — answerable in a spreadsheet, plus a stretch or two beyond it	107

05 Commit to an answer before you turn the page

92 of the 111 chapters stop somewhere before their central result and ask you, in so many words, to write down what you expect.

This is not a quiz and nothing depends on getting it right. It is there because the interesting quantity in each chapter is not the number itself but **the distance between the number and what a reasonable person expected**. A guess you have not written down is one you will quietly revise the moment you see the answer, and you will come away believing you knew all along.

So the prompts are worth honoring, and they are worth honoring in writing. The chapters are built around the moment they turn, and the turn only lands on someone who had committed to the other side of it.

06 One real dataset, one real question, and no code to write

Each chapter takes **one real dataset and one real question**, and goes from a single record to a conclusion you could defend. It does not skip the part where the obvious answer turns out to be wrong.

The data is not simplified for teaching. Every chapter folder carries its own `data/` directory holding the actual files, and each of those directories also holds the script that built them, and beside both a stamp listing every one of those files with its size, its checksum

and the date it was last written. Where a source could not be obtained cleanly, the chapter says so rather than substituting something tidier.

You will not be asked to write any code. This is deliberate and it is not a concession. The code is written for you and it works, and every number on these pages came out of a file, by a route somebody can check. But writing it is not the exercise. The exercise is interpretation: saying what a table shows, what it does not show, and what would have to be true for its conclusion to hold. Making students author code turns a course about evidence into a course about syntax, and the people most likely to be turned away by that are exactly the people who came for the substance.

The chapters do not depend on one another and may be read in any order. They cross-reference each other often, because the same problems recur, but none is written to assume you have read another. Start with whichever question you care about.

07 Four questions, asked of every dataset in the book

Before a document here argues anything, it describes the data it is about to use — answering the same four questions in the same order.

The question	What it catches
What is the source?	A file built for somebody else's purpose, and nobody's obligation to keep it right
Why is it appropriate to this question?	A unit of analysis that cannot carry the question being asked of it
What are the challenges — with the data, or with measuring the thing at all?	A quantity that was never observable, dressed as one that was
Is it clean, or does it need cleaning?	Decisions made on your behalf, upstream, that move the answer

They are not boilerplate; each one protects a different part of a conclusion. One example of why they earn their place. The traffic-stop chapter divides stops by the residents of San Francisco and finds Hispanic drivers stopped 0.86 times as often as white drivers, the lowest of any large group. Divide the same stops by the people who actually drive to work, and the figure is 1.01: essentially even. Neither denominator is right — neither bottom number counts the people at risk of being stopped, because nobody ever counted them. Asking the third question is what surfaces that, and a chapter that hid it behind one confident ratio would be worse than useless.

The same few failures recur across the whole shelf, whatever the source. A code outlives the place it names. A denominator goes unstated. A probable match gets treated as a certain one, or a category is rewritten while the ground stays still. You will meet each of these inside a chapter, at the moment it matters, next to the data that shows it.

And 108 of the 111 chapters carry a late section on what their evidence will and will not bear. That is not a disclaimer. It is the part of the argument that tells you how far you may take it.

08 What this book is not

It is **not a statistics text**. There is arithmetic throughout, and occasionally a simulation or a margin of error: how far off an estimate could be and still be doing its job. But no chapter's purpose is to teach a procedure, and several exist to show a well-regarded procedure producing a confident wrong answer.

It is **not a programming manual**. The code is base R, it is complete, it is on the page, and you are never asked to write any. If you want to learn R, this is not the book, though reading these chapters will not hurt.

It is **not a survey of American politics**. The topics are chosen for what they reveal about evidence, not for coverage. Whole subjects that belong in any politics course are absent here because no dataset makes their difficulties visible, and some minor-looking subjects get a chapter because one does.

And it does not claim that data settles arguments. Most of these chapters end with a question still open – and say which one, and say what evidence would be needed, and say why this file cannot supply it.

09 Sources

Every figure quoted above comes from the chapter it belongs to, read out of that chapter's own tables.

- The two traffic-stop denominators: the policing chapter's stop counts and its American Community Survey denominators, for San Francisco.
- The counts of chapters and briefs, of chapters with a data directory, of prediction prompts and of closing limits sections were obtained by opening the 111 chapter files in the folder above this one and counting.