

Hispanic Origin Is Not a Race on the Census Form

Two questions about one person on the 2020 census form, and the two boxes people use to answer neither

84-355 Democracy's Data

August 2026

On the 2020 census form, question 8 asked whether the person was of Hispanic, Latino or Spanish origin. Question 9 asked what the person's race was. Between them the Bureau printed a note telling respondents to answer both, and explaining why:

For this census, Hispanic origins are not races.

That sentence is a decision, made by the federal government, about the relationship between two things people are routinely asked to be one thing about. It is also a claim you can test, because the form that separates them also records the answers together. This chapter is about what happened when a large minority of respondents declined to accept the separation.

01 The file American districts are drawn from

The Constitution orders a count — an “actual Enumeration” every ten years. But a count requires categories, and categories are made by people.

The file those answers arrive in is called **P.L. 94-171**. A 1975 statute requires the Census Bureau to deliver it to every state within a year of Census Day, specifically so that legislatures can draw districts. Everything downstream runs off it: how many seats a state gets, where the lines go, which jurisdictions must translate their ballots.

Every legal claim about racial fairness in American elections is stated in these categories. A §2 vote-dilution claim requires showing that a minority group is large and compact enough to form a majority in a district — a claim about a number computed from this file. Language-assistance coverage under §203 is determined in these terms. Federal money moves on them.

So if the categories fit people badly, the misfit does not stay in the data. It propagates into the districts, the lawsuits and the appropriations. **A measurement error here is a representation error later**, and there is no second instrument to check it against.

The file is one row per county.

County	State	Total	White	Black	Asian	Two or more races	Hispanic
Allegheny County	Pennsylvania	1,250,578	946,321	163,678	58,541	65,706	34,325

Allegheny County had **1,250,578** people on 1 April 2020. Not an estimate and not a projection — the number the Bureau certified, and the number Pennsylvania's districts were drawn from.

Notice that the last column overlaps every other one. Nothing has gone wrong; that is the two questions arriving in one table.

Item	Value
Source	2020 Census Redistricting Data (P.L. 94-171) Summary File
Reference date	1 April 2020
Counties	503
States	6
People covered	68,883,526
Share of the U.S. population	21%
Tables kept	P1(race), P2(Hispanic origin by race), P5(group quarters)

Six states rather than fifty, because these six together are about 300 MB and they pull race and ethnicity apart in usefully different directions.

Whatever else is wrong with this data, arithmetic is not. Both of the decompositions the file claims can be checked against every county in it rather than assumed, and both hold exactly:

Test	Holds	Counties tested
one race + two or more races = total, in every county	TRUE	503
Hispanic + not Hispanic = total, in every county	TRUE	503

Nothing that follows is a data-quality problem. The count is exact. It is the categories that are in question.

02 Who decides what the boxes are

Not the Census Bureau. The categories on the form come from the **Office of Management and Budget**, a White House office that sets standards every federal agency has to follow. Its Statistical Policy Directive No. 15, in the version governing 2020, is where the boxes were fixed. The Bureau administers a question somebody else wrote.

The directive settles almost everything this chapter is about. It sets **five** minimum race categories: American Indian or Alaska Native, Asian, Black or African American, Native Hawaiian or Other Pacific Islander, and White. It sets **two** ethnicity categories, Hispanic or Latino and not Hispanic or Latino. It says respondents must be allowed to report more

than one race. And where the two are asked as separate questions, it says **ethnicity is asked first**.

So the sentence printed between questions 8 and 9 was not the Bureau's editorial choice. The separation is federal policy, and the ordering of the two questions is federal policy too.

The directive also says what the categories are not. In its own words the race and ethnicity categories are socio-political constructs, and are not to be read as scientific or anthropological. That disclaimer sits in the standard that produces the numbers, which is stranger than it first sounds. The instrument states that the thing it measures is not a natural kind, and then measures it to the person.

Two of the definitions are worth reading in full, because a chapter about boxes that fit people badly should say what the boxes were defined to hold. Black or African American means "a person having origins in any of the black racial groups of Africa." **White means "a person having origins in any of the original peoples of Europe, the Middle East or North Africa."**

Read that second one twice. Under the standard governing the 2020 census, a respondent whose family came from Lebanon, Egypt or Iran was White, and the form offered no other place to put them. That is the same structural problem this chapter measures for Hispanic respondents, produced by a different clause of the same directive, and it is the problem the 2024 revision was written to fix.

Why the question is on the form at all

The decennial census asks about race in order to enforce the Voting Rights Act. Since 1980 the Bureau has asked every person counted to identify their race and whether they are of Hispanic or Latino origin. Self-identification is the rule: the answer is whatever the respondent says it is.

That is the answer to a question students reasonably ask, which is why a headcount for apportionment collects race at all. It does not, for apportionment. Every seat in the House would be allotted identically if question 9 did not exist. The question is there because Section 2 and Section 203 cannot be enforced without it, and it is delivered in P.L. 94-171 for the same reason.

Hold that against the closing argument of this chapter. The categories fit a large population badly, and the reason they are collected is to protect that population.

03 Why this file, and what it took to open it

The obligation is statutory, and it shows in the format. The Bureau does not publish P.L. 94-171 because someone requested it. A 1975 statute requires delivery to every state within a year of Census Day, and the layout is agreed with the states beforehand so that redistricting software can read it on arrival. Most of the files in this book were made for somebody's convenience; this one was made against a deadline that a legislature's calendar depends on.

What actually arrives is less friendly than that sounds. It is pipe-separated text with no column headers, split across four pieces that join on a record number. A county's total population is the sixth field of the first piece, and nothing in the file says so. The 27

columns of the extract this chapter reads were located in a published technical specification and then pulled out by position.

It is the right file because it is the operative one, not because it is the best instrument.

The American Community Survey asks these same two questions every year of a sample, with more detail and more recent data. For almost any question about how Americans describe themselves it would be the better source.

But it is not the file districts are drawn from. When a court is told that a minority group is large and compact enough to make a majority somewhere, the number in front of it came out of P.L. 94-171. And because the ACS uses the same OMB categories, it is not an independent check on them. It inherits the misfit rather than measuring it.

What would genuinely have been better is something the Bureau holds and does not release: the write-in text behind *Some Other Race*, where people said in their own words what the boxes could not accommodate. Those strings are coded into categories before publication. This chapter can show that a hatch was used. It cannot show what was written in it.

The hard part is not counting. It is that the thing being counted is a classification. A population can be enumerated; a race cannot, because the categories are not out in the world waiting to be observed. They are set by the Office of Management and Budget, revised in 1997 and revised again in March 2024, and the second revision merges race and Hispanic origin into a single question by 2029.

The seam this chapter measures is therefore not a permanent feature of American life. It is a product of a two-question form that is being retired. That is an argument for reading it now, and an argument against extending the series forward.

A second difficulty is quieter. To protect respondent confidentiality the Bureau injected noise into the 2020 tables before release, and the file carries no column recording how much. State totals are exact and counties are close, but the smallest county here holds 64 people, and the blocks that districts are actually assembled from are smaller than that by orders of magnitude.

It needed cleaning, and one of the cleaning decisions is a small lesson in itself. The build script downloads the six state archives, keeps only the rows at county level, and joins the three pieces on their record number. Then it throws away the file's own county-name field in favour of the Census Gazetteer's, matched on the FIPS code. That code is held as text rather than as a number. None of these six states begins with a zero, but the national county list it is matched against is full of states that do.

Before writing anything the script re-runs both decompositions and checks Allegheny County's total against the figure the Bureau published. A reader rebuilding this from scratch would have to repeat all of it, and would meet the name field first.

Getting the archives, by contrast, took nothing at all. They sit at fixed public addresses on the Bureau's own server: no key, no form, and nobody standing between the agency that produced the count and the person reading it. So there are no other people's cleaning decisions baked in before the download begins.

The one concession was size. These are large files one state at a time, and taking all fifty would mean moving gigabytes to answer a question that six states answer, so six is what is kept here.

That is a choice about which data to get, and it shows up as a limit on the argument. Every count below is a count of 68,883,526 people. The national figure appears only as a de-

nominator, the bottom number you divide by. Nothing here can be added up into a statement about the country without going back for the other forty-four.

04 What arrives, and what a row of it becomes

Everything above describes the file. This section shows it: four exact slices of the download, not summaries and not retyped versions but the actual bytes. The rest of this section is read out of them.

Here is Allegheny County as the Census Bureau delivers it. This is one line of `pageo2020.p1`, the geographic header, folded to fit the page:

Field	Value	What it means
1	PLST	—
2	PA	—
3	050	summary level — 050 means county
4	00	—
5	00	—
6	000	—
7	00	—
8	0000003	logical record number — the join key
9	05000000US42003	—
10	42003	FIPS code: 42 Pennsylvania, 003 Allegheny
11	1	—
12	2	—
13	42	—
14	01779798	—
15	003	—
16	H1	—
17	01213657	—
50	38300	—
51	1	—
52	430	—
53	99999	—
85	1890571829	—
86	37043982	—
87	Allegheny	—
88	Allegheny County	the county's name
89	A	—
91	1250578	—
92	602095	—
93	+40.4697574	—

Group quarters are not in that line at all. Those are the prison and dormitory counts this chapter uses later. They are in a third file, pa000032020.p1, 15 fields wide, joined back on the same record number:

Field	Value	What it means
1	PLST	file identifier
2	PA	state abbreviation
3	000	—
4	03	—
5	0000003	logical record number — joins back to the two files above
6	34557	—
7	10853	—
8	2151	—
9	185	—
10	8485	—
11	32	—
12	23704	—
13	19266	—
14	22	—
15	4416	—

And here is what the build makes of the three of them together — the whole clean row for Allegheny County, with the position each number came out of. The third column is not annotation: the chunk that draws this table refuses to print it unless every value still matches the raw line beside it.

Column	Value	Where it came from
fips	42003	geo, field 10
county	Allegheny County	Gazetteer, joined on fips
state	Pennsylvania	the archive it came out of
total	1,250,578	seg 1, field 6
one_race	1,184,872	seg 1, field 7
white	946,321	seg 1, field 8
black	163,678	seg 1, field 9
aian	1,909	seg 1, field 10
asian	58,541	seg 1, field 11
nhpi	357	seg 1, field 12
other_race	14,066	seg 1, field 13
two_or_more	65,706	seg 1, field 14
hispanic	34,325	seg 1, field 78
not_hispanic	1,216,253	seg 1, field 79
nh_white	938,252	seg 1, field 81

Column	Value	Where it came from
nh_black	161,554	seg 1, field 82
nh_aian	1,305	seg 1, field 83
nh_asian	58,318	seg 1, field 84
nh_nhpi	304	seg 1, field 85
nh_other	6,064	seg 1, field 86
nh_two	50,456	seg 1, field 87
gq_total	34,557	seg 3, field 6
gq_correctional	2,151	seg 3, field 8
gq_juvenile	185	seg 3, field 9
gq_nursing	8,485	seg 3, field 10
gq_college	19,266	seg 3, field 13
gq_military	22	seg 3, field 14

Seven decisions turned the first three blocks into that row, and each one is a rule somebody chose.

A code decides which rows exist. 050 in field 3 is the only thing marking Allegheny County's row as a county. Pennsylvania's geographic header holds 529,085 rows, and 67 of them are counties. Everything else is thrown out on the strength of a three-digit code that the file never defines: every block, every tract, every borough. That is also why the archive is 57 MB compressed for one state: the counties are a rounding error inside a file built for drawing districts out of blocks.

The join key is not the FIPS code. Segment 1 does not carry 42003. It carries 0000003, a logical record number, and that number is meaningful only inside Pennsylvania's own archive: 0000003 is Allegheny County here and Honolulu County in Hawaii's. **A key that is unique inside one file and meaningless across files is the most common way a join goes quietly wrong.** That is why the build joins piece to piece before it joins anything to anything else.

The identifier is kept as text. 42003 could be stored as the number 42,003 without any loss. Do that to Alabama and 01001 becomes 1001, which is not a county. Nothing in these six states begins with a zero; the national list this file is matched against is full of states that do, so the column is character from the first line that reads it.

144 numbers arrive and 18 are used. The rest is the 63-way race combination table, published in full at county level for every county in the United States: White and Black; White and Asian; White, Black and Native Hawaiian, and so on. This chapter needs the six single races, the two summary lines and the Hispanic split, so it takes 18 fields by position and leaves 126 behind. **That is not a data-quality decision. It is a decision about what question is being asked,** and a chapter about multiracial identification would take the other 126 and leave these.

The file's own name field is discarded. Field 88 says Allegheny County, and the build does not use it. County names throughout this chapter come from the Census Gazetteer, matched on the FIPS code.

In these two states the swap changes nothing at all. The name field is filled in for every one of Pennsylvania's 67 counties and all 5 of Hawaii's, and all 72 of those names match

the Gazetteer exactly.

It is a habit rather than a repair. **One name list used everywhere in this book is worth more than the nearest name list used each time**, because names are what you join on when the codes run out.

The counts are converted from text to numbers, and that is a real step. Every field in the raw line is characters. 1250578 is a string until somebody decides it is a quantity. The build decides that for 18 columns and refuses to decide it for the FIPS code, and the difference between those two decisions is the whole of the previous point.

Then it checks the arithmetic it was given. Before writing anything the build asserts that every county's total equals one-race plus two-or-more, and equals Hispanic plus not-Hispanic. Both hold in all 503 rows. That is worth doing and it is worth knowing what it cannot catch: it verifies that the Bureau's tables are internally consistent, which they were always going to be. It says nothing about whether the field the build called `white` is the field the specification calls `white`.

Reading one field too far along produces a file that passes every check in this build. The only defence against that is the last section's kind of evidence: Allegheny's total matches the population the Bureau announced, to the person.

05 Two answers to how white a place is

Because Hispanic origin is asked separately, a Hispanic respondent also answers the race question — so a person can be, and millions are, both Hispanic and white. Which means there are two different and equally correct answers to a question that sounds like it has one.

State	Total population	% white (race question)	% non-Hispanic white (both questions)	Gap
New Mexico	2,117,522	51.0	36.5	14.5
Texas	29,145,505	50.1	39.7	10.4
New York	20,201,249	55.2	52.5	2.7
Pennsylvania	13,002,700	75.0	73.5	1.5
Hawaii	1,455,271	22.9	21.6	1.3
Mississippi	2,961,279	56.0	55.4	0.6




Figure 1. Two correct answers to “how white is this state?”, in six states. The dark bar counts people who answered the race question “white” and the origin question “not Hispanic”. The whole bar counts everyone who answered the race question “white”, whatever they said about origin. The light segment between them is the seam: people who are both. It runs from 0.6 points in Mississippi to 14.5 points in New Mexico, and it is almost a mechanical function of how Hispanic a state is.

New Mexico is 51.0% white by the race question and 36.5% white once you account for the ethnicity question – a gap of 14.5 points, a ratio of 1.40.

By one measure New Mexico is a majority-white state. By the other it is not close. Both numbers are correct. They answer different questions, and almost nobody reporting a figure like this says which one they used.

Across these six states the correlation between the gap and the Hispanic share is 0.982, because the gap *is* the Hispanic-and-white population. Mississippi, where very few people are both, shows 0.6 points.

So far this is only an accounting consequence of asking two questions instead of one. It becomes something else when you look at the people who answered the race question by rejecting every box on it.

06 The box for people who fit nowhere

State	% choosing ‘Some Other Race’	% Hispanic
New Mexico	15.0	47.7
Texas	13.6	39.3
New York	10.9	19.5
Pennsylvania	3.9	8.1

State	% choosing 'Some Other Race'	% Hispanic
Mississippi	1.9	3.6
Hawaii	1.8	9.5

15.0% of New Mexicans and 13.6% of Texans looked at the race question and chose "Some Other Race."

Who are they? The file answers directly, because it reports each race crossed with Hispanic origin.

Group	People	Share
Chose 'Some Other Race'	7,072,769	
of whom Hispanic	6,684,740	94.5%
of whom not Hispanic	388,029	5.5%

94.5% of the people who rejected every race box are Hispanic.

Read that again. The leftover category on the race question is in practice an ethnicity box. It is the box for people who fit nowhere else. In these six states, 7,072,769 people were asked their race, found the form did not offer one they recognized, and the overwhelming majority of them are the same people the *next* question was about.

The form asks a question a large minority of respondents do not think is well posed, and then records their refusal as though it were an answer.

That would be a contained problem if it were the only such box. It is not.

There is a second unusual column in this file: **two or more races**. Across these six states it moves enormously, from 25.3% in Hawaii to 3.7% in Mississippi. It is normally read as a story about mixed-race families, with Hawaii as the obvious place to look.

Before you read on, commit to a number. 8,576,792 people in these six states marked two or more races on a form that asked about race and about Hispanic origin in two separate places. **What share of those 8,576,792 also said they were of Hispanic origin?** Write down a percentage before you look at the next table.

07 The second hatch, used by the same people

Group	People	Share
Reported two or more races	8,576,792	
of whom Hispanic	6,083,462	70.9%
of whom not Hispanic	2,493,330	29.1%

Most readers write down something in the teens or twenties, because the mixed-race-families reading predicts a number like that. **70.9% of everyone reporting two or more races in these six states is Hispanic.**

Put the two questions side by side. Every person in these six states enters on the left with the race they gave, and leaves on the right with the answer they gave to the separate question about Hispanic origin.




Figure 2. Every race answer in these six states, split by the answer to the separate question about Hispanic origin. Each band on the left enters with the race a person gave, and leaves on the right with what they said about origin. Bands are ordered by how Hispanic the race answer turned out to be. The two at the top are the two residual categories, the boxes for what is left over: “Some Other Race alone” at 94.5% Hispanic and “Two or more races” at 70.9%. They carry almost all of their people into the same node. The next race answer down is 48.0%, and every one below it behaves nothing like the top two.

Those are national-ish totals for the six states together. Both hatches behave the same way inside every one of them, and the comparison worth making is against how Hispanic each state actually is:




Figure 3. Both hatches behave the same way inside every state, and neither behaves like the state. The open square is the share of a state's own population that is Hispanic; the two markers to its right are the share of the people using each residual category who are Hispanic. In Mississippi, only 3.6% of the state is Hispanic and 87.4% of the people there who chose "Some Other Race" are. The distance from square to marker is the size of the misfit, and it is large everywhere.

In every state the red circle sits far to the right of the square. The "Some Other Race" box runs from 80.2% Hispanic in Hawaii to 97.1% in Texas. Those same states' own Hispanic shares run from 3.6% to 47.7%. The second hatch does the same thing everywhere except Hawaii, the one state where reporting two races mostly means what it appears to mean.

And the counties:

County	State	Population	% two or more races	% Hispanic
Starr County	Texas	65,920	47.1	97.7
Brooks County	Texas	7,076	46.8	88.2
Zapata County	Texas	13,889	45.5	93.6
Maverick County	Texas	57,887	43.2	94.9
Hidalgo County	Texas	870,781	40.6	91.9
Cameron County	Texas	421,017	40.4	89.5
Webb County	Texas	267,114	40.2	95.2
Presidio County	Texas	6,131	39.9	81.4
Jim Hogg County	Texas	4,838	36.7	88.5

County	State	Population	% two or more races	% Hispanic
Kenedy County	Texas	350	36.3	74.6

The ten most multiracial counties in this file are all in Texas, and all of them are between 74.6% and 97.7% Hispanic. Starr County, on the Rio Grande, is 47.1% two-or-more-races and 97.7% Hispanic. Hawaii's most multiracial county, Hawaii County at 29.5%, does not make the list.

They are not scattered across Texas either. Shade every county in Texas and New Mexico by the share reporting two or more races:




Figure 4. Where the multiracial counties actually are. Every county in Texas and New Mexico, shaded by the share of its people reporting two or more races, with the 10 counties from the table above outlined in red. They are not scattered. They sit in a line along the Rio Grande, which is the Mexican border, in a state where the same counties are between 74.6% and 97.7% Hispanic.

Across all 503 counties the correlation between Hispanic share and multiracial share is 0.924; for "Some Other Race" it is 0.891.




Figure 5. The multiracial share of a county is largely a function of its Hispanic share.

One dot per county across all 503 in the file. Hawaii's counties are picked out in red, because Hawaii is the state most people name when asked where multiracial America is. They sit off the trend, low on the horizontal axis and high on the vertical one. They are the exception rather than the rule. Everywhere else the correlation is 0.924.

So the obvious reading is wrong twice over. "Two or more races" is not primarily a story about mixed-race families, and it is not primarily about Hawaii. **It is the second escape hatch from the same badly-posed question**, used by the same population, on the same form. The 2020 questionnaire invited people to write in detail. Hispanic respondents who had previously ticked "Some Other Race" alone often wrote entries that the Bureau coded as two races, and the multiracial count in these six states rose accordingly.

That has a hard consequence. **Comparing "two or more races" between 2010 and 2020 compares instruments, not people.**

08 Counted where you sleep

The form makes one more decision about a person, and it is not about what they are. It is about where they are.

County	State	Population	In a correctional facility	% of the county
Forest County	Pennsylvania	6,973	2,653	38.0
Issaquena County	Mississippi	1,338	493	36.8
Hudspeth County	Texas	3,202	996	31.1
Concho County	Texas	3,303	1,012	30.6
Garza County	Texas	5,816	1,731	29.8
Jones County	Texas	19,663	5,073	25.8

38.0% of Forest County, Pennsylvania is in a state correctional facility – 2,653 people out of 6,973.

Those people cannot vote in Pennsylvania. They are counted as residents of Forest County for the purpose of drawing districts, and represented by a legislator they had no part in choosing. The communities they came from, mostly urban ones, lose the population that would have counted toward their own representation.

Note that every county at the top of that list is small. The percentage is driven as much by the denominator as by the facility.

The same residence rule catches students, without the disenfranchisement:

County	State	Population	In student housing	% of the county
Tompkins County	New York	105,740	13,084	12.4
Claiborne County	Mississippi	9,135	1,005	11.0
Waller County	Texas	56,794	5,814	10.2
Oktibbeha County	Mississippi	51,788	5,160	10.0

The two lists do not overlap, and the two categories point in opposite directions:




Figure 6. The same residence rule, pointed in opposite directions. The counties with the largest shares of their population in a correctional facility (left of the line) and in student housing (right of it). The two lists do not overlap. Everyone counted on the left is a resident of a county they cannot vote in; everyone counted on the right can vote where

they are counted. Every county here is small, so the percentage is driven as much by the denominator as by the institution.

Which is the useful comparison. Prison gerrymandering combines counting people where they are held *and* denying them the vote there. Colleges have the first without the second, and almost nobody objects. **The objection rests on the combination**, and saying so is more precise than objecting to either half.

09 Discovered, or constructed, or scheduled

The dispute underneath all of this is not really about accuracy. It is about whether the census *discovers* race or *constructs* it, and the federal government has already answered in the standard itself: the categories are socio-political constructs, not scientific or anthropological ones. That language is OMB's rather than the Bureau's, which matters, because it means the disclaimer binds every agency using these categories and not just the one that collects them.

The scholarly answer is more settled than the political one. Nobles (2000), in a history of the American and Brazilian censuses, shows racial categories being added, merged and dropped in response to political pressure in both countries. The book argues that census bureaus have never been bystanders in that process. Kenneth Prewitt (2013), who ran the Bureau, goes further and argues the race question should be dropped outright. The counter-argument is practical rather than philosophical: no count, no civil-rights enforcement. Abolishing the measure abolishes the remedy.

That argument was, in the ordinary run of things, unresolvable. In March 2024 it was partly resolved by fiat. The Office of Management and Budget revised Statistical Policy Directive No. 15, its first revision since 1997. The revision merges race and Hispanic origin into a **single** question and adds a Middle Eastern or North African category, with agencies required to comply by 2029.

That will end the two-answers problem drawn in Figure 1. It will also make the comparison in Figures 2 to 5 permanently impossible to extend forward. **Both effects are certain, and reasonable people weight them differently.** That is the actual trade, and it has been made. The seam this chapter measures has a scheduled end date.

There is also the noise. For 2020 the Bureau applied *differential privacy*: small deliberate adjustments designed so that no individual's answers can be reconstructed from published tables. State totals are exact and county figures are close, but districts are drawn from blocks, and the smaller the geography the more the injected noise matters. Alabama sued to stop it and lost — its motion for a preliminary injunction was denied and the case dismissed in September 2021. Kenny and colleagues (2021), evaluating the system on actual redistricting tasks, found the distortion modest for most purposes and non-trivial for some, which is why the argument continues rather than ends.

10 What the count will bear

The strengths of this source are the strongest in this field. It is the legally operative count, published by the agency that performed it, and it is the file states actually redistricted from. It is internally consistent, and both of its decompositions were tested above against every county rather than assumed. And because it reports race crossed with Hispanic origin, it can be used to interrogate its own categories, which is something most datasets cannot do at all.

The limits are equally definite. This is six states rather than fifty, chosen precisely because race and ethnicity pull apart in them, so every share above is unrepresentative by design. Differential privacy noise is present and unquantified here: it was applied before publication and the file carries no column saying by how much. The categories are not comparable across decades, because the 2020 form invited write-in detail the 2010 form did not and the coding changed with it.

Concretely: 2020 added write-in areas and examples under the White and Black categories, and examples under American Indian or Alaska Native. It collected specific Hispanic origins within the broader category. It also dropped the word “Negro” that the 2010 form still carried.

Two more limits. These are counties, while the geography where both the privacy noise and the districting actually bite is two levels smaller. And it is a count of residents rather than citizens, which is what apportionment uses and what attempts to change have so far failed in court.

What can be said, then, is this. The count is exact, certified, and it breaks apart without error. It still supports two different and equally correct answers to “how white is New Mexico?”: 51.0% and 36.5%. And, more sharply: the race question has two leftover categories, and in these six states 94.5% of the people in one and 70.9% of the people in the other are Hispanic.

The form has two escape hatches and one population is using both. That is a measurement of misfit, not a description of who Americans are. The box is used by Hispanic respondents *because* the form separates race from origin. So the finding is evidence about the form, rather than about the people answering it.

Three questions stay open. Whether a combined question is an improvement: it will fit respondents better and destroy the trend series, and there is no version of the change without a cost. Whether “Some Other Race” should have been retained at all — the Bureau has repeatedly proposed removing it, and the population using it is precisely the population §2 and §203 most concern. And how much differential privacy costs at block level, which is the only level at which the question really matters and the only level this file cannot show.

The instruction printed on the 2020 form was an answer to a hard question, and it was not an unreasonable one. In these six states alone, 7,072,769 people read it and reached for the box marked *Some Other Race*, and 8,576,792 more wrote something the Bureau coded as two races. Both boxes now sit in the file that districts are drawn from, carrying a population the form could not name into lawsuits about whether that population is fairly represented.

11 Sources

Data

- U.S. Census Bureau, *2020 Census Redistricting Data (P.L. 94-171) Summary File*, legacy format, tables P1, P2 and P5, for Pennsylvania, Texas, New York, Hawaii, Mississippi and New Mexico. https://www2.census.gov/programs-surveys/decennial/2020/data/01-Redistricting_File--PL_94-171/ County names joined from the Census Gazetteer, because the file's own header populates the name field inconsistently.
- U.S. Census Bureau, *2020 Census Questionnaire*. The two questions quoted at the head of this chapter are its questions 8 and 9, and the sentence between them is the note the Bureau printed to explain why both had to be answered.
- County outlines in Figure 4 are drawn from the `maps` package's county boundary data, matched to the file by FIPS code.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`p194171_counties.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `p194171_counties.csv` carries both `white` and `nh_white`. Take the difference for every county and sort it. Where does the choice between those two columns change the answer most?
- `other_race` and `two_or_more` are the boxes people reach for when none of the offered ones fit. Work out their combined share county by county. Which places does the form fit worst?
- `hispanic` and `not_hispanic` should add up to `total`. Check that they do. Then explain to somebody who has not read this chapter why Hispanic origin is missing from the race columns.
- The file also counts people in group quarters. Find the counties where they are a large share of the total, and say what that does to a district drawn from this file.

Literature and law

- Nobles, M. (2000). *Shades of Citizenship: Race and the Census in Modern Politics*. Stanford: Stanford University Press.
- Prewitt, K. (2013). *What Is Your Race? The Census and Our Flawed Efforts to Classify Americans*. Princeton: Princeton University Press.
- Office of Management and Budget (1997). "Revisions to the Standards for the Classification of Federal Data on Race and Ethnicity." Statistical Policy Directive No. 15, *Federal Register* 62: 58782 (30 October 1997) – the standard governing the 2020 census: five race categories, two ethnicity categories, two separate questions with ethnicity asked first, more than one race permitted, and the definitions quoted above.
- Office of Management and Budget (2024). "Revisions to OMB's Statistical Policy Directive No. 15." *Federal Register* 89: 22182 (29 March 2024).
- Voting Rights Act § 2, 52 U.S.C. § 10301, and § 203, 52 U.S.C. § 10503 – the enforcement the race and Hispanic-origin questions exist to serve.

- National Conference of State Legislatures, *Redistricting Law 2020* (Denver: NCSL, October 2019), ch. 1, for the OMB standards as they were applied to the 2020 form, the changes from 2010, and the account of why the census has asked these two questions since 1980.
- Kenny, C. T., Kuriwaki, S., McCartan, C., Rosenman, E. T. R., Simko, T. and Imai, K. (2021). "The Use of Differential Privacy for Census Data and Its Impact on Redistricting: The Case of the 2020 U.S. Census." *Science Advances* 7(41): eabk3283.
- *Alabama v. U.S. Department of Commerce*, 546 F. Supp. 3d 1057 (M.D. Ala. 2021).

Further reading

- Scott, J. C. (1998). *Seeing Like a State: How Certain Schemes to Improve the Human Condition Have Failed*. New Haven: Yale University Press. **On the design of this material.** The structure follows Moon Duchin's PUBPOL 2130, *Data and the State*, at Cornell, adapted from Python to base R: race, then ethnicity as a second dimension, then the question of what the categories accomplish.

How the figures were made

The legacy-format P.L. 94-171 files for the six states are read, tables P1, P2 and P5 kept, the county names joined, giving the single `p194171_counties.csv` this chapter draws on. Every number printed above, including the two decomposition checks, is computed from that file at the moment this document was rendered.

REBUILD THIS DATA YOURSELF, WITH AN AI ASSISTANT

I want to rebuild a public dataset from its original source, without any starting files. Please do the following and show your work.

THE SOURCE. U.S. Census Bureau, 2020 Census Redistricting Data (P.L. 94-171) Summary File, legacy format – the file the Bureau is required by law to hand every state so it can draw districts. One zip per state under

https://www2.census.gov/programs-surveys/decennial/2020/data/01-Redistricting_File-PL_94-171/

No account or key is needed. Fetch six states – Pennsylvania, Texas, New York, Hawaii, Mississippi and New Mexico – each from a folder of that name (two-word states use an underscore), for example

Pennsylvania/pa2020.pl.zip. Together they are roughly 300 MB down for a few hundred rows out; that ratio is normal for this file. Each zip holds four pipe-delimited files with no headers: a geographic header, one row per geographic unit, and three data segments that join to it on the LOGRECNO column. Counties are summary level 050. In the first segment, Table P1 (race) starts at field 6 and Table P2 (Hispanic or Latino origin by race) at field 77; in the third segment, Table P5

(group quarters) starts at field 6. Read every code as text – FIPS codes lose their leading zeros the moment they become numbers.

WHAT TO BUILD.

1. One county table covering all six states: total population, the six race-alone counts, two or more races, the Hispanic and non-Hispanic breakdown of each, and the group-quarters counts (correctional, college, military, nursing).

CHECK YOUR WORK. The copies this book used were fetched in August 2026. If your rebuild matches them, you should find:

- 503 county rows: 67 in Pennsylvania, 254 in Texas, 62 in New York, 5 in Hawaii, 82 in Mississippi and 33 in New Mexico
- Honolulu County, Hawaii: 1,016,508 people, of whom 436,853 are Asian alone

These are the final legal files of the 2020 census; nothing in them will change before the 2030 cycle.