

Reading the ANES Cumulative File

Seventy-six years of survey answers, and three columns that run backwards

84-355 Democracy's Data

August 2026

The American National Election Studies has asked Americans about politics in every presidential election since 1948, and in most midterms since. The cumulative file gathers seventy-six years of it into one table: 73,745 interviews, 1,030 columns, 163 megabytes of numbers. It is the longest continuous record of American political behaviour that exists.

It is also a file in which **three of the columns this chapter uses run in the opposite direction from the one you would assume**, and nothing in the data says so. Two of the three fooled the person who built this chapter. The resulting figures were not broken — they were readable, plausible, and wrong, which is the only kind of error worth writing a chapter about.

01 Where the file comes from, and what it is for

ANES is run by Michigan and Stanford under a National Science Foundation grant. Its purpose is not to report the news and never has been: it exists so that the same question, asked the same way, can be compared across decades. That single design decision is what makes the cumulative file possible and what makes it strange.

Why it fits this book. Nothing else reaches back to 1948 with individual respondents attached. Totals like vote counts and party divisions tell you what a country did. ANES is one of the few sources that records what individual people said about it, and it is the only one that has been doing so since Truman.

What it costs to get. The file is requested, not handed over. The download page returns 403 to a script and 200 to a browser, so no build script can fetch it unattended; a person has to go and get it. It is not redistributed with this chapter, both because it is 163 MB and because ANES asks that people take it from ANES. The build script looks in two places and stops with the URL if it finds nothing.

What has already been done to it. A great deal, and this is the important part. The cumulative file is not a pile of surveys; it is a *reconciliation*. ANES staff recoded seventy-six years of differently-worded, differently-ordered, differently-scaled questions so that one VCF number means one thing throughout. The file is comparable across time only because people made it so, case by case.

The codebook is the record of those decisions. ANES states plainly that the data cannot be used with any codebook but its own. That is not a formality. It is the whole warranty.

02 What arrives

Five random rows — not the first five, because the file is ordered by study year and the first five are all interviews from 1948.

ANES Time Series Cumulative Data File, 20260205 CSV distribution
73,745 rows x 1,030 columns

The first 14 of 1,030 column names, exactly as they arrive:

Version VCF0004 VCF0006 VCF0006a VCF0009x VCF0010x VCF0011x
VCF0009y VCF0010y VCF0011y VCF0009z VCF0010z VCF0011z VCF9999

Five random rows, the columns this chapter reads, values untouched:

VCF0004	VCF0301	VCF0310	VCF0050a	VCF0050b	VCF0806	VCF0838	VCF9039	VCF9040	VCF9041	VCF9042	VCF9043
1970	4	3	NA	1	4	NA	NA	NA	NA	1.0000	
1972	6	3	3	NA	0	NA	NA	NA	NA	1.0000	
1984	3	2	3	3	9	2	NA	NA	NA	1.0000	
1994	2	2	NA	3	2	2	2	3	2	2	0.9981
2016	7	3	3	3	7	2	2	1	2	2	1.5380

Every -1, 0 and 9 above is a code, not a measurement.
The codebook is the only way to know which is which.

Every value above is a number, and only some of them are measurements. On this evidence alone you cannot tell which is which, and no amount of staring at the file will tell you. That is what the codebook is for.

03 The codes that are not measurements

Take one column. VCF0806 asks where the respondent stands between a government insurance plan and a private one, on a scale of one to seven. The column also contains -1, 0 and 9.

Those are not positions on the scale. They record *inapplicable*, *no post-election interview*, and *don't know*. Together they are **20.80% of every answer in the column** — 10,063 of 48,381.

Now average the column four ways.

What was removed	Mean	Answers used	Distance from correct
Nothing removed, the file as it arrives	3.966	48,381	+0.137
Only 9 (don't know) removed	3.380	43,337	-0.449
Only 0 and -1 (inapplicable) removed	4.431	43,362	+0.602
Every non-scale code removed	3.829	38,318	—

The first row is the trap and the middle two are the lesson. Doing nothing gives 3.966, and the correct answer is 3.829. That is off by only 0.137 on a seven-point scale, which is small enough to survive any sanity check you might apply to it.

It is small for a reason that should frighten you. The 9s are larger than the largest real answer, so they drag the mean up. The 0s and -1s are smaller than the smallest, so they drag it down. There are 5,044 of the first and 5,019 of the second, and **the two errors very nearly cancel**.

Clean half the problem and the cancellation breaks. Remove only the don't-knows and you get 3.380. Remove only the inapplicables and you get 4.431. Both are three to four times further from the truth than leaving the column entirely alone.

Before you read on, commit to a view. You have inherited a script that recodes *some* of a survey's missing-value conventions and misses others. Is its output better than the same script with the recoding stripped out altogether, or worse? **Say which, and say what you would need to look at to find out.**

04 The axis nobody was asked about

The most reproduced figure in the study of public opinion plots an attitude against **political awareness**. No respondent has ever been asked how aware they are. There is no such question and no such column.

Awareness is *built*. The usual materials are the interviewer's judgment of how informed the respondent seemed, and the respondent's own account of how closely they followed the campaign. Two defensible builds from the same file:

How it was built	Respondents	Years	Distinct levels
Interviewer rating only (VCF0050a, falling back to VCF0050b)	41,635	1966–2016	5
Interviewer rating plus campaign interest (VCF0310)	38,773	1966–2016	9

Neither is more correct than the other. They cover different people, they have different resolution, and a reader who is shown only the finished figure has no way to know which was used or that a choice existed.

The part that nearly ruined this chapter

The interviewer's rating is coded like this, and the coding exists only in the codebook:

Variable	Code	What the code means
VCF0050a — respondent level of political info	1	Very high
	2	Fairly high
	3	Average
	4	Fairly low
	5	Very low

One is the most informed respondent, not the least. The scale descends.

Read it the intuitive way and every figure in the next section comes out inverted: the party gap appears to *shrink* as awareness rises. That inverted result is not obviously broken. It reads as a finding: *information moderates partisanship, exposure to argument brings people together*. It is comfortable enough that somebody might publish it. It held across every study year tested, which made it look robust rather than wrong.

Nothing in the CSV says which end is which. The direction of a scale is not a property of the data; it is a fact about the data, kept somewhere else.

05 Two ways awareness works

With the scale the right way round, here is opinion by awareness, by party.




Figure 1. Mean position by constructed awareness, pooled across every study year in which each item was asked. Left: government health insurance, 1 = a government plan, 7 = a private plan. Right: abortion, 1 = never permitted, 4 = always permitted. Party is the seven-point identification collapsed to leaners.

Hovering an awareness band gives every party's mean and its number of respondents. Hovering a party name below dims the others in **both** panels, so the two items can be compared directly. That is the comparison the figure exists to support.

On health insurance the lines **spread**. Among the least aware the gap between Republicans and Democrats is 0.79 points; among the most aware it is 2.04 points, and it widens at every step in between.

Abortion does something better. Among the least aware the gap is 0.06 — indistinguishable from nothing. Among the most aware it is 0.62, **and the sign has flipped**. There is no partisan difference on abortion among Americans who pay little attention to politics. The difference is manufactured further up.

This is what the awareness axis is for. A party gap that grows with attention is not a fact about what Democrats and Republicans are like. It is a record of how far down the elite disagreement has travelled.

A second prediction. Both panels widen. Suppose you found an issue where the lines converge — where the most aware partisans agree and the least aware disagree. **What would have to be true of the country's politicians for that to happen?**

06 The same trap, twice more

The racial resentment battery is four questions asked together since 1986. All four use the same response scale, one for agree strongly through five for disagree strongly. Two of them mean the opposite of the other two:

Variable	Item	Agreeing means
VCF9039	Conditions make it difficult for blacks to succeed	less resentment
VCF9040	Blacks should not have special favors to succeed	more resentment
VCF9041	Blacks must try harder to succeed	more resentment
VCF9042	Blacks gotten less than they deserve	less resentment

Average the four without reversing the middle two and the directions cancel almost perfectly: every party lands within five hundredths of every other, in every year, for thirty-eight years. That is a spectacular finding: *no partisan difference in racial attitudes, ever*. It is arithmetic on mixed-direction items, nothing more.

Reversed properly, the same four questions give this.




Figure 2. Mean racial resentment by party identification, 1986–2024. Higher is more resentful. The four items are reversed where needed so that all point the same way; without that step every line lies on top of every other. Hovering gives the Republican-minus-Democrat gap at each year as well as the three levels. The chapter's claim is about the gap and about which side moved, and neither is reliably read off a pair of printed lines.

The gap between Republicans and Democrats was 0.34 points in 1986 and 1.44 in 2024. It more than quadrupled. And the movement is asymmetric in a way the figure makes unmissable: Republicans end close to where they started, while Democrats fall by 0.82 points. The parties did not walk away from each other. One of them walked.

07 What the file will not let you ask

A cumulative file is a set of decisions about what deserves to be carried forward, and those decisions bound every question you can put to it.

A question must have been asked three or more times to be included at all. Anything asked once or twice is not here, and that includes most questions about anything new. The file is excellent on what has been asked for decades and silent on everything else, and that silence looks exactly like the thing not mattering.

Coverage is uneven, and the gaps are not where you would guess.

Variable	Question	Years	Studies	Answers
VCF0009z	Sample weight	1948–2024	33	73,745
VCF0301	Party identification, 7 point	1952–2024	32	73,083
VCF0310	Interest in the campaign	1952–2024	30	70,124
VCF0050b	Interviewer rating of political information (b)	1966–2016	20	37,983
VCF0050a	Interviewer rating of political information (a)	1968–2016	14	29,773

Variable	Question	Years	Studies	Answers
VCF0806	Government health insurance, 7 point	1970–2024	16	48,381
VCF0838	Abortion, 4 point	1980–2024	17	48,086
VCF9039	Conditions make it difficult for blacks to succeed	1986–2024	12	39,802
VCF9040	Blacks should not have special favors to succeed	1986–2024	13	41,083
VCF9041	Blacks must try harder to succeed	1986–2024	12	39,802
VCF9042	Blacks gotten less than they deserve	1986–2024	13	41,083

Only 9 of these 11 variables reach 2024. **The interviewer rating stops in 2016**, and it is the raw material of the awareness axis. The measure behind the most reproduced figure in the field is missing for the most recent three studies. A chapter wanting to extend Figure 1 to the present could not do it from this file.

And question wording drifts. ANES documents the changes, because a question asked identically in 1972 and 2024 may still not mean the same thing to the person hearing it. The file gives you comparability; it cannot give you equivalence.

08 What this chapter cannot tell you

The awareness bands pool every study year together, which means Figure 1 mixes eras whose elite politics differed. A version done properly year by year would show the pattern arriving at different times for different issues, and that is the actual argument in the literature; this is the shape, not the history.

Nothing here is weighted. VCF0009z is in the file and unused, so every number in this chapter describes the people ANES interviewed rather than the country. For the comparisons made here, which are gaps between groups within the same survey, that matters less than it would for a level. It is still a choice, and you should know it was made.

Racial resentment is a measure with a large critical literature attached to it, and reasonable people disagree about what it captures. This chapter takes no position on that. It shows what happens to four questions when two of them are pointed the wrong way, which is a claim about arithmetic, not about the concept.

09 Sources

- American National Election Studies. *ANES Time Series Cumulative Data File (1948–2024)*, version of 5 February 2026. Ann Arbor, MI and Palo Alto, CA: University of Michigan and Stanford University. <https://electionstudies.org/data-center/anes-time-series-cumulative-data-file/>
- ANES. *Time Series Cumulative Data File Codebook and Appendices*, same release. Every variable coding quoted in this chapter is read from the codebook, including the one that runs backwards.
- Zaller, John R. (1992). *The Nature and Origins of Mass Opinion*. Cambridge University Press. The awareness axis, and the mainstream and polarization patterns, are his.

- Tesler, Michael (2016). *Post-Racial or Most-Racial? Race and Politics in the Obama Era*. University of Chicago Press. Figure 2 is a much-simplified version of a relationship he documents at length.
- Converse, Philip E. (1964). "The Nature of Belief Systems in Mass Publics," in David E. Apter, ed., *Ideology and Discontent*. Free Press.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`awareness.csv`, `codes.csv`, `coverage.csv`, `resentment.csv`, `zaller.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `codes.csv` shows the same quantity computed with and without the non-answer codes removed, with an `error` on each. Work out how big the mistake is. Would you have noticed it in a published table?
- `coverage.csv` says which years each variable actually spans. Find a variable you would want for a fifty-year trend and check whether it runs the whole way. What breaks if it does not?
- `zaller.csv` splits each item by party and by awareness. Read across the awareness levels within one party. Do better-informed partisans agree more with their own side, or less?
- `resentment.csv` runs one measure by party across the years. Find where the two parties diverge. Then check `coverage.csv` to see whether the question was even asked the same way throughout.

WHY THIS DATA CANNOT BE FETCHED AGAIN

The file behind this chapter comes from a source no script can fetch for you, and it cannot be passed along to you either.

WHY NOT. The chapter is built on the American National Election Studies Time Series Cumulative Data File, 1948–2024, in the version of 5 February 2026 – every ANES election survey since 1948, recoded by staff so that the same variable number asks the same question across seventy years of studies. Its download page,

<https://electionstudies.org/data-center/anes-time-series-cumulative-data-file/>

answers 403 to a script and 200 to a browser: the wall judges the client, not the request, so no assistant can fetch the file unattended. Behind the wall sit a free registration and a 22 MB zip holding a 163 MB CSV. ANES asks that everyone take the file from ANES, so it is not redistributed here.

WHAT EXISTS INSTEAD. Register yourself – it is free and takes a few minutes – and download the CSV distribution in a browser, along with the two codebook PDFs that ship beside it. From that moment on an assistant can do everything: read the file, look codes up in the codebook, redo every number. The only step no assistant can do for you is the download.

WHAT YOU CAN STILL CHECK. Once you hold the file, these numbers say whether it is the same object:

- 73,745 rows and 1,030 columns in the CSV
- the government health insurance scale, VCF0806, carries 48,381 non-empty values; its mean is 3.966 as the column arrives and 3.829 once every code that is not a real answer is removed
- both interviewer ratings of political information, VCF0050a and VCF0050b, stop in 2016, eight years before the file ends