

Sixty Thousand People, and Sixty-Eight Vermonters

What the largest political survey in America buys, where it runs out, and the difference between being precise and being right

84-355 Democracy's Data

August 2026

In 2024 the Cooperative Election Study interviewed 60,000 Americans. That is roughly sixty times a standard national poll. The thing usually said about a sample that large is that its **margin of error** is negligible: how far off the answer could be from the luck of who happened to get drawn.

Nationally, it nearly is. The trouble is that a national sample is not a national sample everywhere. **The same 60,000 respondents are 3,618 Texans and 68 Vermonters**, and the second number decides what the survey can say about Vermont.

01 Before you read on

60,000 people, spread across 51 jurisdictions.

How many of them live in the smallest state cell? Write down a number.

One number. Most people put it an order of magnitude too high.

02 The file, and the door

This is the 2024 CES Common Content: 60,000 respondents, 694 variables, 175 MB. Of them, 43,749 named one of the two major presidential candidates; the other 16,251 named somebody else, did not vote, or did not answer.

It cannot be fetched by a script. The file sits on Harvard Dataverse, which answers an automated request with a success code and an empty body – a success-shaped refusal this book's survey access chapter measures directly. A person with a browser fetched it, and it is preserved in `./ces-class/data/raw/`, which this chapter reads rather than keeping a second copy.

Two sibling chapters supply what the CES does not. The certified 2024 returns come from historical campaigns. The **crosswalk** comes from data sources, a table that matches one set of codes to another. It turns the numeric state codes the CES stores into the abbreviations the returns use. Neither file is typed here.

03 Nationally, it works

The CES puts the two-party Democratic share at **47.79%**. The certified national result was **49.25%**.

It misses by 1.46 points.

That is a good result for a survey and a bad result for the usual explanation of survey error. A sample of 43,749 carries a national margin of about half a point. The miss is three times that, so most of it is something other than sampling.

That miss is roughly three times the national sampling margin. Whatever produced it is not the sample size.

04 Where it runs out

State	Respondents	Margin ±	CES says	Result was	Error
Vermont	68	15.3	66.5	66.4	+0.1
Wyoming	77	13.7	24.6	26.5	-1.9
North Dakota	85	13.3	50.7	31.3	+19.4
Alaska	106	15.8	45.6	43.2	+2.4
District of Columbia	110	10.5	79.6	93.3	-13.7
...					
Texas	3,618	2.7	42.9	43.1	-0.2
California	3,429	2.8	59.0	60.4	-1.4
Florida	3,106	3.5	46.6	43.4	+3.2

Vermont has 68 respondents. Texas has 3,618, which is 53.2 times as many. The margin of error runs from 2.71 points in the best-sampled state to 15.81 in Alaska.

The median state's margin is 6.20 points. A survey described in every write-up by one national sample size is, at the level most people want to use it, an ordinary small poll fifty-one times over.

The weights cost more than half of it

The respondent counts above are not the precision the survey actually has. **Weighting** means counting some answers more heavily than others so the sample resembles the country. It costs precision to do that, and the cost has a name: the **design effect**, the factor by which the effective sample is smaller than the headcount.

The median state's design effect is **2.78**. A state cell of 300 respondents carries roughly the precision of a simple random sample of about 110. Every margin above already accounts for that; the point is that the number in the press release does not.

05 Precision is not accuracy

Now grade every state against what it actually did.




Figure 1. Every state twice: the margin of error it claims, and the error it actually made against the certified result. Filled points are red where the result falls outside the margin.

Two correlations say the whole thing.

A state's sample size predicts its **margin of error** at -0.75 . It predicts **how wrong the state actually was** at -0.25 .

Sample size buys precision. It barely buys accuracy. The hollow points fall away steeply to the right; the filled ones scatter almost as widely at 3,000 respondents as at 300.

06 The margin does not cover the error

If sampling were the only thing going wrong, about 95% of states would have the certified result inside their margin. Sampling error is what a margin of error describes, and nothing else.

45 of 51 do, which is 88.2%, against the 48.4 states that 95% coverage predicts.

The median state misses by 2.49 points. The worst, North Dakota, misses by 19.41.

That shortfall is everything a margin of error does not describe. It covers who answered and who did not, whether the people in a small state's cell resemble that state, and whether a person who says they voted for somebody did. The weighting chapter takes one of those apart on a single poll.

07 What this chapter cannot tell you

It cannot tell you the CES is bad. It is among the best survey instruments in American political science, and a 1.46-point national miss is a good result. The chapter is about what a large sample does and does not buy, and the CES is the right file precisely because it is large and careful.

One question, one year. Presidential vote in 2024, which is the most-asked question in the survey and the easiest to grade. Nothing here says how the CES performs on questions with no certified answer, which is most of them.

The benchmark is not the same population. The certified result counts everyone who voted. The CES estimate counts people who said they voted a particular way, weighted to a model of the electorate. Some of the gap in Figure 1 is that difference rather than survey error.

And the design effect is estimated, not published. The Kish formula used here reads the weights and returns the precision they imply. A design-based standard error from the survey's own documentation would differ, probably not by much, and this chapter does not have one.

08 What it does establish

A survey of 60,000 people estimates the national two-party vote to within 1.46 points, and estimates the median state to within 2.49.

Its smallest state cell is 68 people. The largest is 53.2 times bigger, the median state's margin is 6.20 points, and the weights cost a factor of 2.78 in effective sample before any of that.

And the number that a bigger sample reliably improves is the margin of error, not the error. -0.75 against -0.25. A survey can be more and more precise about a number that stays the same distance from the truth.

09 Sources

- **Cooperative Election Study**, Common Content 2024, Data Release 2, Ansolabehere, Schaffner and Pope. doi:10.7910/DVN/X11EP6, CCES24_Common_OUTPUT_vv_topost_final.csv. 60,000 respondents. <https://doi.org/10.7910/DVN/X11EP6> The copy read here is the one the ces-class chapter keeps.
- **Not scriptable.** Harvard Dataverse answers automated requests with a success code and an empty body.
- **Certified 2024 returns and the state crosswalk** are read from the historical campaigns and data sources chapters rather than retyped.
- **Build.** data/build-data.R verifies the vote coding against the certified national result before proceeding, then writes the two tables this chapter reads. Every number above is computed at the moment this document was rendered.

Checks

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`states.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `states.csv` has `n` next to `n_eff` and `deff`. Sort by the ratio of the two. What is a large sample actually worth once the design is accounted for?
- `states.csv` also has `moe`, `error` and `covers`. Count the states where `covers` is `FALSE`. Then ask what a margin of error is promising, and to whom.
- Sort `states.csv` by `error` rather than `moe`. Is the state with the widest margin the state that got the answer most wrong?
- Find the smallest `n` in the file. Work out what you would be willing to say about that state, and write the sentence you would actually publish.

Check	Passed
the whole common content arrived	yes
every row carries the common weight	yes
the sibling chapter's county file, which carries the state crosswalk, is present	yes
every respondent's state code resolves to a state	yes
the sibling chapter's returns are present	yes
the sibling chapter's national series is present	yes
exactly two major-party rows for 2024	yes
the vote coding is not reversed – the estimate is near the result	yes
the largest state cell is at least twenty times the smallest	yes
weighting costs precision everywhere	yes
the margin of error does not account for most of the error	yes
sample size buys precision far more than it buys accuracy	yes