Everyone Guesses Forty

What Americans think each party is made of, how wrong that is, and the one thing they are most wrong about

84-355 Democracy's Data

August 2026

Out of every hundred Republicans, how many do you think earn more than $250,000 a year?

The average American answer is **38.2**. The real figure is **2.2**.

That is the most quoted number from a study of what people believe the two parties are made of, and it is one of eight. This chapter shows all eight, then shows the thing the famous figure cannot: **who was doing the guessing**, and how much that changes the answer.

01 The study

In March 2015 YouGov put eight questions to 1,000 Americans. For each party it named four groups that party is stereotyped around, and asked what percentage of that party's supporters belong to it. Four Democratic groups, four Republican ones, 8,000 answers in all.

The study was run for two researchers, Douglas Ahler and Gaurav Sood, who published the results in 2018. **The survey itself is theirs and YouGov's, and nobody else has it.** What is public is a copy of part of it, which they deposited so their results could be checked. That is a different object from the study, and the last section of this chapter measures the difference.

02 Everyone is wrong, and by a lot

Party	Out of every 100 supporters, how many are...	Really	People said	Gap
Democratic	Union member	10.5%	39.3%	+28.8
Democratic	Lesbian, gay or bisexual	6.3%	31.7%	+25.4
Democratic	Atheist or agnostic	8.7%	28.7%	+20.0
Democratic	Black	24.0%	41.9%	+17.9
Republican	Earning $250K+	2.2%	38.2%	+36.1
Republican	Aged 65 or over	21.3%	39.1%	+17.8
Republican	Evangelical	34.3%	41.6%	+7.3
Republican	Southern	35.6%	40.4%	+4.7

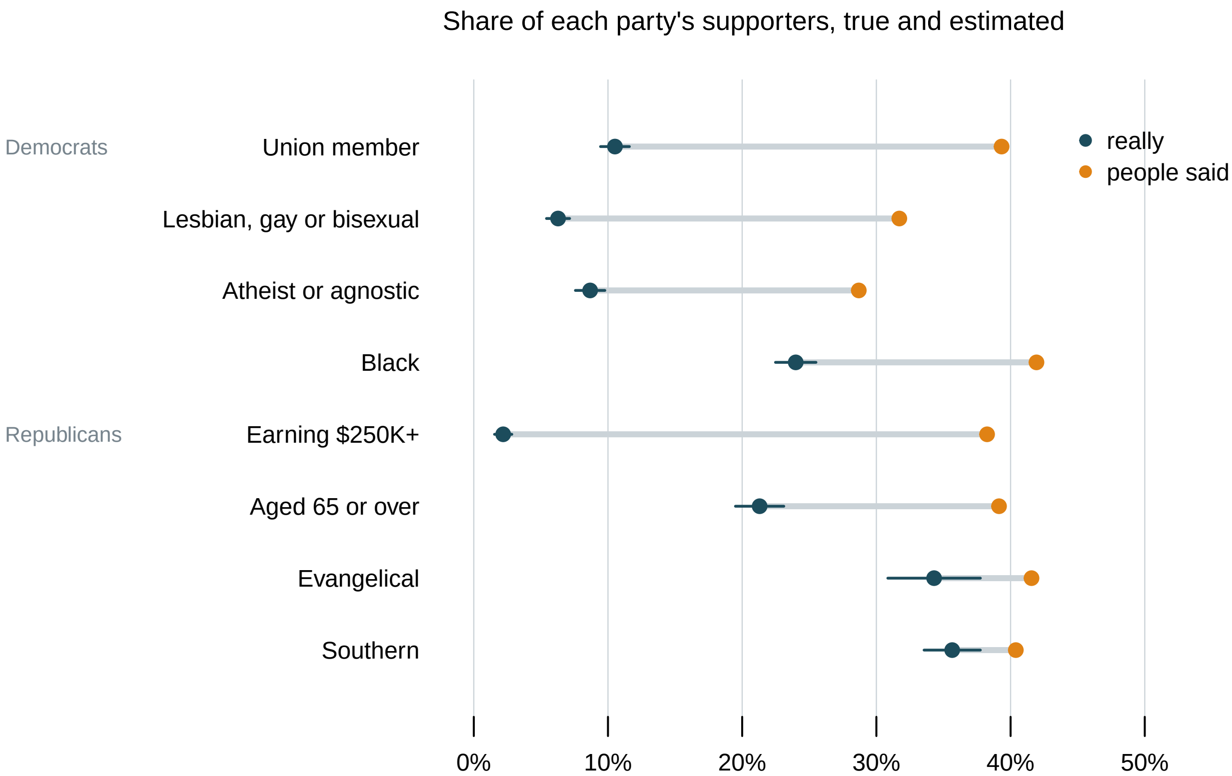

Figure 1. The eight items. The dark dot is the true share, the light dot is the average guess, and the bar joining them is the gap.

Every single group is overestimated. The average gap is **19.8 points**.

The sizes are not alike. Earning $250K+ is out by **36.1 points**, which is the number everybody quotes. Southern is out by **4.7**, which nobody quotes. They are in the same figure.

The pattern in those sizes is worth having before anything else. The groups people are most wrong about are the rarest ones. Lesbian, gay or bisexual Democrats are 6.3% of the party and people said 31.7%. Southerners are 35.6% of the Republican party and people said 40.4%. Where a group really is common, the guess is close.

03 And you are worse about the party you are not in

This is the study's actual argument, and Figure 1 cannot show it. Everyone in that figure is averaged together: Democrats, Republicans and people who are neither, all guessing about both parties at once.

The file records each person's own party, so the same eight items can be asked twice.

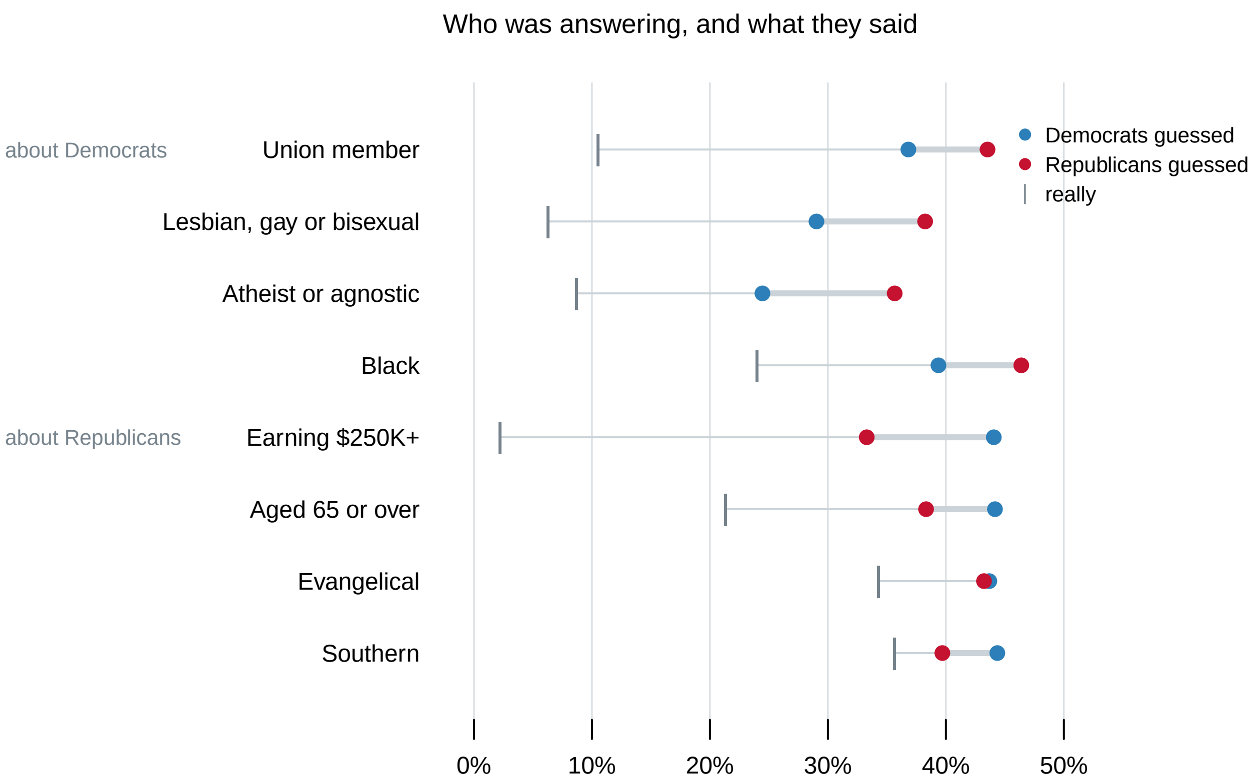

Figure 2. The same eight items, split by who answered. The blue dot is what Democratic respondents said, the red dot what Republican respondents said, and the grey tick is the truth. The top four rows ask about groups of Democrats and the bottom four about groups of Republicans.

The far dot changes colour halfway down, and that is the finding.

Who was answering	Asked about	Average overestimate (pts)
Democrats	their own party	+20.1
Republicans	their own party	+15.3
Democrats	the other party	+20.7
Republicans	the other party	+28.6

Asked about their own party, people overestimate by **17.7 points**. Asked about the other one, **24.6**. The worst single cell in the study is Democrats estimating how many Republicans are earning over $250,000 a year: out by **41.9 points**.

You cannot meet a political party. You learn what is in the other one from people describing it to you, and what you end up with is a picture assembled out of the descriptions.

Counting leaners as partisans is a decision here, and it is worth saying what it costs. People who call themselves independent and then admit to leaning are counted as partisans above, which is the party-id chapter's subject. Push them back out and the two numbers become 17.3 and 23.3. A smaller gap, the same direction, the same conclusion.

> **Before you read on, write two numbers down.** The average gap in Figure 1 is about twenty points. **First: of the 8,000 answers behind it, what share do you think were round numbers ending in a zero? Second: if you swapped the average guess for the middle guess, the value half the answers fall below, how much of that twenty-point gap do you think would survive?** Commit to both before turning the page.

04 What the average is an average of

Each point in Figure 1 is a thousand people reduced to one number. Here is what the thousand look like.

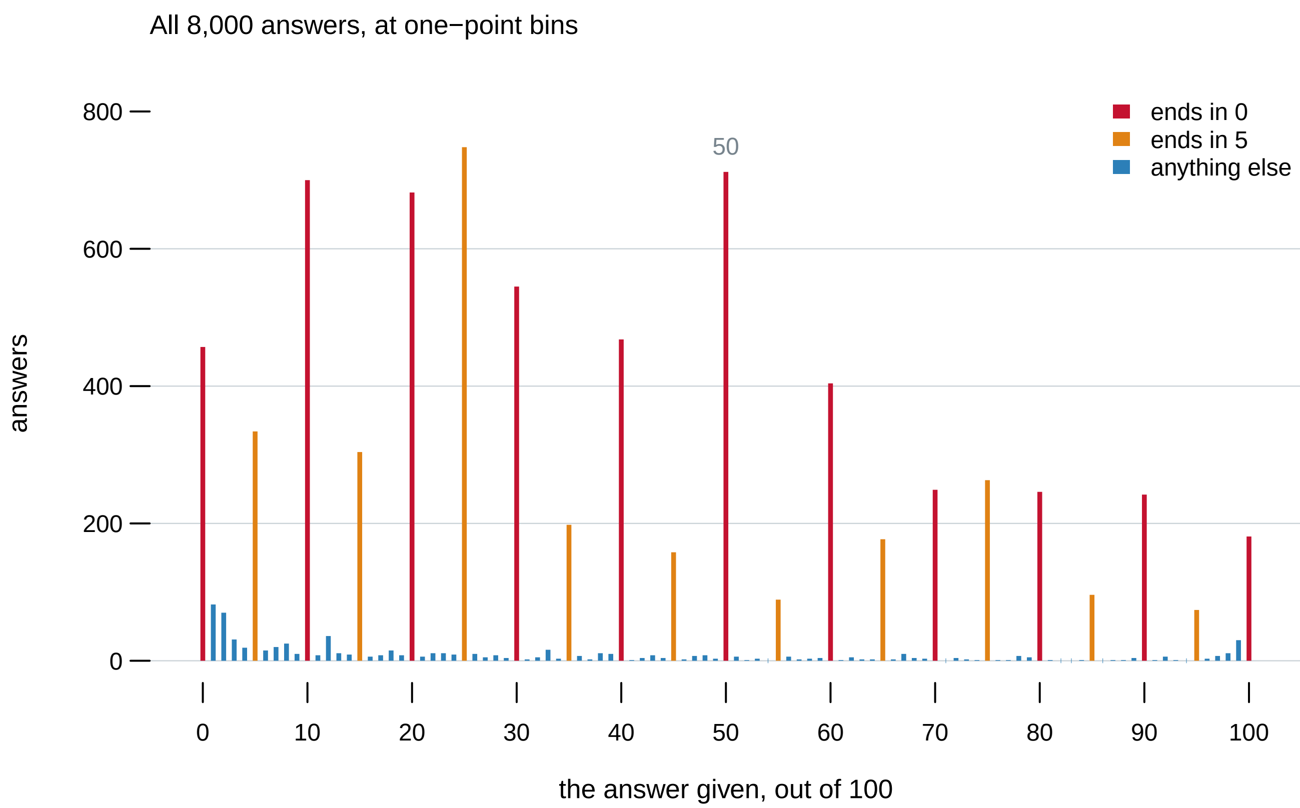

Figure 3. Every answer, at one-point bins. In the browser you can click any single item and see its own thousand answers against its own true share.

61.1% of the answers end in a zero. 91.6% end in a zero or a five. Only 8.4% use any of the other eighty values on the scale, and exactly 1 answer in the whole study is not a whole number: somebody typed 0.5.

Item	Ends in a zero	Ends in a zero or a five	Exactly 50
Black	60.1%	90.6%	8.6%
Union member	60.8%	92.9%	7.9%
Atheist or agnostic	62.2%	90.8%	8.8%
Lesbian, gay or bisexual	58.2%	88.6%	5.9%
Evangelical	63.7%	94.1%	8.9%
Earning $250K+	59.9%	90.0%	8.5%
Aged 65 or over	62.9%	92.9%	12.0%
Southern	60.8%	92.8%	10.6%
all eight pooled	61.1%	91.6%	8.9%

And 8.9% of all answers are exactly **50**. No group here has a true share anywhere near a half. Fifty is what a person types when they do not know and the box will not let them say so.

That pile is lumpy and it leans. A few people answering 90 pull the average up and nothing pulls it back, so the average sits well above the middle of the answers.

The same gap, measured off the mean and off the middle

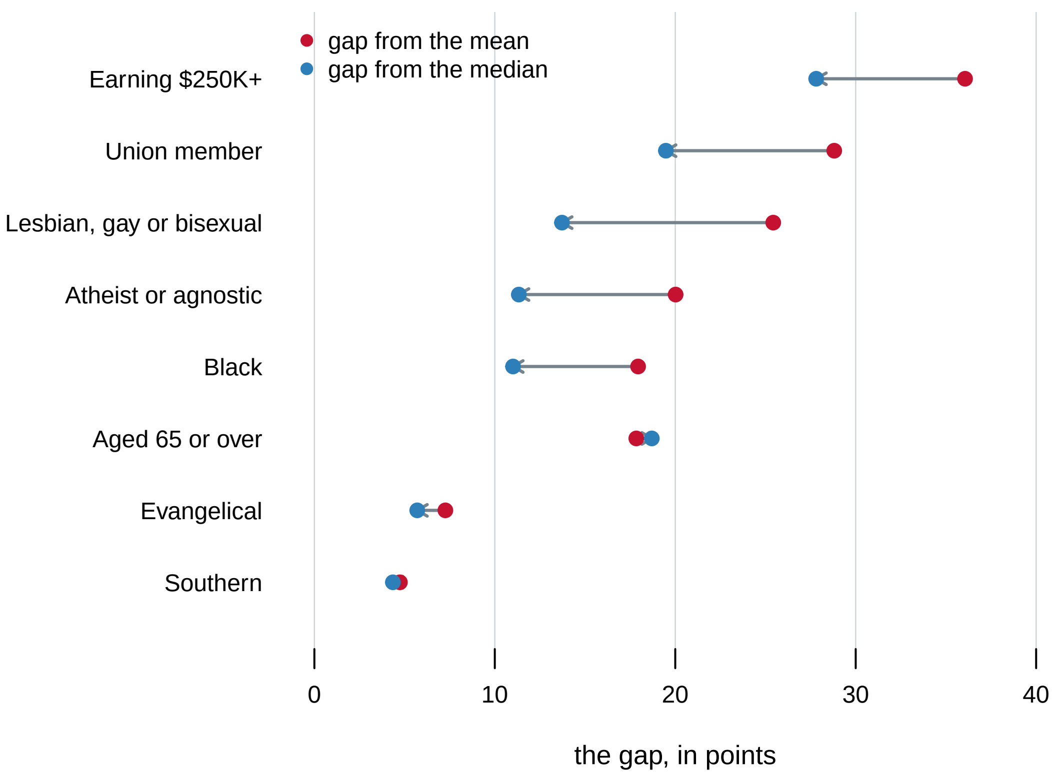

Group	Really	Mean guess	Middle guess	Gap from the mean	Gap from the middle
Earning $250K+	2.2	38.2	30	+36.1	+27.8
Union member	10.5	39.3	30	+28.8	+19.5

Group	Really	Mean guess	Middle guess	Gap from the mean	Gap from the middle
Lesbian, gay or bisexual	6.3	31.7	20	+25.4	+13.7
Atheist or agnostic	8.7	28.7	20	+20.0	+11.3
Black	24.0	41.9	35	+17.9	+11.0
Aged 65 or over	21.3	39.1	40	+17.8	+18.7
Evangelical	34.3	41.6	40	+7.3	+5.7
Southern	35.6	40.4	40	+4.7	+4.3

Figure 4. The same eight gaps, twice. Measured from the average guess, and from the middle one.

Measured off the middle, the average gap is 14.0 points rather than 19.8. Seven of the eight shrink and one grows. Neither number is the wrong one. The average answers *what does the country believe on average*, which is what the study asked. The middle answers *what does a typical person believe*. They are 5.8 points apart and the difference is never mentioned.

05 Everyone guesses forty

The true shares span **33.5 points** across these eight groups. The average guesses span **13.2**.

People are not answering at random. The guesses do line up with reality, at a correlation of 0.633, on a scale where 1 is a perfect straight line. But they are squeezed into a band, and the band sits high: six of the eight average guesses fall between 38 and 42.

That is the chapter's title and it is also a warning about the measure. Anyone answering near forty about a group that is really at 2.2% is 36.1 points wrong automatically. **A list of these gaps sorted from largest to smallest is very close to a list of the groups sorted from rarest to commonest.** The ranking is telling you about the groups as much as about the people guessing.

06 The survey, and the copy of it we have

Everything above comes from a file of 26 columns. The survey it was taken from had many more.

That is not a complaint. It is the ordinary condition of published research data: a researcher deposits what their results rest on. The questionnaire as administered, the panel the thousand people were drawn from, everyone who started and gave up — none of that is deposited, and none of it is anywhere a reader can go. **YouGov does not publish this study, so the deposit is the only route to any of it.**

The interesting part is that the deposit ships its own codebook, so the two can be set against each other.

Column	Which side it falls on
pew_churatd	described, and not in the file
pew_prayer	described, and not in the file
pew_religimp	described, and not in the file
ideo5	in the file, and not described
newsint	in the file, and not described
newsint_r	in the file, and not described
pid3	in the file, and not described
pid7	in the file, and not described
votereg	in the file, and not described

3 columns are described in the codebook and are not in the file. 6 are in the file and are not in the codebook, and one of those is `pid7`, **the column the whole of Figure 2 rests on**. The codebook says as much in its own first line: it is not comprehensive.

The rule of thumb. When a finding rests on a public file, the file is not the study. Ask what was collected, then ask what was deposited, and do not assume the second is the first with fewer rows.

07 If you ran this in a room

Here is the block, ready to append to the class survey. Put it after the questions that ask who *you* are, because these ask who you think everybody else is, and that order keeps the second set from colouring the first.

Field	Question
percep_dem_atheist	Out of every 100 Democrats, how many do you think are atheist or agnostic?
percep_dem_black	Out of every 100 Democrats, how many do you think are Black?
percep_dem_lgb	Out of every 100 Democrats, how many do you think are lesbian, gay or bisexual?
percep_dem_union	Out of every 100 Democrats, how many do you think belong to a labour union?
percep_rep_250k	Out of every 100 Republicans, how many do you think earn more than $250,000 a year?
percep_rep_65plus	Out of every 100 Republicans, how many do you think are 65 or older?
percep_rep_evangelical	Out of every 100 Republicans, how many do you think are evangelical Christians?
percep_rep_south	Out of every 100 Republicans, how many do you think live in the South?

> The last eight questions ask you to estimate. Nobody expects you to know these — guess. Give a whole number between 0 and 100. Please do not look anything up: a wrong guess is useful data and a correct look-up is not.

Set each item as a short answer, validated as a number between 0 and 100, shuffled within the block, all eight required. Two things there are decisions rather than formatting, and this chapter is the evidence for both.

The wording asks for a count out of a hundred, not a percentage. The original asked for a percentage and 61.1% of the answers came back ending in a zero. Asking somebody to count out of a hundred will not remove that, but it invites a finer answer from anyone who does not think in percentages.

Score the room against two things, not one. Against the true shares is whether the room is wrong. Against the national guesses is the better question: is a room of students less wrong than the country, or wrong in a different direction?

And the actuals are a decade old. They describe the parties as of 2015. Recomputing them from a current survey is a real exercise, and the distance between the 2015 figure and today's is a finding rather than a correction.

08 What this chapter cannot tell you

No room has answered these questions yet. Everything here is the national study. The class comparison is set up and not run, and nothing here is a result about any group of students.

Why people answer near forty is not settled. It could be that they picture the party through its stereotypes, which is the study's argument. It could be that somebody with no information gives a number in the middle of the scale, which would produce the same pattern without anyone being misled about anything. This chapter can show the compression. It cannot say what causes it.

The true shares are estimates too. Each was computed by the authors from the American National Election Studies, the Current Population Survey or a Pew report, and each carries a standard error. The widest is Evangelical: 34.3%, honestly somewhere between 30.9% and 37.7%. For most items that is small against a twenty-point gap. It still means the gap is a distance between two estimates, and "the true share" is doing work the word "estimate" would not do.

The panel is not the country. YouGov recruits people who volunteer to take surveys online, and the case weight is what stands between that and a national statement. The surveys-source chapter sets out what a sample like this can establish, and the ces-class chapter shows what happens when the weights are turned off.

Eight groups is a small and chosen list. The authors picked groups each party is stereotyped around, which is the right list for their question and a list built to produce large gaps. A different eight, say Democrats who are veterans or Republicans who rent, would give different numbers, and this chapter has no way to say how different.

09 Sources

- The data is a **YouGov survey of 1,000 American adults, fielded in March 2015** for Douglas J. Ahler and Gaurav Sood. YouGov does not publish it.
- The copy read here is their replication deposit for "The Parties in Our Heads: Misperceptions About Party Composition and Their Consequences," *The Journal of Politics* 80(3), 2018: 964–981, on Harvard Dataverse, doi:10.7910/DVN/CLMQ8E. 26 columns of the study, plus the eight true shares and the published averages. https://doi.org/10.7910/DVN/CLMQ8E

- The true shares in that deposit were computed by the authors from the American National Election Studies, the Current Population Survey and a Pew Research Center report on religion.
- The averages recomputed here from the 8,000 individual answers reproduce the published ones to five decimal places, which is the check this chapter runs before it writes anything else.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`by_party.csv`, `coverage.csv`, `dist.csv`, `heaping.csv`, `items.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `items.csv` puts `actual_pct` beside `perceived_pct` for every group asked about. Which group is misjudged by the widest margin, and is it guessed too high or too low?
- `by_party.csv` splits every guess by who is doing the guessing. Do people judge their own party more accurately than the other one? Check every item before you answer.
- `heaping.csv` counts the answers that are multiples of ten, multiples of five, and exactly 50. Work out what share are multiples of five. What does that tell you about how precise a survey percentage really is?
- `items.csv` gives both a mean and a median for the same question, and they are far apart. Look at `dist.csv` to see why. Which one would you report?

REBUILD THIS DATA YOURSELF, WITH AN AI ASSISTANT

I want to rebuild a public dataset from its original source, without any starting files. Please do the following and show your work.

THE SOURCE. A YouGov survey of 1,000 American adults, fielded in March 2015 for Douglas Ahler and Gaurav Sood. YouGov does not publish it, and you cannot have it. What you can have is the copy of part of it the two authors deposited so their results could be checked: the replication archive for "The Parties in Our Heads: Misperceptions About Party Composition and Their Consequences", The Journal of Politics 80(3), 2018, on Harvard Dataverse at doi:10.7910/DVN/CLMQ8E. Seventeen files, all public, no account and no key. You want four of them:

pcomp_yougov_data 1,000 respondents from a March 2015 YouGov study, one row each, with eight estimates, a case weight and party identification

fig_1_data_actual the true share of each group in each party, with a standard error on each

fig_1_data_perc the published perceived means, to check yourself
against
codebook what every column holds

HOW TO FETCH THEM. List the files with the Dataverse API:
https://dataverse.harvard.edu/api/datasets/:persistentId/?persistentId=doi:10.7910/DVN/CLMQ8E
That returns JSON with a numeric id for each file. Download one with
https://dataverse.harvard.edu/api/access/datafile/ followed by the id.
The catch is silent and it will cost you an afternoon if you miss it.
That address answers 303, not 200, and hands back a redirect to
storage. A client that does not follow redirects writes a zero-byte
file and reports success. Follow redirects, then check the file size
before you parse anything.

WHAT TO BUILD.

1. One row per person per item: 8,000 answers, each with the case weight and the group and party it belongs to.
2. The weighted mean of each item, which is what the paper published. Weight, or you will be close and wrong.
3. The median of each item beside that mean, and the gap from the true share computed both ways.
4. How round the answers are: the share ending in a zero, the share ending in a zero or a five, and the share that are exactly 50.
5. The eight items split by the respondent’s own party, from pid7, so you can compare what people say about their own party against what they say about the other one. Note that pid7 is in the file and not in the codebook; set the two against each other and you will find three columns described and absent, and six present and undescribed.

CHECK YOUR WORK. The copies this book used were fetched in August 2026.
If your rebuild matches them, you should find:

- 1,000 respondents and 8,000 answers, every one between 0 and 100
- your weighted means landing on the published ones to five decimal places; the worst of the eight items here first differs from the published figure in the sixth decimal place
- an average gap of 19.8 points measured from the mean and 14.0 measured from the median

> – 61.1% of answers ending in a zero, 91.6% ending in a zero or a five, and 8.9% of them exactly 50
> – the largest single gap on the Republican $250,000-and-over item: a true share of 2.19% against a mean guess of 38.25%
> – people overestimating their own party by 17.7 points and the other party by 24.6, counting leaners as partisans
> This archive was deposited in 2017 and has not changed since, so a number that does not match is yours rather than theirs.