

Bellwether Counties, 1960-2024

Bellwethers, 1960-2024, and why a perfect record predicts nothing

84-355 Democracy's Data

August 2026

On 13 November 2020, nine days after the election was called, the *Wall Street Journal* published a list of nineteen counties. Each had voted for the winner of every presidential election since 1980 — ten in a row, without a miss. The article's subject was that almost all of them had just failed at once.

The list was not unusual. Every four years the same genre appears. *The ten counties that show Trump is in serious trouble* in the *New York Times*, a month before the 2020 election. A *Washington Post* feature on Vigo County, Indiana. And a small library of local news stories in which a county explains what it means to be a barometer of the nation. The counties are real, the records are real, and the arithmetic is checkable.

This chapter checks it. The question is not whether these nineteen counties had perfect records — they did, and this file reproduces the *Journal's* list exactly, from a source the *Journal* did not use. The question is what a perfect record is worth.

01 What a bellwether would have to be

A bellwether is a county whose vote matches the national outcome, election after election. The interesting claim is not historical but predictive: that knowing a county has matched the last ten winners tells you something about the eleventh.

For that to be true, the record has to be evidence of a property — something about the county that produces the match. The alternative is that the record is evidence of nothing except that a great many counties exist.

There are 3,081 counties in this file. Suppose each one were a fair coin, landing on the national winner half the time and independently of the others. Then the number with a perfect record over k elections would be $3,081 \times 0.5^k$, which is not a small number until k is large.

Elections	Perfect records expected from coin-flipping	Perfect records observed
the last 3 (2016-2024)	385	32
the last 5 (2008-2024)	96	14
the last 7 (2000-2024)	24	1
the last 10 (1988-2024)	3.0	0
the last 13 (1976-2024)	0.38	0


Elections	Perfect records expected from coin-flipping	Perfect records observed
the last 17 (1960-2024)	0.02	0

Two things in that table matter, and they point in opposite directions.

Over a short run, chance alone manufactures bellwethers by the hundred. Thirty-two counties called the last three elections right. Nobody should be impressed: coin-flipping would have produced around 385. A county with a three-election record is not a barometer, it is a rounding error.

Over a long run, there are fewer than chance predicts, not more. At every length in the table the observed count sits *below* the coin-flip expectation. Seven elections: 24 expected, 1 observed. Over all 17, **0 counties in the country have a perfect record**. Not a handful. None.

That second fact is the one that does the work, and it is not what the folk theory predicts. If some counties really were demographically representative barometers, they would beat the coin, and the top of the distribution would be *fatter* than chance, not thinner. It is thinner.



All 3,081 counties by how many of the 17 elections they called correctly, against what independent coin-flipping would produce.

Figure 1. Every county, by how many elections it called correctly. The bars are the 3,081 counties in the file; the dashed line is what 3,081 independent fair coins would produce. The observed distribution is *narrower* than the coin model at both ends: fewer counties with very good records, fewer with very bad ones. At 17 of 17 the coin model predicts 0.02 counties and the country contains 0.

02 Where these numbers come from, and the wall in front of them

Before going further it is worth saying what had to happen to ask this question at all, because the answer is the least reassuring thing in the chapter.

No federal agency counts votes. County-level presidential returns exist as fifty-one separate state publications, and for elections in the 1960s and 1970s most of them exist only on paper. There is no official national file.

Every long county series in circulation is therefore a *compilation*. The two that political scientists use, CQ Press's *Voting and Elections Collection* and Dave Leip's *Atlas*, are commercial products behind a subscription. Grofman and Chen used CQ Press. The *Journal* and the *Times* were working from Leip.

This file uses the one long series that is free: the replication archive Carlos Algara and Sharif Amlani published with their 2021 *Electoral Studies* paper, which carries county presidential returns from 1868 to 2020. It is openly licensed and posted for reuse on Harvard Dataverse. It is also, itself, built out of CQ Press and the ICPSR historical returns — **a compilation of compilations, none of which any government produced.**

And it cannot be fetched by a script.

```
$ curl -sS -D- \
    "https://dataverse.harvard.edu/api/access/datafile/5028532"
```

```
HTTP/2 202
server: awselb/2.0
content-length: 0
x-amzn-waf-action: challenge
```

Harvard Dataverse sits behind a web application firewall that answers automated requests with a JavaScript challenge. Note what comes back: **status 202, Accepted**. Not 403, not 401, not an error of any kind. A script that checked whether the request succeeded would conclude that it had, and would write an empty file. The other builds in this book all fetch their own data; this one is the single exception: the work here refuses to run rather than pretend otherwise. Somebody has to open a browser once.

Nothing here is a permission problem. The file is free, the licence is open, the authors want it used. The obstruction is plumbing, it is invisible, and it is the reason to be suspicious of the sentence “the data show” whenever nobody says which data.

Two checks, because a file that arrives by hand is a file nobody audited

A hand-carried download is exactly the circumstance in which the wrong file would go unnoticed, so the build script checks it against two sources it *can* fetch.

Check	Result
Against MEDSL's independently compiled county returns, 2000–2016	15,562 county-elections compared, 12 disagreements about the winner (0.077%)
Against 26 states' own certified 2024 returns, from county-returns/	1,909 counties compared, every vote total identical
Reproducing the Wall Street Journal's list from scratch	searched for counties perfect over 1980–2016: found 19, and they are the same 19

The middle check is the book auditing itself: `county-returns/` is the standing project to replace the GitHub compilation this chapter uses for 2024 with the states' own certified publications, and twenty-six jurisdictions are finished. On all 1,909 of their counties the compilation and the states agree to the vote.

The 12 disagreements in the first check are worth a moment, because none of them is arithmetic.

County	State	Year	Algara & Amlani say	MEDSL says
Jackson	MO	2004	D	R
Jackson	MO	2008	D	R
Jackson	MO	2012	D	R
Jackson	MO	2016	D	R
Cayuga	NY	2004	R	D
Chemung	NY	2012	R	D
Dutchess	NY	2016	R	D
Montgomery	NY	2012	R	D
Ontario	NY	2012	R	D
Rockland	NY	2004	R	D
Sullivan	NY	2004	R	D
Lake	TN	2004	D	R

7 are in New York, where **fusion voting** lets one nominee appear on several party lines, so the two compilers add up different subsets of the ballot. Four are Jackson County, Missouri, where **Kansas City reports separately** and one compiler keeps the city out of the county. One, in Tennessee, is a straight transposition — the Democratic and Republican totals are swapped in one of the files. These are unit decisions and a typo, not measurement error, and the same fusion problem bit the `house-competition` build in a different costume.

None of them touches a bellwether: the best record among these counties is 14 of 17.

03 One row is one county

Field	Value
FIPS code	18167
county	Vigo
state	IN
carried by, 1960 through 1976	D D R R D
1980 through 1996	R R R D D
2000 through 2016	R R D D R
2020	R
2024	R

That is Vigo County, Indiana — the most written-about bellwether in the country, home of Terre Haute, and for decades the standard illustration of the genre. The file holds one such row for each of 3,081 counties, and it holds nothing else: no demographics, no turnout, no margins. Just who carried the county.

Three things about the panel are decisions rather than facts.

Alaska is not in it. Alaska has no counties. It reports presidential votes by State House district, so there is no county-level return to carry, and every count in this chapter is of the other fifty jurisdictions. The `county-returns` project makes the same call for the same reason.

Counties that did not exist throughout are dropped. 99 of the 3,180 FIPS codes appearing anywhere in the series are missing from at least one election — created, merged, renamed, or in Virginia's case repeatedly reorganised between cities and counties. A county cannot have a seventeen-election record if it did not exist for seventeen elections.

Six counties are dropped for a tie. In a tied county there is no winner to match. They are Ware County, Georgia; Boone and Gallatin Counties, Kentucky; Clarke County, Mississippi; Horry County, South Carolina; and Ferry County, Washington.

04 The nineteen


Now the test. The *Journal's* nineteen counties were selected in November 2020 on a record running through 2016. Whether that record was worth anything is a question with a real answer, because two elections have happened since.

Before you read on, write a number down. Nineteen counties, each with a perfect ten-election record when the list was published. Two elections have been held since. **How many of the nineteen still have a perfect record?** Commit to a count before you look at the table — and, if you want the harder version of the question, commit also to whether the ones that failed did so in the *same* election as each other.

County	Record 1980-2016	2020	2024	Still perfect
Bremer, IA	10 of 10	R	R	no
Warren, IL	10 of 10	R	R	no
Vigo, IN	10 of 10	R	R	no
Washington, ME	10 of 10	R	R	no
Shiawassee, MI	10 of 10	R	R	no
Van Buren, MI	10 of 10	R	R	no
Hidalgo, NM	10 of 10	R	R	no
Valencia, NM	10 of 10	R	R	no
Cortland, NY	10 of 10	R	R	no
Otsego, NY	10 of 10	R	R	no
Ottawa, OH	10 of 10	R	R	no
Wood, OH	10 of 10	R	R	no
Westmoreland, VA	10 of 10	R	R	no
Essex, VT	10 of 10	R	R	no
Clallam, WA	10 of 10	D	D	no
Juneau, WI	10 of 10	R	R	no
Marquette, WI	10 of 10	R	R	no
Richland, WI	10 of 10	R	R	no
Sawyer, WI	10 of 10	R	R	no

The national winner in 2020 was a Democrat. **18 of the 19 voted Republican**, and their records ended that night. One survived: Clallam, WA, which voted for Biden as the country did.

The national winner in 2024 was a Republican. Clallam, WA voted Democratic. **0 of the nineteen survive.**



The nineteen counties the Wall Street Journal published in November 2020, one square per election.

Figure 2. The nineteen counties, 1960–2024. Each row is a county and each square an election, green where the county matched the national winner. The block from 1980 to 2016 is solid green by construction – that is the stretch the counties were selected on, and a record cannot fail a test it was chosen to pass. The two columns to the right of it are the only ones carrying information, and they are almost entirely purple.

05 They were never nineteen predictors

The temptation is to read 18 simultaneous failures as 18 separate surprises — nineteen barometers that all happened to break. That reading is arithmetically impossible.

If the nineteen were independent, this could not happen. Nineteen coins, each landing on the national winner half the time, produce eighteen or more failures in a single election with probability 0.00004 — about four in a hundred thousand. The event is not rare. It is nearly certain, once you notice what the counties had in common.

Quantity	Value
Of the nineteen, the share that voted Republican in 2020	95%
Of all 3,081 counties, the share that voted Republican in 2020	83%
Probability 18 of 19 independent fair coins fail together	0.00004

95% of the nineteen voted Republican in 2020. So did 83% of every county in America.

The bellwethers did not do anything unusual. They did what the typical American county did, which is what they had been doing all along. And in 2020 the typical American county voted for the loser of the national popular vote, because most counties are small and rural, and the national majority no longer lives in them.

That is the whole mechanism. A bellwether is not a county that tracks the nation. It is a county whose *partisan lean happens to sit near the tipping point* of the elections you looked at. When the tipping point moves, they all fail together, because they were all the same bet.

06 Why the bet stopped paying

The reason bellwethers have become scarce is not that counties got worse at predicting. It is that counties stopped changing their minds.



The share of counties that changed party from the previous election.

Figure 3. How much the map moves. The share of counties voting differently than they had four years earlier. In 1968, 56% of counties changed party. In 2020 it was 2.5% and in 2024, 2.8%.


A county that never changes its mind cannot be a bellwether, because sooner or later the nation changes its mind without it. When 56% of the map is in motion, an enormous number of counties will happen to move with the country; when 2.8% is in motion, almost none will.

Grofman and Chen put a sharper edge on this. A county has to be right when Democrats win *and* when Republicans win, and those two abilities have come apart.

Six elections	Won by	P(correct D wins)	P(correct R wins)	Correlation	Perfect records
1960-1980	3D / 3R	0.560	0.773	-0.638	78 (2.53%)
1968-1988	1D / 5R	0.550	0.790	-0.597	375 (12.17%)
1980-2000	2D / 4R	0.489	0.782	-0.644	241 (7.82%)
1992-2012	4D / 2R	0.368	0.801	-0.738	60 (1.95%)
2000-2020	3D / 3R	0.221	0.817	-0.848	9 (0.29%)
2004-2024	3D / 3R	0.221	0.841	-0.890	2 (0.06%)

Read the two middle columns. In 1960-1980 a randomly chosen county called Democratic victories right 56% of the time and Republican victories 77% of the time. By 2004-2024 the first number has fallen to 22% while the second has *risen* to 84%.

The average county has not got worse at predicting. It has got worse at predicting one party and better at predicting the other, which is a different thing and is fatal to bellwethers. A perfect record requires being right in both kinds of year, and the correlation between those two abilities has gone from -0.64 to -0.89. Counties that are good at one are now almost mechanically bad at the other.



Two conditional probabilities that used to be close together and no longer are, measured

Figure 4. The scissors. For each six-election window, how often a randomly chosen county called the winner when a Republican won and when a Democrat won. The two lines start 21 points apart and end 62 points apart. A bellwether has to be high on both lines at once, and after about 1992 there is no such place to stand.

The two effects are not the same size, and the gap between them is the point. Average accuracy fell from 0.667 in 1960-1980 to 0.531 in 2004-2024, down about 20%. Over the same stretch perfect records fell from 78 counties to 2, down 97%.

A fifth off the average wiped out almost all of the perfect records, and that disproportion is what the correlation buys. Being right on average is a statement about one election at a time; a perfect record is a statement about every election at once, and it requires the two abilities in Figure 4 together. When they were mildly opposed a county could hold both. At -0.89 there is almost nowhere left that is good at both, so the top of the distribution empties out far faster than the middle of it moves.

07 The definition nobody states


One more thing decides the answer, and it is never specified in the newspaper version: what “the national winner” means.

In 2000 and 2016 the popular vote and the Electoral College named different men. Grofman and Chen argue that the popular vote is the sensible standard for this question – a county is a sample of the country’s voters, not of its electors. The newspaper lists all use the Electoral College. Both are defensible and they do not agree.

“the national winner” means	Elections won	Perfect records	Best record	Held by
Electoral College winner	8D / 9R	0	16	Vigo, IN; Blaine, MT; Valencia, NM
National popular-vote winner	10D / 7R	0	16	Socorro, NM

Under either definition the answer to “how many perfect bellwethers are there” is **0**. But the near-misses are different counties entirely. Three counties come within one election under the Electoral College: Vigo, IN, Blaine, MT, Valencia, NM. They share **no members** with the one that comes within one under the popular vote, Socorro, NM.

Three of those Electoral College near-misses are counties Grofman and Chen named in 2023 as the last survivors, when the series ended at 2020. They are still the last survivors. What changed is that the fourth, Kent, DE, missed 2024.



The same counties scored against two different definitions of “the national winner”.

Figure 5. Two definitions, two different sets of counties. The same 3,081 counties and the same seventeen elections, scored against the Electoral College winner and against the popular-vote winner. The distributions are not small variations on each other, because switching the definition flips the scoring of two whole elections at once. Neither definition produces a single perfect record, which is the one thing they agree on.

08 What this chapter cannot tell you

The panel holds who carried each county and nothing else, and the limits follow from that directly.

It cannot explain why any county votes as it does. There are no demographics here, no incomes, no turnout, no margins. Zimny-Schmitt and Harris (2020) found the bellwethers to be whiter, older, less educated and lower-income than the country. That is already enough to sink the “demographically representative” story, and nothing in this file could have told you it.

A margin is not a winner. Every county here is reduced to D or R, so a county carried by nine votes counts exactly as much as one carried by ninety thousand. That is the right unit for the bellwether question, which is asked in those terms, and it is the wrong unit for almost any other question. Six counties were dropped for a literal tie, and a county decided by eleven votes sits in the file as a clean call.

The long series is a compilation of compilations. Its own sources are CQ Press and ICPSR. The 0.077% disagreement rate against MEDSL is a floor on the uncertainty rather than a measurement of it, because two compilations sharing a source will agree more than either agrees with the truth. Nobody audited the 1964 returns of every county in America, and this chapter cannot.

Counties are not stable objects. 99 FIPS codes were dropped for not existing throughout, and treating a county as the same thing in 1960 and 2024 is a convention. Boundaries moved, cities were carved out and folded back, and the seventeen-election panel quietly assumes none of that mattered.

Alaska is absent because it has no counties, so every proportion here is of the other fifty jurisdictions.

09 What a perfect record is worth

3,081 counties, 17 elections, and **0 perfect records** — under either definition of the national winner. Coin-flipping alone would have produced 0.02, so the finding is not that bellwethers are rarer than chance in some technical sense. It is that there are none, and there is no reason to have expected any.

The nineteen counties the *Wall Street Journal* published in November 2020 had genuinely perfect records over the stretch they were chosen on. This file reproduces that list exactly, without using the *Journal's* data. Two elections came after, the only two the list did not already contain. 18 of the counties failed at once, and the last one failed four years later. 0 survive.

They failed together because they were never nineteen independent readings on the country. They were nineteen copies of one bet: that a county sitting near the national

tipping point in the 1980s would still be sitting there. When the tipping point moved, it moved out from under all of them at once. In 2020 95% of them voted Republican, which is almost exactly the 83% of *all* American counties that did. Their famous predictive power amounted to being ordinary.

The habit worth taking from this is not scepticism about counties. It is a question to ask of any record presented as evidence: **out of how many tries?** A perfect ten-election record is remarkable in one county chosen in advance and unremarkable in a country that contains 3,081 of them. The *Journal* did not choose in advance. Nobody does. The list was assembled by searching 3,081 counties for the ones that had already won, and a search like that is guaranteed to find some. That is why the answer it returns carries no information about the future at all.

Tufte demonstrated this in 1974, and again with Sun in 1975, on the elections through 1968. Hopkins reviewed the evidence again in 2017. Grofman and Chen called it “flogging a dead horse” in 2023 and wrote the paper anyway, because the horse keeps getting up. It will be back in 2028.

10 Sources

Data

- Algara, Carlos and Sharif Amlani. 2021. *Replication Data for: Partisanship & Nationalization in American Elections: Evidence from Presidential, Senatorial, & Gubernatorial Elections in the U.S. Counties, 1872–2020*. Harvard Dataverse, V1. <https://doi.org/10.7910/DVN/DGUMFI>. File `dataverse_shareable_presidential_county_returns_1868_2020.Rdata`; used for 1960–2020. **Must be downloaded through a browser**; there is no address a script can call.
- MIT Election Data and Science Lab, *County Presidential Election Returns 2000–2016*, <https://github.com/MEDSL/county-returns>. Used only to check the file above.
- McGovern, Tony. *US County Level Election Results 08–24*, https://github.com/tonmcg/US_County_Level_Election_Results_08-24. Used for 2024, and checked against the twenty-six states’ certified returns assembled in `labs/county-returns/`.
- The party of the national winner in each election is not in any of these files. It is typed in by hand as two columns, one for the Electoral College and one for the popular vote.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`county_winners.csv`, `crosscheck.csv`, `national.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `county_winners.csv` has one column per election and a D or an R in each. Count the counties that match the national winner every single time. Then work out how many you would expect by chance alone.
- Take the counties with a perfect record and look at what they did in the most recent year. A perfect record up to now — what did it actually predict?

- `crosscheck.csv` documents places where two compilations disagree about who won a county, with the reason spelled out. Read one. How would you have caught that without somebody writing it down?
- `national.csv` records both the electoral college and popular vote winner each year. Find the years they differ. What does that do to the whole idea of a county that 'picks the winner'?

Academic literature

- Grofman, Bernard and Haotian Chen. 2023. "Understanding the Factors that Affect the Incidence of Bellwether Counties: A Conditional Probability Model." *Political Research Quarterly* 76(1): 119-126.
- Amlani, Sharif and Carlos Algara. 2021. "Partisanship & Nationalization in American Elections: Evidence from Presidential, Senatorial & Gubernatorial Elections in the U.S. Counties, 1872-2020." *Electoral Studies* 73: 102387.
- Deutsch, Karl W. and William G. Madow. 1961. "A Note on the Appearance of Wisdom in Large Bureaucratic Organizations." *Behavioral Science* 6: 72-78.
- Hopkins, David A. 2017. *Red Fighting Blue: How Geography and Electoral Rules Polarize American Politics*. Cambridge University Press.
- Tufte, Edward R. 1974. *Data Analysis for Politics and Policy*. Prentice-Hall, chapter 3.
- Tufte, Edward R. and Richard A. Sun. 1975. "Are There Bellwether Electoral Districts?" *Public Opinion Quarterly* 39(1): 1-18.
- Zimny-Schmitt, Daniel and Michael Harris. 2020. "An Inquiry of Bellwether Counties in US Presidential Elections." <https://doi.org/10.2139/ssrn.3677602>.

Journalism this chapter tests

- McCormick, John. "Bellwether Counties Nearly Wiped Out by 2020 Election." *Wall Street Journal*, 13 November 2020. The source of the nineteen.
- Wasserman, David. "The 10 Bellwether Counties That Show Trump Is in Serious Trouble." *New York Times*, 6 October 2020.

WHY THIS DATA CANNOT BE FETCHED AGAIN

The core file behind this chapter comes from a source no assistant can fetch for you – though you can fetch it yourself, once, in a browser.

WHY NOT. The long series is Algara and Amlani's county-by-county presidential returns, 1868-2020, published openly on Harvard Dataverse at <https://doi.org/10.7910/DVN/DGUMFI>. Dataverse sits behind a bot wall. A scripted request is not refused: it comes back with a success-shaped status code and zero bytes attached, plus a challenge that only a real browser's JavaScript can answer. A tool that did not check would record the empty download as a win. The file is free and openly licensed; the obstacle is plumbing, not permission.

WHAT EXISTS INSTEAD. Open the DOI above in a browser and download one file by hand:

`dataverse_shareable_presidential_county_returns_1868_2020.Rdata` (about 1.5 MB, in R's data format). Everything else arrives by script: the 2024 county results from the `tonmcg` compilation on GitHub, and MEDSL's independently compiled 2000–2016 county returns, which exist here to cross-check the hand-carried file. There is no free official substitute, because no federal agency counts votes county by county.

WHAT YOU CAN STILL CHECK. If you obtain the file, these numbers say whether you are holding the same object:

- the download is exactly 1,459,576 bytes
- kept to counties that are present, with a clear winner, in all seventeen elections from 1960 through 2024, the panel holds 3,081 counties
- the two independent sources that overlap for 2000–2016 disagree about a county's winner in just twelve county-years
- the 2024 file from the compilation adds 3,160 county rows