What Candidates Would Not Say

Fourteen issues, 1,662 candidates, and a code that means nobody mentioned it

84-355 Democracy's Data

August 2026

In 2014, 1,662 people ran in a congressional primary. A team at the Brookings Institution read every one of their campaign websites and wrote down where each candidate stood on fourteen issues of the day. Immigration. The Affordable Care Act. Guns, abortion, the minimum wage, climate change.

222 of those candidates said nothing about any of the fourteen.

That is not a gap in the data. It is the data. The people who built this file gave themselves four codes for each issue, and the fourth one is **"Candidate Provides No Information."** Read across all fourteen issues, the fourth code is the most common answer in the file. The median candidate stated a position on 6 of the 14 and was silent on the other 8.

This chapter is about what that means. A campaign website is a document a candidate chose to write. A blank on it is a decision, not an absence.

01 Where these numbers come from

The **Primaries Project** was a Brookings research effort led by Elaine Kamarck and Alexander Podkul, reported in 2014. Their question was about factions: they wanted to know whether the Tea Party was really taking over the Republican Party, and what the equivalent fight looked like among Democrats. To answer it they needed to place every candidate, not just the famous ones.

So they coded nearly everybody who filed. The result covers 1,662 candidates in the 2014 congressional primaries, and it is the whole field rather than a sample of it.

Quantity	Value
Candidates coded	1,662
House candidates	1,443
Senate candidates	219
Issues coded per candidate	14
Columns in the workbook	48
Primary election dates covered	15, from 4 March to 9 September

The coding was done from the candidates' own websites. That single fact decides what every number below can mean. Nobody interviewed these candidates. Nobody sent them a questionnaire. A researcher opened the campaign site, read what was there, and coded it.

So the file does not record what candidates believed. It records what they published. Those are different things, and the distance between them is the subject of this chapter.

The fourth code is where the difference lives.

Code	What the codebook calls it
1	Candidate Supports
2	Candidate Opposes
3	Candidate Provides Complicated/Complex/Unclear Position
4	Candidate Provides No Information

Codes 1 and 2 are positions. Code 3 is a position the coder could not resolve into a side. **Code 4 is not a position at all.** It says the researcher read the website and found the subject was not raised.

That is a genuinely useful thing to record, and most datasets throw it away. A survey would have offered "no opinion" and a candidate would have picked something. Here nobody was asked, so nobody had to answer.

What it costs to get. Nothing, in every ordinary sense: the workbook and its codebook are free to anyone, with no account, no key, and no fee. What is missing is any promise of permanence. Both files sit in the upload folder of a content management system rather than in an archive, the report they accompanied has already lost its own link, and neither carries a DOI. An address like that owes nobody an answer tomorrow, which is why copies of both are kept beside this chapter.

02 The most common answer is nothing

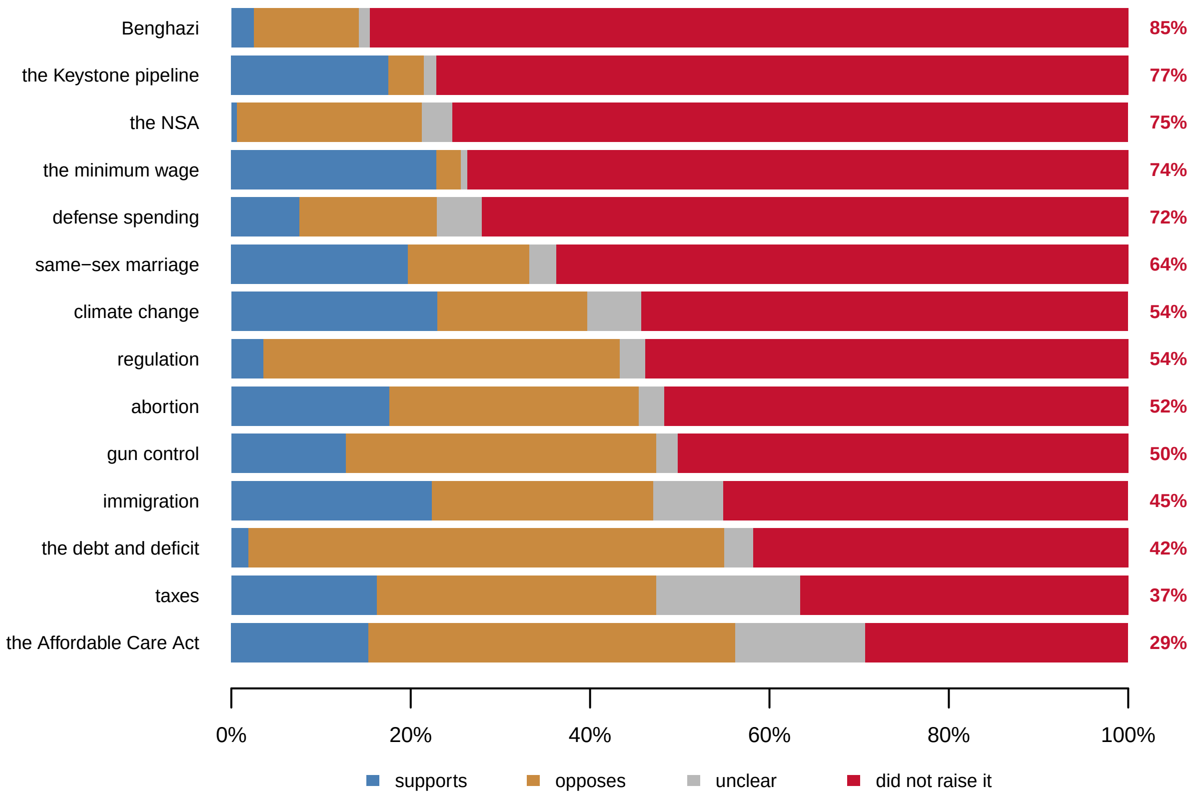

Figure 1. Where 1,662 candidates stood on fourteen issues, sorted by how many stayed off the subject. Every row is the same 1,662 candidates. The figure at the right of each row is the share coded "provides no information."

The range is wide and it is not random. Benghazi is the quietest subject at 84.5 percent, and the Affordable Care Act the loudest at 29.3 percent.

Read the top row and the bottom row together. Benghazi was the story conservative media ran hardest in 2014, and 1,405 of 1,662 candidates did not put it on their website. The Affordable Care Act was the most-discussed subject in the file, and 487 candidates still said nothing about it.

Silence is not evenly distributed across the parties either, but this table cannot tell you that cleanly. A candidate's silence on abortion means one thing in a contested primary and another in a race nobody was watching.

03 How much does one candidate leave out

Before you read on, write a number down. The file codes 14 issues for each of 1,662 candidates. **On how many of those 14 do you think the typical candidate said nothing?** Pick a number between 0 and 14. Commit to it before you look at Figure 2.

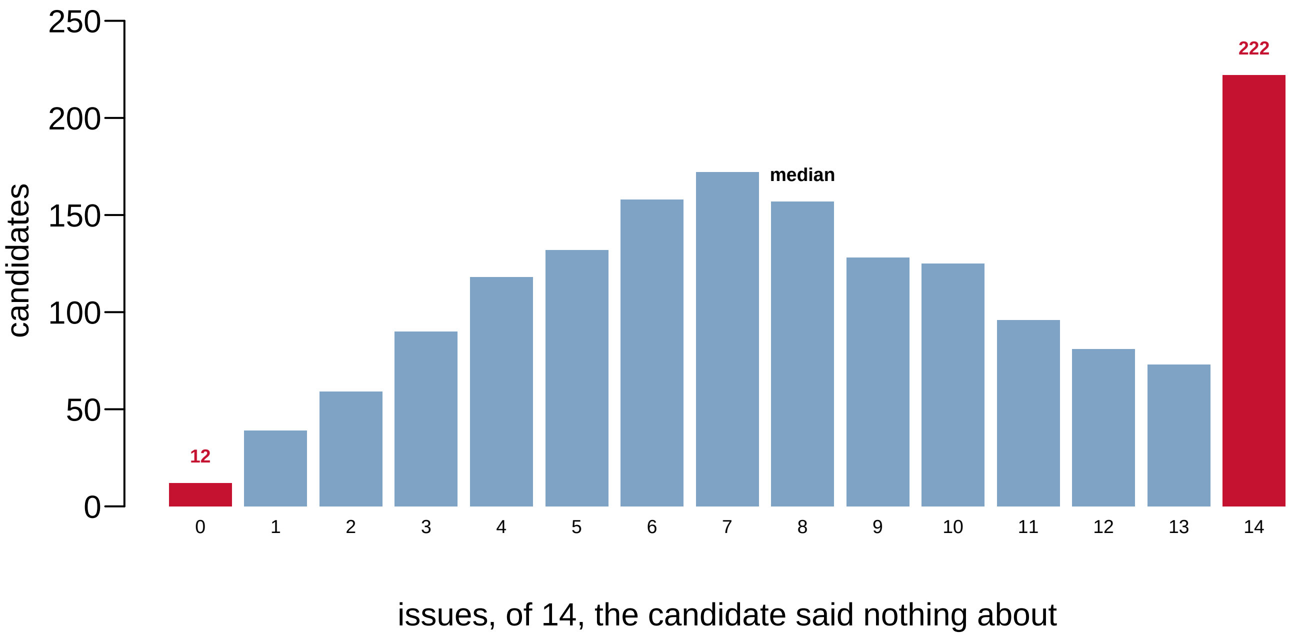

Figure 2. How many of the 14 issues each candidate left unmentioned. The two ends are marked. 12 candidates addressed all 14, which is 0.7 percent of the field. 222 addressed none of them.

Most readers guess something between two and four, because the candidates they can picture are the ones with a full issues page. The median is 8.

The shape matters more than the middle. This is not a hump around a typical candidate who covers most subjects. The distribution leans heavily toward silence, and its single tallest bar sits at the far end, where a candidate's website mentioned none of the fourteen.

Who those 222 are is the obvious next question, and the file will only half answer it. They are not all fringe candidates. Some are, and some are serious people whose sites were a biography and a donate button.

Primary result	House candidates	Mean issues left unmentioned
won	779	7.3
lost	638	9.2
runoff	26	8.5

House candidates who won their primary left 7.3 of the 14 unmentioned on average. The ones who lost left 9.2. **Do not read that as a finding about what wins primaries.** Winners include hundreds of incumbents running unopposed, who have both the least to gain from a policy page and the most staff to build one.

The comparison is here because it shows the variable is not noise. It moves with something. What it moves with, this file cannot say.

04 Two readers, one file, two answers

Everything above depends on reading the workbook correctly, and the workbook has a trap in it that has nothing to do with politics.

The party column holds 15 distinct strings for what are really about a dozen parties. Three candidates are filed under `D` followed by a space. One Republican is `G.O.P.` rather than `R`. One Libertarian is lowercase.

The string in the cell	Candidates	Characters
R	896	1
D	716	1
No Party Preference	19	19
Libertarian	7	11
Independent	6	11
Green	4	5
"D" <- ends in a space	3	2
Peace and Freedom	3	17
Alaskan Independence	2	20
G.O.P.	1	6
libertarian	1	11

Here is the part worth remembering. Whether that trailing space exists depends on which program opens the file.

Reading the same file with	Distinct party strings	What happens to "D"
Python, pandas.read_excel()	15	kept
R, readxl::read_excel()	14	silently trimmed to "D"
R, readxl::read_excel(trim_ws = FALSE)	15	kept

R's reader takes an argument called `trim_ws`, and its default is to trim. So the three dirty cells arrive clean, the defect disappears, and nothing warns you. Python's reader leaves them alone.

Neither program is wrong. They made different reasonable choices about whitespace, and the choice is documented in both. But a student who counts parties in R and a student who counts parties in Python will disagree, and neither will be able to see why.

This chapter reads the column **both ways** on purpose. The clean version is what the tables use. The 3 cells that differ between the two reads are counted, and the count is printed above rather than assumed.

05 The variable the codebook does not define

The Primaries Project existed to study factions. Its faction variable is the least documented column in the file.

Code	Candidates	Party	What the codebook says it means
11	261	Republican	— not defined —
22	251	Republican	Tea Party Republican
33	9	Republican	— not defined —
44	246	Republican	Conservative Republican
55	40	Republican	— not defined —
66	175	Democrat	— not defined —
77	308	Democrat	Establishment Democrat
88	66	Democrat	Moderate Democrat
99	47	Third party	— not defined —
100	259	Democrat and Republican	— not defined —

Ten codes. The codebook prints four of them, and it prints them under the heading *examples* rather than as a key. For codes 11, 33, 55, 66 and 99 there is no published definition anywhere.

Code 100 is the strangest. The codebook says of this variable, in its own words: *"label: PartyCatLabel, but 1 nonmissing value is not labeled"*. That unlabeled value holds 259 candidates, and unlike every other code it contains both Democrats and Republicans.

Altogether 791 (47.6%) of the candidates carry a faction code with no published meaning.

This chapter does not guess what they are. You could make a decent guess. Code 11 is Republican, common, and sits beside a code labelled Tea Party Republican, so "Establishment Republican" is a reasonable bet. A reasonable bet is not a codebook. What the file can tell you without guessing is which side of the aisle each code sits on, and that is the column above.

06 What this file cannot tell you

It cannot tell you what candidates believed. It is a record of published claims. A candidate who opposed the Affordable Care Act privately and never mentioned it is coded identically to one who had no view.

It cannot tell you why a subject was left out. Silence on immigration might be strategy in a district where the issue splits your coalition. It might be a website built in a weekend. It might be a candidate who talks about it constantly at events and never wrote it down. The code is the same in all three cases.

It cannot separate "no experience" from "we could not find out." The prior office variable is a 0/1 flag, and the codebook defines its zero as *"No or No information – never held office and not established are one code"*. Two very different candidates share that number.

It is one election. 2014 was a midterm with an unpopular Democratic president, a health law two years old, and a Republican faction fight in progress. Nothing here establishes that campaign websites are usually this empty.

And the websites are gone. The coding was done in 2014 from live sites. Most of those domains now redirect, expire or resolve to something unrelated. You cannot go back and check a single row. The file is the only surviving record of what it describes, which means its coding decisions cannot be audited by anyone, including its authors.

That last one deserves a moment. Every other source in this part of the book sits on top of a document you can still fetch. This one does not.

07 Why the silence is the point

There is a book about exactly this problem. *Parties on the Ground* (2026), by Kathleen Bawn, John Zaller and four colleagues, studies 55 winnable open-seat House nominations from this same 2013–14 cycle. It is built from 346 interviews with the people who ran them.

Its argument is that primaries do not work the way the reformers who created them intended. Progressives introduced the direct primary to move nominations from party bosses to voters. Bawn and her colleagues find that most voters pay too little attention to use it. The actors who matter are organised groups with intense policy demands, who vet candidates, vouch for them, and supply or withhold the money.

Figure 1 is the public half of that argument. If a voter wanted to do the job the primary was designed to give them, the campaign website is where they would start. On the median candidate they would find 6 of 14 subjects addressed. On 222 candidates they would find none.

The groups in Bawn's account are not reading websites. They are meeting candidates, asking direct questions, and getting answers that are never published. The information exists. It is simply not in any document a voter can open.

So the fourth code is doing more work than a missing-value flag. It is a measurement of how much of a nomination contest happens where the public cannot see it.

Two questions this file raises and cannot settle. Would a voter who read every website in their primary know enough to choose? And is a candidate who tells organised groups more than they tell voters doing something wrong, or just answering the people who actually asked?

08 Sources

Data

- Brookings Institution, *The Primaries Project* database, compiled by Elaine C. Kamarck and Alexander R. Podkul. https://www.brookings.edu/articles/the-2014-primaries-project-data/ — 1,662 candidates in the 2014 congressional primaries, 48 columns. The workbook's own metadata records it as last saved on 5 January 2015. **The fourteen issue positions are Brookings' coding, made by reading each candidate's campaign website during the 2014 cycle.**

- Brookings Institution, *Primaries Project Database Codebook*, December 2014. Distributed beside the data. Documents four of the ten values of the faction variable and states that one is unlabeled. Its own pointer to the companion report no longer resolves.

Both files are committed beside this chapter, because they are served from a content management system's upload folder rather than an archive, and neither carries a DOI.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`by_count.csv`, `by_issue.csv`, `by_outcome.csv`, `candidates.csv`, `factions.csv`, `party_labels.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `by_issue.csv` gives `silent_pct` for every issue. Sort it. Which issues do candidates most avoid, and is avoidance a position?
- `by_count.csv` counts candidates by how many issues they said nothing about. Find the typical candidate. How much of a platform is actually stated?
- `by_outcome.csv` compares `mean_silent` for winners and losers. Does saying less help you win a primary?
- `factions.csv` has a `documented` flag and codes with no label at all. Count the undocumented ones. What do you do with a variable whose codebook is incomplete?

The book this chapter is written against

- Bawn, K., Brown, K., Ocampo, A. X., Patterson, S. Jr., Ray, J. L. and Zaller, J. (2026). *Parties on the Ground: A Study of Nominations for the House of Representatives.* University of Chicago Press. doi:10.7208/chicago/9780226853314. The 55 contests, 346 interviews and the account of vetting are theirs. Its online appendix is announced at https://osf.io/s5982, which was not publicly readable when this chapter was written.

Related reading

- Kamarck, E. C. and Podkul, A. R. (2014). "The 2014 Congressional Primaries: Who Ran and Why." Brookings Institution.
- Hall, A. B. (2019). *Who Wants to Run? How the Devaluing of Political Office Drives Polarization.* University of Chicago Press.
- Boatright, R. G. (2013). *Getting Primaried: The Changing Politics of Congressional Primary Challenges.* University of Michigan Press.

How the figures were made

Every number printed above is computed from the two committed files at the moment this document was rendered. The party column is read twice, once with whitespace trimming turned off, and the difference between the two reads is reported rather than resolved silently. The validation results are reproduced in `checks.csv`.

REBUILD THIS DATA YOURSELF, WITH AN AI ASSISTANT

I want to rebuild a public dataset from its original source, without any starting files. Please do the following and show your work.

THE SOURCE. Brookings Institution, The Primaries Project, compiled by Elaine Kamarck’s team, who read the campaign website of every 2014 congressional primary candidate and coded what it said. One Excel workbook, about 640 KB, in the old .xls format:

https://www.brookings.edu/wp-content/uploads/2016/06/PrimariesDatabase_Clean.xls

and its codebook, a Word document in the same folder, PrimariesCodeBook_clean.docx. Each row is one candidate. No account or key is needed. The file was last saved in January 2015 and will never be updated; the address is a content-management system’s uploads folder, so the risk is the link dying, not the data changing.

HOW TO FETCH IT. Nothing blocks the download, but watch your reader: some tools silently trim whitespace from text cells. Read the party column exactly as stored at least once, because a few cells end in a space, and whether they exist depends on who is looking.

WHAT TO BUILD.

1. A count of the rows and columns, and the distinct strings in the party column — first exactly as the file stores them, then after trimming — reporting how many the trim makes disappear.

2. From the fourteen issue columns, each coded 1 to 4, where 4 means the website never mentioned the issue: how many of the fourteen each candidate is silent on, its median, and how many candidates are silent on all fourteen and on none.

CHECK YOUR WORK. The copy this book used was fetched in August 2026. If your rebuild matches it, you should find:

- 1,662 candidates and 48 columns
- 15 distinct party strings exactly as stored; a whitespace-trimming reader sees only 14
- the median candidate is silent on 8 of the 14 issues
- 222 candidates are silent on all fourteen, 12 on none

The workbook has not changed since 2015 and should not change again; any difference means you are holding a different file.