

The Age of Members of Congress

The denominator that decides the answer

84-355 Democracy's Data

August 2026

The Constitution sets a floor on the age of American officeholders and no ceiling at all. A representative must be 25, a senator 30, a president 35. Above those numbers the document is silent: no retirement age, no maximum term, no mechanism for removing somebody who has simply stayed, short of an election or a death.

So the age of American government is not set by anything. It is a residue — what is left over after nominations, primaries, incumbency, party recruitment, and several thousand private decisions about when to stop. In the 119th Congress that residue is a median member aged **59** and a median senator aged **65**, both at or near the highest figures in a series that begins in 1789.

Say that in company and somebody will answer within a sentence. *Of course they are older. Everybody is older. Life expectancy has gone up thirty years.*

It is a good objection, offered in good faith, and it is testable. Testing it turns out to reverse it. The reversal has nothing to do with Congress and everything to do with which row of a government table you happen to read.

Age is not a proxy for competence and nothing here treats it as one. It matters for a narrower and more defensible reason. **Age is a rough guide to how long ago somebody's politics were formed, and to how much of the future they will personally be present for.**

A legislature whose median member entered public life forty years ago is making binding decisions about a country. Most of that country's residents were not born then. Whether that is a problem is a political judgment. Whether it is true is a question about data.

01 Three files that were never meant to meet

This chapter runs on three sources with nothing in common except a calendar.

The first is Voteview's member table (Lewis, Poole, Rosenthal, Boche, Rudkin and Sonnet): one row per member per Congress from the first to the 119th, maintained at UCLA as academic infrastructure. It was built to serve DW-NOMINATE, a scaling of roll-call votes that has nothing to do with age; the biographical columns, ~~born~~ among them, are there because a scaling has to know who is who. No statute required any of it, and nothing keeps it current except a discipline's need for it.

Calling that a source understates what it is. Voteview does not receive a member table from anybody. It builds one.

The primary record is roll calls, chamber by chamber over two centuries, in formats that changed as often as the printing did. Turning those into tidy rows meant three things. Deciding that two spellings are one man. Attaching a number to every member, because the government's own records carry no usable one. And coding a party for people who served before parties had settled names.

That labour is the only reason the file opens in a spreadsheet at all. It also means **the cleaning happened decades before this chapter did, by people making judgment calls the file never records**. What arrives looks like a record of who served. It is a research team's considered answer to that question.

One of those calls leaves a mark on a number used below. When a member changes party in the middle of a Congress, Voteview does not amend his row. It issues him a second identifier and a second row, so he sits in the file twice for that Congress under two party codes. It happens 22 times: Jeffords, Specter, Manchin, and the group of Democrats who crossed over after 1994. Nothing in the file flags it.

This chapter dates a member's arrival by the first Congress in which his identifier appears. So a switcher's new identifier reads as a freshman. Billy Tauzin entered the House in 1979, and appears in the 104th twice: once with 16 years of service, and once with none at all.

Those few extra rows, against 51,062, move no median anywhere in this chapter. The point is not the size of the error. It is that the error entered through a decision somebody else made, for a purpose that was not this one. A reader who never opens the build will never learn it happened.

The second and third come from the National Center for Health Statistics, a federal statistical agency doing precisely what it was created to do. Life expectancy at birth arrives from CDC's open-data catalog as a machine-readable annual series. Remaining life expectancy at 65 and at 75 does not: that series sits in Trend Table 4 of *Health, United States, 2019*, a published PDF, and had to be parsed out of one. Two rows of the same table, two completely different amounts of work to obtain — and, as this chapter is about to argue, two different answers.

The parse is the fragile step, and worth being explicit about why. A file pulled from a catalogue can be fetched again next year and compared byte for byte against the copy read here. A table lifted out of a PDF cannot. What would need checking is not a file but a *reading* of one, and the reading lives in a script rather than on the page.

So the build checks it sideways, against the series it can fetch again. There are 26 years where the PDF and the CDC catalogue both publish life expectancy at birth. The two agree exactly in 24 of them, and differ by at most 0.4 of a year in the rest.

That agreement is the entire warrant that the column the chapter turns on was read off the page correctly.

What makes these the right evidence is also what makes them risky. The objection being tested is that Congress is older because everybody is older. Testing it means putting a congressional series and a mortality series on one axis, when the people who built them never imagined the two would meet. The join is calendar year, and it is not automatic. A Congress convenes in an odd year and sits through the even one after it, while life tables are published for calendar years. So this chapter states the correspondence outright rather than letting a merge guess.

Better evidence does not exist. Nobody publishes a life table for the population of people who become senators.

Underneath sits a problem no file can fix. “Old” has no denominator here. There is no bottom number to divide by, no agreed figure that separates old from not. The Constitution supplies a floor of twenty-five, thirty and thirty-five, and no ceiling whatever. So any line at which somebody counts as old is the analyst’s choice, not the country’s.

And the two mortality series describe different populations. That is the entire mechanism of the reversal ahead. One of them is a made-up figure for everybody born in a given year, dominated by deaths in infancy, and nobody in Congress is an infant.

The cleaning is mostly assertion. All three sources are fetched, and what is believed about them is stated before anything is written. That the first Congress convenes in 1789 and the 119th in 2025. And that no person in the file carries two conflicting birth years, a count that comes to 0. That last one is the check that would have caught a merge gone wrong.

Exactly 1 person has a birth year on some of his rows and not on others. The build names him rather than patching him: CLEVELAND, Stephen Grover. Nothing is de-duplicated, because a member who serves twenty years is supposed to appear ten times. One row is a member in one Congress, not a member.

And every fetch is logged with its byte count and a fingerprint, so a reader in ten years can establish whether the file they downloaded is the file this chapter read. That is a question this book keeps asking and that almost no published analysis can answer about itself.

02 One column, and what a birth year buys you

Every age in this chapter comes from a single column of a single file: the Voteview member table, one row per member per Congress since 1789, with a `born` column.

Fact	Value
Rows in the file	51,062
Congresses covered	1 to 119 (1789 to 2025)
Chambers in it	House, Senate, President
Presidents in it	45 people, 47 presidencies
Rows with no birth year	176
What <code>born</code> contains	a year. Not a date.

`born` is a year, so every age below is approximate to within a year. Age here means the year a Congress convened minus the year a member was born. For anybody with a December birthday, that overstates their age on opening day by one. The error is not hypothetical and it is easy to watch happen. Joe Biden was born in 1942 and was inaugurated in January 2021, aged 78:

President	Born	Term began	Age by year arithmetic	Actual age on inauguration day
Biden	1942	2021	79	78

One year off, in a case where the true answer is a matter of public record. Across 51,062 rows those errors run about half a year on average and point in no particular direction, so medians and trends survive them. Individual ages do not, and nothing in this chapter turns on one person's age.

There is a second warning in the same file, quieter. Search it for ADAMS, John and two different men come back: the president, whose birth year is unrecorded here, and a Massachusetts representative born in 1778. **Names are not identifiers.** The file supplies a numeric one, icpsr, and everything below uses it.

And a third, of the kind you only find by going looking for it:

Check	Value
people with a birth year on some rows and not others	1
who that is	CLEVELAND, Stephen Grover
people carrying two different birth years	0

Grover Cleveland has a birth year on his first-term rows and a blank on his second-term rows — the same man, the same file, one number present and then absent. No two rows anywhere in 51,062 disagree about a birth year; exactly one person is half-missing, and nothing announces it. An average would skip the blank rows and return a number that looked entirely reasonable. That is why the build script counts this rather than quietly patching it.

03 Twenty-two columns arrive and four of them are the chapter

Everything above is about what the file gets wrong. This is about what it is *shaped* like, because before a median can be taken of anything the file has to stop being what Vote-view wrote and start being a table about age.

Here is the top of HSal1_members.csv exactly as it comes off the server. The header is one physical line in the file and is far too long for this page, so it and the rows under it are folded at their commas; nothing else is changed.

Column	What it is	Record 1	Record 2	Record 3
congress	the numbered two-year Congress this row describes	1	1	1
chamber	House, Senate or President	President	House	House
icpsr	the member's permanent ID, kept across every Congress they serve	99869	4766	8457
state_icpsr	ICPSR's own state code, not the Census FIPS code	99	1	1
district_code	district number; 0 for senators and at-large members	0	98	98
state_abbrev	two-letter postal abbreviation	USA	CT	CT

Column	What it is	Record 1	Record 2	Record 3
party_code	ICPSR party code: 100 Democrat, 200 Republican	5000	5000	5000
occupancy	which occupant of the seat this is in this Congress	(empty)	(empty)	(empty)
last_means	how the member reached office	(empty)	(empty)	(empty)
bioname	name, as last, first middle	"WASHINGTON, George"	"HUNTINGTON, Benjamin"	"SHERMAN, Roger"
bioguide_id	Biographical Directory ID — the join key to other sources	(empty)	H000995	S000349
born	year of birth — the column this chapter is built on	(empty)	1736.0	1721.0
died	year of death, empty for the living	(empty)	1800.0	1793.0
nominate_dim1	DW-NOMINATE first dimension	(empty)	(empty)	0.589
nominate_dim2	DW-NOMINATE second dimension	(empty)	(empty)	0.307
nominate_log_likelihood	how well the scaling fits this member's votes	(empty)	(empty)	(empty)
nominate_geo_mean_probability	average probability the model gave the votes cast	(empty)	(empty)	(empty)
nominate_number_of_votes	votes the score was estimated from	(empty)	(empty)	(empty)
nominate_number_of_errors	votes the model classifies wrongly	(empty)	(empty)	(empty)
conditional	documented, and empty on every row	(empty)	(empty)	(empty)
nokken_poole_dim1	dimension 1, re-estimated within each Congress	(empty)	(empty)	(empty)
nokken_poole_dim2	the same, for dimension 2	(empty)	(empty)	(empty)

Nothing in that block was written for this chapter. Of its 22 columns, 9 are the roll-call scaling the file exists to serve: ideal points, likelihoods and error counts. Not one of them is read here. The four this chapter needs are `congress`, `chamber`, `icpsr` and `born`. They are in the file so a scaling can tell one member from another.

Then read the block for what is not in it. **There is no year.** A row is dated by a Congress number, and the arithmetic that turns 1 into 1789 is the analyst's, not the file's. **There is no age.** Age is a subtraction nobody has performed.

born arrives as 1736.0. That is a calendar year written as a decimal, because a column with any blank in it reads as numeric. It will not join to anything until somebody makes it a whole number.

Washington's row is empty from bioguide_id rightward, including `born`, which is why 176 rows in the finished file still have no age.

And **bioname contains a comma**, inside quotes. That is the trap that eats a whole file if it is split carelessly, and the build says so in a comment before it reads a byte.

Now the same three rows after the build. This is the clean unit: one row per member per Congress, with the columns that were derived rather than supplied.

Congress	Year	Chamber	lcpstr	Bioname	Born	Age	Entry year	Years since entry
1	1789	President	99869	WASHINGTON, George	NA	NA	NA	NA
1	1789	House	4766	HUNTINGTON, Benjamin	1736	53	1789	0
1	1789	House	8457	SHERMAN, Roger	1721	68	1789	0

Four of those nine columns are not in the download. `year` is the claim $1787 + 2 \times \text{congress}$, checked afterwards against two fixed points rather than trusted.

`age` is `year - born`. It inherits the year-only precision of `born`, so it is right to within twelve months and no better. `entry_year` required grouping the whole file by member and chamber and taking a minimum. `years_since_entry` is the difference. Those two columns are what let a later figure split “older” into arriving older and staying longer.

Washington keeps his row and carries NA four times over. Dropping him would have been tidier and would have hidden the gap.

The counts are the rest of the story. The raw file has 51,062 rows, and so does the clean member file, all 51,062 of them. **Nothing is dropped and nothing is de-duplicated.**

A member who serves twenty years is supposed to appear ten times. One row is a member in one Congress, not a member. Removing duplicates anywhere in here would quietly answer a different question.

What the build does drop is 129 presidential rows, and only from the totals. Presidents get their own file, because a president is not a legislator and averaging him into a chamber would be a mistake of kind.

The last reshaping is a collapse, and it invents a category. From 50,933 member-Congress rows the build writes 357 rows of `derived/age_by_congress.csv`, one per chamber per Congress – House, Senate, and a third value called Congress that appears nowhere in the download. It is the two chambers pooled, so every member-Congress row is counted twice in that file, once in his chamber and once in the whole. That is a defensible choice and an invisible one: a reader who sees three chambers in a file that has two will not learn from the file which one is real.

Congress	Year	Chamber	N	Median age	Mean age	Mean entry age	Mean years since entry
1	1789	Congress	95	45	44.8000	44.8000	0
1	1789	House	66	44	43.9545	43.9545	0
1	1789	Senate	29	48	46.7241	46.7241	0

04 When it happened




Figure 1. Two and a third centuries of congressional age. Median age of each chamber at every Congress, with the shaded band covering the middle half of the chamber. One green dot marks each president at the start of a term. The Senate runs above the House throughout. The recent rise is steep and short. The median member was 50 in 1979, and 60 by 2017.

Quantity	Value
Youngest Congress on record (median age)	42 (1793)
Oldest Congresses on record (median age)	60 (2017, 2019, 2021)
Median age, most recent Congress	59 (2025)
Post-war low point	50 (1979)
Median senator now	65
Median representative now	58

Three things are visible in that picture and only one of them is the thing people expect to see.

First, the rise is real but it is not a long slide. Congress stood at 50 in 1979, the first time it had been that young since 1913, 66 years earlier. It reached 60 by 2017. Almost the entire increase people argue about happened inside a single working lifetime, after 1980.

Second, the Senate has always been older, by about 7 years today and by a similar gap for most of two centuries. That is the chamber doing what it was designed to do: six-year terms, a higher age floor, and until 1913 no direct election at all.

Third, the presidents are a scatter rather than a trend. Individual presidents are single draws, and forty-odd draws across two centuries will not support a line. What can be said is arithmetic: 55 was the average age at which a president took office before 1900, and 64 since 1981.

05 Everyone lives longer now

Here is the objection, stated as strongly as it deserves. Between 1950 and 2018 American life expectancy rose from **68.2** years to **78.7**, a gain of **10.5 years**. Over exactly the same span the median member of Congress aged from 53 to 60, a gain of **7 years**.

A gain of 7 is less than a gain of 10.5. On that comparison Congress has not aged relative to the lifespan it is drawn from; it has aged *less* than the country it governs. The objection does not merely survive, it wins outright, and the chapter could stop here.

It is wrong, and it is wrong in a way worth learning, because the error is not in the arithmetic and not in the source. It is in the choice of which number counts as “how long people live.”

Before you read on, commit to a number. You have two of the three numbers this argument turns on. Between 1950 and 2018 the median member of Congress aged 7 years, while life expectancy at birth rose 10.5 years. There is a third way to measure how long people live, and it is the one that actually applies to a sitting senator. **Over those same 68 years, did it rise by more than 10.5, by something between 7 and 10.5, or by less than 7?** Write down a number of years before you read on.

06 The row further down the same page

78.7 is life expectancy **at birth**. It is the average length of a life for a baby born in 2018, and the reason it rose so far across the twentieth century is mostly that babies and children stopped dying. Antibiotics, sanitation, vaccination, obstetrics: enormous achievements, every one of them irrelevant to a sitting senator, who has already survived all of it.

The quantity that describes an officeholder’s remaining life is **remaining life expectancy at an age they have actually reached** — how many more years a 65-year-old can expect. The National Center for Health Statistics publishes it in the same table as the at-birth figure, one section further down the page.

Measure	1950	2018	Change
Life expectancy at birth	68.2	78.7	+10.5
Remaining life expectancy at 65	13.9	19.5	+5.6
Remaining life expectancy at 75	–	12.3	+1.9
Median age of Congress	53	60	+7

Most readers write down something at or above 10.5, reasoning that if lives got that much longer, the years after 65 must have got longer too. The gain is 5.6 years, which is smaller than the 7 years Congress aged over the same span.




Figure 2. The same 68 years, measured against two rows of one government table. Here are the gains between 1950 and 2018. Life expectancy at birth rose 10.5 years. The median age of Congress rose 7. The remaining life expectancy of a 65-year-old rose 5.6.

Congress aged less than the first and more than the third. Only the third is about people who have already reached the age members of Congress are.

Congress aged 7 years; a 65-year-old's remaining life expectancy grew 5.6. Against the honest yardstick Congress aged *faster* than longevity, not slower. Against the at-birth yardstick it aged slower. Same data, same source, same printed table, opposite conclusions, and the only thing that changed between them was which row was read.

One further number makes the mechanism concrete.

Quantity	Value
Gain in life expectancy at birth, 1950-2018	10.5 years
Gain in remaining life expectancy at 65, 1950-2018	5.6 years
Share of the headline gain that is about life after 65	53%

At most 53% of the famous rise in life expectancy is a rise in how long a 65-year-old lives. The remainder is people who used to die young and now do not. That transforms the average life, and changes a senator's prospects not at all.

The lesson is not really about actuarial tables. It is that a comparison has a denominator, that the denominator is a choice, and that a denominator chosen without thinking can reverse the sign of the answer. It is the same move as "crime per capita" against "crime per capita among people who go outside."

Two caveats belong beside that result rather than after it, because anybody who uses these numbers will run into both. The first is that the at-65 series is **sparse**. It is published for 25 years between 1950 and 2018. Not for every year, and not at all before 1950. Nothing between the published years has been filled in here.

The comparison above can be made for 1950 onwards and cannot honestly be made for 1900. That is a shame. 1900 is exactly where the at-birth series looks most impressive, and where the temptation is greatest.

The second is that the two NCHS publications used here disagree with each other.

Check	Value
years where HUS and CDC both give $e(0)$	26
of those, years where the two agree exactly	24
years where they disagree	2003, 2004
largest HUS-vs-CDC disagreement on $e(0)$, years	0.4

Two NCHS products publish life expectancy at birth. On 24 of 26 shared years they agree to the tenth. On two of them, 2003 and 2004, they do not. The printed trend table runs 77.2 then 77.6. The open-data file runs 77.6 then 77.5.

One rises and one does not, which is what a swapped pair of numbers looks like. Nothing here settles it. It is reported because a reader told that the two sources agree everywhere has been told something false.

07 Arriving older, or staying longer

A median age can rise two ways: new members arrive older, or sitting members stay longer. These are different phenomena with different causes and different remedies, and people assert one or the other constantly without checking which.

Checking is possible because the two add up exactly. For every member, at every Congress:

$$\text{age now} = \text{age when they first arrived} + \text{years since}$$

That is an identity, not a model, so it survives averaging. Average each term across a chamber and its mean age decomposes into exactly two pieces, with nothing left over.

Every Congress, split into the two things that make it old




Figure 3. Mean age of Congress, split into its only two components. The lower band is the average age at which sitting members first arrived; the upper band is the average time they have been there since. The two sum exactly to the mean age, drawn as the dark line, because the split is an identity rather than a model. Between 1950 and 2018 the mean rose +6.10 years, of which +3.86 is arrival age and +2.25 is time served.

Chamber	Rise in mean age	From entering older	From staying longer	Entering older, share
Congress	+6.10	+3.86	+2.25	63%
House	+6.28	+4.25	+2.03	68%
Senate	+5.55	+2.40	+3.15	43%

For Congress as a whole, 63% of the rise between 1950 and 2018 is people arriving older. The popular explanation, that incumbents cling on, is the smaller half of it.

And the two chambers disagree. In the House, 68% of the rise is arrival age. In the Senate it is 43%. So **most of the Senate's ageing really is senators staying**, which is what six-year terms and very high re-election rates would predict. A single national figure conceals two different institutions.

The decomposition also explains the dip in Figure 1. Look at the blue band around 1970.

Period	Change in mean age	Change in arrival age	Change in years since arriving
1950 to 1980	-2.68	-3.50	+0.82
1980 to 2025	+8.71	+7.45	+1.26

Congress got *younger* through the 1960s and 1970s, and it did so almost entirely by electing younger people: mean arrival age fell 3.5 years while time in office barely moved. Then the same lever went the other way. Whatever produced the Watergate-era freshman classes stopped producing them, and the recovery of mean age since 1980 has come mostly through the same channel that lowered it.

08 The median hides the tail

A median describes the middle member and says nothing about the shape around them. The shape is where the argument actually lives, because nobody is worried about a 59-year-old.

The whole chamber, not the middle of it




Figure 4. The full age distribution of the House and Senate together, by era. One curve per era, over all 50,768 member-Congress records with a recorded age. The median and the share aged 70 or over are printed beside each. The dashed line marks age 70. The median moves from 43.0 in the first era to 57.0 in the last, while the share over 70 moves from 0.8% to 12.6%: the right tail travels further than the middle does.

Quantity	1950	2025
Share of Congress aged 65 or over	16.1%	34.9%
Share of Congress aged 70 or over	8.5%	24.1%
Share of Congress aged 80 or over	1.1%	4.4%
Share of the SENATE aged 70 or over	16.5%	35.9%

Between 1950 and 2025 the median moved 6 years and the tail moved much further. The share of Congress over 70 went from 8.5% to 24.1%. In the Senate it now stands at 35.9%, better than one senator in three. The over-80 share was 1.1% in 1950 and is 4.4% now.

Take the two Congresses in that table and compare them at three points instead of one.

Congress	Year	5th percentile	Median	95th percentile
81	1950	36	53	72
119	2025	39	59	79
	moved by	+3	+6	+7

The young end moved 3 years, the middle 6, and the old end 7. **The distribution did not slide; it stretched, and it stretched at the top.** Any statistic that reports only the middle will understate what changed.

Over the full two centuries the picture is different again. The young end moves almost as much as the old one, so the stretching is a feature of the modern period rather than of the whole record. Figure 4 is worth reading twice for that reason.

09 The careers you cannot see

If members stay longer, careers should be getting longer, and they are.

Entered in	Careers	Median age on arrival	Median years served	Served 20 years or more
1789-1848	2,901	41	4	1.1%
1849-1908	4,392	45	4	3.7%
1909-1948	2,578	48	6	12.1%
1949-1985	1,688	44	10	21.1%

Notice what that table had to exclude. It counts only people who entered by 1985 and drops anybody still serving. Both exclusions are necessary and both cost something. **A career you can measure is a career that is over.**

Ask “how long do members serve?” of everyone in the file, and the recent groups come back looking short. Not because service got shorter, but because their careers are still running.

This is right-censoring, and it is the commonest way a tenure statistic gets reported backwards.




Figure 5. Age on arrival against length of career, for careers that are finished. One dot per completed chamber career, nudged apart and coloured by when it began. The group medians are drawn large. Median arrival age moves right across the three groups: 43, 46, 48. Median completed service goes 4, 8, 8 years.

The last of those is a side effect of the cut-off rather than a finding. The most recent group can only contain careers that have already ended.

The cloud tilts the way it should. Of the careers that ran thirty years or more, 72.5% began at 45 or younger; their median arrival age is 42 against 45 for careers in general. **Arriving young is close to a precondition for a long career**, which makes the rise in arrival age a plausible brake on future tenure as well as a cause of present age.

One popular explanation for senators arriving old does not survive contact with the file. Senators do increasingly come up through the House:

Quantity	Value
Senators who sat in the House first, all time	31.7%
Senators who sat in the House first, entering since 1981	42.1%
Mean Senate arrival age since 1981, WITH prior House service	51.6
Mean Senate arrival age since 1981, WITHOUT it	53.3

Senators who came up through the House arrive younger, not older. The pipeline is not what is aging the Senate. Something else is, and this file does not say what.

10 Who leaves the data before you can count them

There is one more reason for care with any statement about how old officeholders become, and it is a selection problem rather than a measurement one.

Era	Member-terms ending in death in office
1900s-1910s	2.90%
1920s-1930s	3.58%
1960s-1970s	1.54%
2010s-2020s	0.57%

Members used to die in office roughly 6.3 times as often as they do now, and every one of those deaths removed an old member from every subsequent Congress in this file.

That biases the historical series downward. Some of the apparent youth of Congress a century ago is not a preference for young people; it is that old members stopped being members by dying. And it is precisely the part of the change that longevity *does* explain, which is why it belongs in the same argument as the denominator above, on the other side of the ledger. The honest version of “people live longer now” is not that it accounts for the rise. It is that it accounts for a piece of the rise that nobody counts, in the historical comparison rather than the modern one.

11 The argument about all of this

The measurement is not seriously disputed. Everything downstream of it is.

That officeholders are older is agreed by everyone who has looked. What people disagree about is the right comparison. To the population’s age structure, to the electorate’s, to the pool of plausible candidates, or to nothing at all. Each of those choices gives a different sense of how unusual the current Congress is. The American population has also aged over this period, for reasons having nothing to do with elections, and no series here measures it. Whether the right question is “how old is Congress” or “how much older than the country is Congress” is a judgment this data cannot settle.

There is also a serious methodological objection to everything above. Median age is a crude summary of a body whose members do not have interchangeable roles. Seniority concentrates power, so the age of a chamber’s *leadership* and of its committee chairs may matter far more than the age of its median member, and neither is measured here. Anybody arguing from Figure 4 should be asked why the median is the right statistic to be arguing from.

And there is a limitation no amount of care fixes: the counterfactual is unobservable. There is no country where the same electorate ran the same elections with a different spread of candidate ages. So any claim that Congress is “too old” is a claim about a world for which nobody has data.

12 What this evidence will bear

The Voteview file is a complete enumeration rather than a sample: every member of every Congress since 1789, with a stable numeric identifier. The life-expectancy series are official NCHS publications. Every number above is recomputed from those files each time this document is rendered.

Its limits bound different claims, and they are worth taking one at a time. **born is a year**, so every age carries about a year of noise, which trends absorb and individuals do not.

176 rows have no birth year at all, concentrated in the first twenty Congresses, along with four presidents. So the earliest part of Figure 1 rests on fewer people than it appears to.

The at-65 series is sparse and starts in 1950. So the most tempting version of the comparison, 1900 against today, cannot be made honestly and is not made here.

Two NCHS sources disagree in two years, which is small and real and reported rather than resolved. **Deaths in office censor the historical series**, in the direction that flatters the “Congress used to be young” reading.

And every age is measured **at the start of a Congress** rather than on any particular day. A member who arrived mid-term is dated to the opening of the Congress they joined.

Three questions stay genuinely open. Whether the rise continues or has peaked is unknown; the last three Congresses are flat, which is one observation rather than a trend. What produced the change in arrival age is the larger half of the whole effect and the half nobody has a settled explanation for. And whether any of it affects what the institution actually does is untouched by this file, because no measure of performance appears anywhere in it.

What can be said is bounded and worth saying. Congress is at or near its oldest recorded point. The rise is recent, mostly after 1980, and mostly a matter of people arriving older rather than staying longer. The Senate is the exception, where the reverse holds.

And “people just live longer now” does not account for it. Between 1950 and 2018 the median member of Congress aged 7 years, while the remaining life expectancy of a 65-year-old grew 5.6. That comparison runs the opposite way from the one everybody reaches for, using the same table, the same source and the same two dates.

13 Sources

Data

- Lewis, J.B., Poole, K., Rosenthal, H., Boche, A., Rudkin, A. and Sonnet, L. *Voteview: Congressional Roll-Call Votes Database*. https://voteview.com/static/data/out/members/HSall_members.csv 51,062 member-Congress records, Congresses 1 to 119.
- National Center for Health Statistics. *Life expectancy at birth, by race and sex, 1900-2018*. CDC open data. <https://data.cdc.gov/api/views/w9j2-ggv5/rows.csv?accessType=DOWNLOAD>
- National Center for Health Statistics. *Health, United States, 2019*, Trend Table 4: life expectancy at birth, at age 65 and at age 75. <https://www.cdc.gov/nchs/data/hus/2019/004-508.pdf> Parsed from the published PDF and checked against the open-data series above.

All three are fetched from the addresses above, reduced to the files this chapter reads, and the counts quoted in the opening pages recorded.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`age_by_congress.csv`, `entry_tenure.csv`, `life_expectancy.csv`, `members.csv`, `presidents.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `age_by_congress.csv` has `median_age` and `mean_entry_age`. Plot both. Is Congress ageing because members arrive older, or because they stay longer?
- Join `age_by_congress.csv` to `life_expectancy.csv` by year. Compare the rise in median age against the rise in `e0_annual`. Is Congress ageing faster than the country?
- `entry_tenure.csv` has `entry_age` and `years_served` for every career. Group by entry decade. Has the path into Congress changed?
- `age_by_congress.csv` splits by chamber. Compare House and Senate. Which is older, and is the gap growing?

REBUILD THIS DATA YOURSELF, WITH AN AI ASSISTANT

I want to rebuild a public dataset from its original sources, without any starting files. Please do the following and show your work.

THE FIRST SOURCE. Voteview, the congressional roll-call database. One CSV, about 6 MB, no account or key:

https://voteview.com/static/data/out/members/HSall_members.csv

One row per member per Congress, every Congress since 1789, presidents included. Birth and death arrive as YEARS, not dates, so every age computed from it is approximate to a year. Congress *n* begins in the year 1787 plus twice *n*.

THE SECOND SOURCE. CDC/NCHS, life expectancy at birth, annual back to 1900, from the CDC's open-data portal:

<https://data.cdc.gov/api/views/w9j2-ggv5/rows.csv?accessType=DOWNLOAD>

A small CSV: three race groups by three sex groups by year.

THE THIRD SOURCE, AND THE CATCH. Life expectancy AT AGE 65 – the honest comparison for an officeholder, who has already survived everything that kills children – is NOT in the open-data portal. It lives in a published PDF: "Health, United States, 2019", Trend

Table 4, at <https://www.cdc.gov/nchs/data/hus/2019/004-508.pdf>

The table has a real text layer and extracts with any PDF-to-text tool. It covers selected years from 1950 to 2018 only, and skips 2011.

WHAT TO BUILD.

1. Every member's approximate age in every Congress they served, and each president's age on taking office.
2. One life-expectancy table joining the annual at-birth series with the selected-year at-65 and at-75 series, so either can be held up against the ages in the first table – and the choice between them can be seen to matter.

CHECK YOUR WORK. The copies this book used were fetched in August 2026. If your rebuild matches them, you should find:

- 51,062 rows in the Voteview file, Congresses 1 through 119
- 176 member-Congress rows carry no birth year, and 45 presidents are in the file
- 1,071 rows in the at-birth file, and 25 published years of at-65 values, 1950 through 2018

Voteview re-estimates continuously and the CDC extends its series; larger counts mean the sources moved on, not that you erred.