

Ideal Points from Supreme Court Agreement

Building an ideal-point score for the Supreme Court out of nothing but agreement

84-355 Democracy's Data

August 2026

“The most conservative Court since the 1930s.”

Sentences like that are written constantly, and every one of them quietly requires a number. Nine people who never stand for election, answer no surveys and give no interviews about their views are nonetheless described in print by adjectives that nobody defines and everybody uses: *moderate*, *conservative*, *the swing vote*. Those adjectives decide how nominations are fought and how confirmation votes are cast.

The Supreme Court is the easiest place in American politics to put a number behind them. Nine people, every vote recorded, every case public, and a database that has counted all of it since 1946. So the question this chapter asks is not where the justices sit, which everybody thinks they know. It is whether a number that places them can be built **without anybody supplying an opinion about ideology** — and, if one can, what exactly that number is made of.

It matters because the Court is the only branch that never has to explain itself in electoral terms, which makes the informal vocabulary unusually load-bearing. Putting a number behind it changes the argument. Segal and Cover (1989) derived justices' values from pre-confirmation newspaper editorials, that is, from *perceptions* recorded before any vote was cast, and found they predicted votes well. Martin and Quinn (2002) built the standard dynamic score directly from votes, allowing positions to move over time. Both turned a rhetorical claim into a measurable one, and both are now routine inputs to research on judicial behavior (Epstein, Landes and Posner, 2013).

01 The votes, not the scores

Everything here comes from the Supreme Court Database, held at Washington University in St. Louis: every justice's vote in every case the Court has decided since 1946. It is taken from the justice-centered file organized by citation, 2024 release 01.

It is unlike most of the sources in this book, because no statute required it and no organization produced it while doing something else. Harold Spaeth started it as a research project, the National Science Foundation paid for it, and academics have maintained it ever since. It exists because somebody decided the Court's behavior was worth writing down in a form that could be counted. The releases are numbered and archived, so naming one identifies exactly what was read, which is more than a versioned news graphic

can offer. It is also, in 2026, served over `http` and not `https`, a small fact about maintenance that will eventually become a large one.

Choosing this file rather than an easier one is most of the chapter's argument. Martin-Quinn scores are published, updated every term, and would answer *where does each justice sit* in a single download. They are not used, for two reasons.

The practical one is that the site sits behind a bot check and will not answer a script. The real one is that **a published score is the answer, and this chapter is about how the answer gets made.** The database supplies the votes those scores are computed from, so the ordering recovered below can be checked, broken on purpose and argued with rather than accepted on authority. The data best suited to a question is not always the data closest to it.

The difficulty is that the most important column in the file is not an observation. `direction` records whether a decision came out in the conservative or the liberal direction. That is a coder's reading of a written opinion against a codebook, a rule specifying issue area by issue area which outcome counts as which. Where the codebook supplies no rule the field is left blank.

That is not a defect anybody could repair. Nothing measures the ideological direction of a legal holding. Every number ever published about the Court's ideology rests on somebody having read the cases and decided.

What can be asked is whether the rule was applied the same way across 78 terms of American legal history. That is a question about a codebook rather than about a Court.

`direction` is not the exception in this file. It is the rule, made visible. Issue area, the constitutional provision at stake, which party prevailed, the reason the Court took the case. Each of those is likewise a person's reading of a document, entered against a rule somebody wrote. **The Supreme Court publishes opinions. It publishes no data.**

Everything countable about the Court was made countable by someone outside it, decades after some of the cases were decided. The resulting file arrives looking like a record of what happened. It is a record of what a coding team concluded had happened.

That is the price of having a dataset at all here, and it is not avoidable. It is also entirely invisible from inside the CSV, where a judgment and an observation are both just a number in a column.

The cleaning is short and consequential. The release is downloaded whole, and the terms from 2020 to 2023 are kept. Every case whose `direction` is blank is dropped from the agreement counts, because "voted the same way" has no meaning when there is no direction to compare.

The build then writes the four small tables this chapter reads. One choice in that build does more work than the rest: for every pair of justices it records the number of cases as well as the number of agreements, rather than the percentage alone. That is why the pair who never sat together is visible in what follows instead of silently absent. Nothing was filled in at build time. The one fill in this chapter is made in the open, on the page, and undone again a few pages later.

02 Sixty-one columns, and nearly all of them are opinions

That claim, that a judgment and an observation look the same inside a CSV, is worth seeing rather than taking on trust. Here is the header of the download, SCDB_2024_01_justiceCentered_Citation.csv, folded at its commas:

Position	Column name	What the codebook says
1	caseId	identifier for the case
2	docketId	identifier for the case as docketed
3	caseIssuesId	identifier for one legal issue within the case
4	votId	identifier for one justice's vote — the grain of this file
5	dateDecision	the day the decision came down
6	decisionType	how the Court disposed of it: signed opinion, per curiam, and so on
7	usCite	citation in the United States Reports
8	sctCite	citation in the Supreme Court Reporter
9	ledCite	citation in the Lawyers' Edition
10	lexisCite	citation in LEXIS
11	term	the Court's term, which begins in October
12	naturalCourt	which stable set of nine justices was sitting
13	chief	the Chief Justice at the time
14	docket	the docket number the Court assigned
15	caseName	the case's name
16	dateArgument	the day it was argued
17	dateRearg	the day it was reargued, if it was
18	petitioner	coded category of who brought the case
19	petitionerState	the petitioner's state, where one applies
20	respondent	coded category of who was sued
21	respondentState	the respondent's state, where one applies
22	jurisdiction	how the case reached the Court — appeal, certiorari, and so on
23	adminAction	the federal agency below, if any
24	adminActionState	that agency's state, if any
25	threeJudgeFdc	whether a three-judge district court heard it
26	caseOrigin	the court where the case began
27	caseOriginState	that court's state
28	caseSource	the court the Court took it from
29	caseSourceState	that court's state
30	lcDisagreement	whether the record shows disagreement in the court below
31	certReason	why the Court agreed to hear it — a reason somebody had to read for
32	lcDisposition	what the court below did
33	lcDispositionDirection	whether the decision below was liberal or conservative

Position	Column name	What the codebook says
34	declarationUncon	whether a law was declared unconstitutional
35	caseDisposition	what this Court did with it
36	caseDispositionUnusual	whether the disposition was out of the ordinary
37	partyWinning	whether the petitioner won
38	precedentAlteration	whether the Court altered its own precedent
39	voteUnclear	whether the vote itself was hard to code
40	issue	the specific legal issue, from a coded list
41	issueArea	the broad area that issue falls in
42	decisionDirection	whether the decision was liberal or conservative
43	decisionDirectionDissent	whether the dissent ran the other way
44	authorityDecision1	the basis the Court decided on
45	authorityDecision2	a second basis, where there was one
46	lawType	the kind of law at issue — constitutional, statutory, other
47	lawSupp	the specific provision
48	lawMinor	the provision in the Court's own words
49	majOpinWriter	who wrote the majority opinion
50	majOpinAssigner	who assigned it
51	splitVote	whether the vote was recorded in more than one way
52	majVotes	how many justices were in the majority
53	minVotes	how many were not
54	justice	the justice's numeric ID
55	justiceName	the justice's name
56	vote	how this justice voted, as a code
57	opinion	whether this justice wrote an opinion
58	direction	whether this justice's vote was liberal or conservative
59	majority	whether this justice was in the majority
60	firstAgreement	a justice this one joined
61	secondAgreement	a second justice this one joined

61 columns. Read down them and sort each into two piles. Things that could be checked against the U.S. Reports by anyone with a library card, and things that exist because a person read an opinion and reached a conclusion. The first pile is small — `usCite`, `dateDecision`, `docket`, `majVotes`, `justiceName`. Nearly everything else is in the second: `certReason` (why the Court took the case), `issueArea`, `partyWinning`, `precedentAlteration`, `lcDispositionDirection`, `declarationUncon`, and the column this chapter turns on, `direction`.

Here is one justice-vote, the finest grain the file has, printed one field per line:

Field	Value
term	1946

Field	Value
caseName	HALLIBURTON OIL WELL CEMENTING CO. v. WALKER ...
issueArea	8
decisionDirection	2
majVotes	8
minVotes	1
justiceName	HHBurton
vote	2
opinion	1
direction	1
majority	1

Every value in that row except the name and the date is an integer standing for a word.

issueArea 8 is *economic activity*. *direction* 1 is *conservative* and 2 is *liberal*, and which is which depends on the issue area, by a rule the codebook sets out case type by case type. In a criminal case a vote for the defendant is liberal; in an economic case a vote for the government usually is. *vote* 2 is not “two votes”; it is a code meaning the justice dissented from the opinion of the Court. *majority* 1 means the same thing again, in a different vocabulary.

That last redundancy is the point. The file says the same fact three times in three encodings, and a reader who summed *vote* across a term, or averaged *issueArea*, would get a number with no referent whatever. **None of these columns is a quantity, and every one of them will do arithmetic if you ask.**

Where the codebook has no rule, *direction* is blank rather than guessed. That is the honest choice, and it has a cost. Those cases cannot enter a pairwise agreement count, because “voted the same way” means nothing when there is no way to compare. The build drops them, and that is the only filtering it does.

Here is the clean side — one row per pair of justices, over the 2020 to 2023 terms:

A	B	Agree	Cases	Pct
BMKavanaugh	JGRoberts	225	234	96.2
EKagan	SSotomayor	216	234	92.3
KBJackson	SSotomayor	100	110	90.9

Two names, a count of cases they both voted in, and a count of times they agreed. That is all this chapter’s central figure is built from, and **the unit of analysis has moved twice** to get here: from a justice-vote, to a case, to a *pair of justices*. Nothing in the source file is a pair. The pairs were built by taking every justice against every other, finding the cases in which both cast a directional vote, and counting matches. That is why the *cases* column had to be kept beside *agree*, rather than reduced to a percentage on the spot. A pair who never sat together would otherwise have arrived as a percentage of nothing, and disappeared.

03 Ten people, and everything the database knows about them

Justice	Cases with a coded direction	% of votes coded conservative	% of the time in the majority
Alito	230	70.0	82.0
Thomas	232	67.7	79.3
Barrett	222	61.3	90.9
Kavanaugh	234	60.3	95.5
Roberts	234	59.8	94.3
Gorsuch	232	57.8	81.3
Breyer	122	38.5	72.1
Kagan	234	33.8	74.0
Jackson	110	33.6	78.3
Sotomayor	234	32.1	70.7

10 justices sat between the 2020 and 2023 terms. For each of them the file records how many cases they voted in, what share of their votes the database codes **conservative**, and how often they were in the **majority**.

Two of them, Jackson and Breyer, have about half as many cases as everybody else. One retired during the window and the other replaced him. Nothing follows from that yet; it is worth noticing anyway.

04 Who agrees with whom

The raw material for any ideal-point score is embarrassingly simple: **how often does each pair of people vote the same way?**

Justice	Justice	Voted the same way	Shared cases	% agreement
Kavanaugh	Roberts	225	234	96.2
Kagan	Sotomayor	216	234	92.3
Jackson	Sotomayor	100	110	90.9
Thomas	Sotomayor	123	232	53.0
Alito	Sotomayor	123	230	53.5
Thomas	Breyer	65	121	53.7

Kavanaugh and Roberts agreed 96.2% of the time. At the other end, Thomas and Sotomayor agreed 53.0%, barely better than a coin.

That is the entire dataset: 45 numbers, one per pair. **No ideology has been supplied to it.** Nobody has told this table who is liberal, what any of the cases were about, or even that the people involved are judges. Figure 1 draws all 45 at once.




Figure 1. All 45 agreement rates, in the order the file is in and in the order recovered from them. Every cell is one pair of justices; the darker the cell, the more often the two voted the same way, on an absolute scale running from 53.0% to 96.2%. The red-outlined cell is a pair with no rate at all. The right-hand panel is the identical 45 numbers with the rows and columns re-ordered by the scaling described below, which is not applied to the left-hand panel and changes nothing about the data.

Two things are already visible in that picture and neither was put there. **There is a block structure**, a set of justices who agree with each other far more than with anybody else, and **there is one cell with nothing in it**.

05 The score falls out

Now the move that makes DW-NOMINATE, Martin-Quinn and every other ideal-point score work. Treat agreement as **closeness** and disagreement as **distance**. Then ask: if these 10 people had to be arranged along a single line so that people who agree sit near each other, what line would that be?

Justice	Position
Thomas	-0.218
Alito	-0.217
Gorsuch	-0.132
Barrett	-0.120
Roberts	-0.074
Kavanaugh	-0.073

Justice	Position
Jackson	0.151
Kagan	0.206
Breyer	0.231
Sotomayor	0.246

Read the order. Thomas, Alito, Gorsuch, Barrett, Roberts, Kavanaugh, then a gap, then Jackson, Kagan, Breyer, Sotomayor.

That is the ideological ordering of the Roberts Court as any court reporter in America would give it, and **it came out of 45 agreement rates and nothing else.**

One warning about that column of numbers, and it governs every figure below. The method recovers *distances*, not coordinates. Multiply every position by -1 and every pairwise distance is unchanged, so the fit is unchanged. A scaling solution can be flipped, or in more than one dimension turned, at no cost whatever.

Which end comes out negative is a convention of the arithmetic, not a finding: **the order and the spacing mean something; the direction does not.** The same is true of Martin-Quinn, of NOMINATE, and of every ideal-point estimate anybody will ever quote at you. The orientation is always supplied by the analyst after the fact.

The scaling was handed pairwise agreement rates and nothing else — no case subjects, no vote directions, no party of appointing president. Set its output beside each justice's share of conservative votes, a column it never touched, and the two line up almost exactly. The correlation is -0.975. Its size, 0.975, is the finding; its sign, by the paragraph above, is not.

That is what an ideal-point score *is*. A summary of who agrees with whom, which turns out to line up with something the culture already had a word for.

Quantity	Value
Span of the 6 on one side	0.145
Span of the 4 on the other	0.095
Distance across the gap between them	0.224

And it is not a spectrum. The 6 justices on one side span 0.145; **the single step across the middle is 0.224, larger than that entire bloc.** This is not a line with people spread evenly along it. It is two clusters with a chasm between them.

At this point the score looks like a measurement of nine people. There are 45 agreement rates behind it, a recovered ordering that matches conventional wisdom without having been told it, and a correlation of 0.975 in magnitude with a variable the model never saw. It is worth pausing over how much of that would survive if one of the 45 numbers turned out not to exist.

06 One of the forty-five is not a number

Justice	Justice	Voted the same way	Shared cases	% agreement
Jackson	Breyer	0	0	undefined

Jackson and Breyer never sat on a case together. She replaced him. Their agreement rate is not low and it is not high. It is **undefined**, and no amount of additional data will ever fix it, because the data that would fix it was never generated. That is the red-outlined cell in Figure 1.

The scaling cannot take a hole, so it was given something instead.

Quantity	Value
Pairs with a real agreement rate	44
Pairs with none	1
Value used to fill the gap	0.732
What that value is	the mean agreement across every other pair

That fill is a modeling decision, made here, in one line of code. It is about the most conservative guess available, since it assumes nothing whatever about whether the two would have agreed. It is still a guess, sitting in the middle of the matrix, carrying exactly the same weight in the arithmetic as the 44 real observations around it.

07 Refusing to guess

So refuse. Drop both justices entirely and re-run the scaling on the 8 who overlap completely. No fill, no guess, no cell without data in it.

Before you look at the table, say what it will show. That fill is an invented number sitting in the middle of the matrix the entire ordering came out of. Take it away and how much of the ordering survives? Does it stand exactly as it did? Does it come apart? Or something between — and if between, how many of the 8 justices end up somewhere other than where the filled version put them? Commit to a count, and to one more answer: whichever way the individual ranks fall, do the two blocs survive intact?


Rank	With the gap filled	Refusing to guess	Changed
1	Thomas	Alito	swapped
2	Alito	Thomas	swapped
3	Gorsuch	Gorsuch	
4	Barrett	Barrett	
5	Roberts	Kavanaugh	swapped

Rank	With the gap filled	Refusing to guess	Changed
6	Kavanaugh	Roberts	swapped
7	Kagan	Kagan	
8	Sotomayor	Sotomayor	


2 adjacent pairs change places: Thomas and Alito; and Roberts and Kavanaugh. Nothing else moves at all.

Most readers commit to one of the two extremes — that removing an invented cell changes nothing, or that it brings the whole ordering down. Both are wrong, and being wrong in either direction is what makes the actual answer worth having: 4 positions move out of 8, and every one of the moves is between neighbors.

10 justices, gap filled



8 justices, refusing to guess



the direction of this axis is arbitrary: only order and spacing mean anything

Figure 2. The same Court, scaled with the guess and without it. Recovered positions for all 10 justices with the undefined pair filled by the mean agreement rate, and for the 8 justices whose pairs are all real. The blocs and the chasm are in the same places in both; 4 positions move, in 2 adjacent swaps. The direction of the axis is arbitrary in both panels and carries no meaning.

This is the turn, and it cuts two ways at once.

The recovered order is not a fact about the justices at fine resolution. Adjacent positions moved because of a decision about a cell that contains no data. Anybody reporting which of Thomas and Alito sits further out is reporting a fill choice as much as the Court's behavior.

And the blocs do not move. Six on one side, the rest on the other, in both versions, with the chasm in the same place. The correlation with conservative voting comes out at 0.990 in size, on 8 justices with no guess in the matrix at all. That is slightly stronger than the 0.975 obtained with one. **The coarse structure is in the data. The fine structure is in the assumptions.**

Published scores make dozens of choices of this kind. Martin-Quinn makes distributional assumptions and lets positions drift over time; NOMINATE fixes a scale and a functional

form. **The difference is not that this one has assumptions and those do not. It is that this one is visible, and reversible in a single line.**

08 The number that surprises people more

Justice	% of votes coded conservative	% of the time in the majority
Kavanaugh	60.3	95.5
Roberts	59.8	94.3
Barrett	61.3	90.9
Alito	70.0	82.0
Gorsuch	57.8	81.3
Thomas	67.7	79.3
Jackson	33.6	78.3
Kagan	33.8	74.0
Breyer	38.5	72.1
Sotomayor	32.1	70.7

Kavanaugh is in the majority 95.5% of the time, and the lowest figure is Sotomayor at 70.7%.

Set that beside the recovered line. The two justices nearest the middle of it win almost every case; the justices at either end lose most often. **"Moderate" here is not a description of anybody's opinions. It is a description of how often you are on the winning side**, and it is a mechanical consequence of where you sit.

09 Does the middle still win when it is close?

Justice	% in majority, all cases	In majority, of the 30 one-vote cases	% in majority, close cases
Kavanaugh	95.5	22	73.3
Roberts	94.3	22	73.3
Sotomayor	70.7	17	56.7
Kagan	74.0	16	53.3
Barrett	90.9	14	46.7
Alito	82.0	14	46.7
Thomas	79.3	13	43.3
Gorsuch	81.3	13	43.3

There were 30 decisions in this window with a one-vote margin. The table shows the 8 justices who sat for the whole run of them. The 2 who sat for part of the window cannot

be given a comparable denominator, meaning a comparable bottom number to divide by. That is the missing-cell problem from earlier, turning up somewhere else entirely.

The answer is yes, and much less so. Kavanaugh and Roberts are still at the top, but at about 73% rather than 96%. And the ordering below them scrambles.




Figure 3. How often each justice wins, against where the scaling puts them. Share of the time in the majority across all cases (black) and across the 30 one-vote cases alone (orange), plotted against recovered position. The black points form an arch: the closer to the middle of the line, the more often you are on the winning side, peaking at Kavanaugh. In orange the arch flattens by 22.2 points at its peak and partly turns over. Hollow points are the 2 justices who sat for only part of the window and have no comparable close-case denominator.

The 95.5% figure is mostly a fact about unanimous and lopsided cases, which are most cases. Restrict attention to the decisions that were actually contested and the Court looks a great deal less like a single pivotal justice deciding everything. The highest close-case rate is 22.2 points below the highest all-case rate, and the order below the top two turns over. Roberts and Kavanaugh keep the lead. But Sotomayor and Kagan, at the far end of the line, are in the majority in more of the close cases than any of the 4 justices between them and the middle.

10 Seventy-eight terms, and what they are made of

The same database reaches back to 1946, so the long series comes free with the file.

Quantity	Value
Terms in the series	78
First	1946
Last	2023

Quantity	Value
Highest share of decisions coded conservative	65.1% (2021)
Lowest	21.8% (1963)




Figure 4. Share of decisions coded conservative, every term from 1946 to 2023. The gray line is the term-by-term figure and the red line a five-term moving average; the dotted rule is half. The series runs from a low of 21.8% in 1963 to a high of 65.1% in 2021, and every point in it is a coder's judgment about the direction of a decision.

Before believing any story in that line, notice exactly what is plotted. It is the share of decisions that the database's coders labeled conservative: a human judgment, applied case by case, for 78 years, by people following a codebook. It is a good codebook, and a case coded conservative in 1962 may have almost nothing in common with one coded conservative in 2022.

That is not a reason to discard the series. It is a reason to say what it is. **Every long-run measure in social science is somebody's coding decision repeated for decades**, and the ones worth using say so.

11 The argument about this

On circularity. The strongest defense of what was done above is that the ordering was not put in. The scaling was given 44 numbers and no labels, and the dimension it found happens to be the one the culture already names. That is a genuine finding about the structure of judicial voting. The strongest objection is that agreement and ideology may simply be the same phenomenon observed twice, in which case a correlation of 0.975 is a tautology with extra steps. Both positions are defensible, and neither is refutable from this table.

On what a score can be used for. Martin and Quinn (2002) built the standard version precisely because a static score cannot show drift, which is the question people most want answered about long-serving justices. Segal and Cover (1989) went the other way and measured ideology *before* any votes existed, from newspaper editorials, so that votes

could be predicted rather than summarized. **The three designs answer three different questions**, and treating them as competing estimates of one quantity is a category error.

On the majority rates. The arch in Figure 3 looks like a finding about power, and the orange points show it is substantially a finding about which cases exist. Whether “the middle wins” is a claim about the Court or about its docket is not settled here — and the docket is chosen by the same nine people.

On whether “conservative” is constant. The database codes each decision’s direction case by case, following a codebook, for 78 years. It is a human judgment repeated across an enormous span of American legal history, and Figure 4 is only as comparable across its two ends as that codebook is.

On agreement itself. Two justices can reach the same result from incompatible doctrines. Agreement is not agreement about reasons, and the reasons are most of what law professors argue about.

12 What this evidence will bear

The Supreme Court Database records every vote in every case since 1946, coded consistently by a documented codebook and maintained as a versioned academic resource. The window scaled here is a complete census of the Court’s decisions rather than a sample. And the method is transparent enough to be broken on purpose, which is what the second half of this chapter does to it.

The limits bound different claims, and each has already appeared above in its proper place. **One cell has no data in it**, and the fill chosen for it moves 4 positions in the recovered order, so no fine ranking here is safe. **Two justices sat for part of the window**, so their statistics rest on different denominators from everybody else’s, which is why they are hollow in Figure 3 and absent from the close-case table. **One dimension was chosen** and found to fit well, which is not the same as discovering that judicial disagreement is one-dimensional. **“Conservative” is a coder’s judgment** applied for 78 years, and **agreement is not agreement about reasons**.

And the scaling is **a snapshot**, in which nobody can move. It cannot show a justice drifting, which is exactly what Martin-Quinn was designed to capture.

Three questions stay genuinely open. The recovered dimension may be ideology, or it may be something else that moves with it: deference to precedent, institutional caution, the identity of the litigants. Agreement rates cannot settle that. Whether coding direction is stable across seventy-eight years is a question about the codebook rather than about the Court. And whether the Court’s high agreement rates reflect consensus or docket selection is unanswerable here, because the justices choose which cases to hear, which makes the denominator endogenous to the thing being measured.

What can be said is worth the trouble it took. A left-right ordering of the Supreme Court can be recovered from 44 agreement rates that carry no information about ideology, cases or appointing presidents. It lines up with a column the method never saw, at a size of 0.975. The blocs in it are in the data, and survive refusing to guess about the one pair who never met. The fine ordering is in the assumptions, and does not.

Every score has a fill. Somewhere inside every published ideal-point estimate is a cell with no data in it and a decision about what to put there. The question to ask of any such

number is not whether it made assumptions, because all of them do. It is **which ones, and what happens when you change them.**

13 Sources

Data

- **The Supreme Court Database**, Washington University in St. Louis, 2024 Release 01, justice-centered by citation — every justice's vote in every case since 1946. <http://scdb.wustl.edu/data.php> (the site serves over `http` only). Agreement and scaling use the 2020–2023 terms; cases where the database records no ideological direction are excluded, because “voted the same way” is undefined without a direction to compare.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`agreement.csv`, `by_term.csv`, `close_cases.csv`, `justices.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `agreement.csv` gives every pair of justices a `pct` agreement. Sort it. Who agrees with whom, and do the blocs match the usual labels?
- `justices.csv` has `pct_conservative` and `pct_in_majority`. Plot one against the other. Who is most often on the winning side, and is that the same as being in the middle?
- `close_cases.csv` lists the divided decisions with the `majority_bloc` named. Count how often each justice appears. Who actually decides the close ones?
- `by_term.csv` runs `pct_conservative` back decades. Follow it. Does the Court move gradually, or in steps when its membership changes?

Academic literature

- Martin, A. D. and Quinn, K. M. (2002). “Dynamic Ideal Point Estimation via Markov Chain Monte Carlo for the U.S. Supreme Court, 1953–1999.” *Political Analysis* 10(2): 134–153.
- Segal, J. A. and Cover, A. D. (1989). “Ideological Values and the Votes of U.S. Supreme Court Justices.” *American Political Science Review* 83(2): 557–565.
- Epstein, L., Landes, W. M. and Posner, R. A. (2013). *The Behavior of Federal Judges: A Theoretical and Empirical Study of Rational Choice*. Harvard University Press.

Further reading

- Poole, K. T. and Rosenthal, H. (1997). *Congress: A Political-Economic History of Roll Call Voting*. Oxford University Press. (The same scaling idea, applied to legislatures.)

REBUILD THIS DATA YOURSELF, WITH AN AI ASSISTANT

I want to rebuild a public dataset from its original source, without any

starting files. Please do the following and show your work.

THE SOURCE. The Supreme Court Database, maintained at Washington University in St. Louis – every case the Court has decided since 1946, coded by researchers from the written opinions, because the Court itself publishes opinions and no data. Release 2024_01, the justice-centered file organised by citation:

http://scdb.wustl.edu/_brickFiles/2024_01/SCDB_2024_01_justiceCentered_Citation.csv.zip

About 1.7 MB zipped. Each row is one justice's vote in one case – more than 80,000 rows. Two coded columns matter most: direction (1 conservative, 2 liberal, blank where the coders could not say) and majority (1 dissent, 2 majority).

HOW TO FETCH IT. The site serves over plain http only; an https address fails outright. That is the source, not a problem with your tooling – say so and proceed. The file is UTF-8; read it as such, or the curly apostrophes in case names silently corrupt.

WHAT TO BUILD.

1. Per term, 1946 through 2023: how many cases were decided and what share went in the conservative direction.
2. For the 2020 through 2023 terms: every pair of justices and how often the two voted the same direction, with the number of cases they shared – and, per justice, how often each was in the majority.

CHECK YOUR WORK. The copy this book used was fetched in August 2026.

If your rebuild matches it, you should find:

- 78 terms, 1946 through 2023
- ten justices sat in the 2020–2023 window, making 45 pairs; the busiest of them each heard 234 cases there
- Kavanaugh is in the majority 95.5% of the time and Roberts 94.3%, the two highest; Sotomayor, at 70.7%, is lowest

A new release lands roughly yearly at a new address; this one stays put, so a difference means you fetched a different release, not that you made an error.