

Section III — Administrative Data

Records kept to run a system, about people who never asked to be described

Democracy's Data

Last updated September 13, 2026 at 10:53a EDT

This section is about records that were made while their makers were doing something else. A county officer certified an election total — declared it final and lawful, over a signature — because the law required a total. A campaign treasurer filed what the campaign raised because a statute set a deadline. A clerk recorded a roll-call vote, each member's name called and each answer written down, so the House would know whether the bill passed. A police officer wrote down a traffic stop to account for a shift.

None of them was measuring anything. Each was running a system, and the file is the paperwork the system left behind.

The family is large, and this is the largest section of the book. It holds certified election returns and the rosters of who actually held each seat. It holds voter registration files, every roll-call vote since 1789,¹ campaign-finance filings and lobbying registrations. It holds records of attention, like a year of Wikipedia pageviews. And it holds records of ordinary people: the first names written on Social Security cards, traffic stops, jury strikes (the prospective jurors each side removed from a panel without having to give a reason), and the mortgage-risk grades a federal agency stamped on neighbourhoods in the 1930s.²

One thing unites all of it. Nobody who appears in these files answered a questionnaire to get there, and nobody who wrote them was trying to describe anyone. That single fact is the source of everything this kind of data can do, and of everything it cannot.

¹ Kept, vote by vote, by the Voteview project at UCLA: <https://voteview.com/>.

² The Home Owners' Loan Corporation's maps, digitised by the University of Richmond's Mapping Inequality project: <https://dsl.richmond.edu/panorama/redlining/>.

01 What a record can do that a count or a survey cannot

Sections I and II both produced numbers by trying to describe people. The census counts everyone and misses some. A survey asks a sample and then weights it, counting some answers more than others, until the sample looks like the country. Both operations end in an estimate, and an estimate always carries a gap between the number and the world.

An administrative record involves neither operation. Nothing was sampled and nothing was weighted, so a certified total carries no margin of error — no published statement of how far off it could be. **The record is not an estimate of the event. The record is the event, written down as it happened.**

The purest case is the roll call. When a member of Congress votes, the vote is the act, not a report of one. Nobody self-reported it, nobody misremembered it, and there was no

filing deadline to miss. That makes it about the cleanest data in this book, and it is worth meeting a clean source early so you have something to compare the messy ones against.

But the advantage has edges, and two of them matter enough to state now. First, no record holds an opinion. A voter file can say you voted; only a survey can say why, which is what Section II was for.

Second, some of these records are claims their subjects made about themselves. The Federal Election Commission checks a filing for completeness and audits a small number of committees, but it does not verify that a campaign spent money the way the campaign says it did. It publishes what was filed.³ A disclosure is excellent evidence about what somebody was willing to say on a form with their name on it, and weaker evidence about what happened.

³ The Commission's audit reports, and the narrow grounds on which an audit is opened — publicly funded presidential committees, or a vote of the commissioners: <https://www.fec.gov/legal-resources/enforcement/audit-reports/>.

02 Made while doing something else

Every record here shares one property, and it is the single most important thing to remember here: **the record was made for another purpose, and its columns do that job and no other.** A voter file exists to hand each person the correct ballot, and the voter-file chapter finds most of its columns doing exactly that and nothing else. A stop record exists so an officer can account for an enforcement action. A lending grade existed so a government corporation could decide where to insure mortgages. The description of a person is a by-product, and the by-product is what you are reading.

The last cluster of this section reads three such records. Here is what each was for:

Record	Written by	In order to	One row is
Traffic stops	A police officer, at the roadside	Record an enforcement action	One stop
Peremptory strikes	A court clerk, during jury selection	Record who was seated and who was removed	One juror in one trial
HOLC security maps	A federal agency, grading mortgage risk	Tell lenders which neighbourhoods were safe to lend in	One graded area

Now read the last column — the one this whole section keeps returning to:

Record	Silent about
Traffic stops	Everyone who was driving and was not stopped
Peremptory strikes	Why a juror was struck
HOLC security maps	Everyone the grade was applied to

Each file counts the events an institution acted on. None of them counts the population those events were drawn from, because counting that population was nobody's job. Nobody at the roadside was recording the drivers who went past.

So an administrative record usually arrives as a **numerator**: the top number of a fraction, with nothing underneath it. Any claim about a *rate* needs a denominator, the bottom number you divide by. That bottom number is not in the file and has to be brought in from somewhere else, usually from Section I. **Choosing it is not a preliminary to the analysis. It is the analysis.**

03 The shape of this section

The section runs in seven clusters. Each is led by a chapter about one kind of record: what it is, who writes it and under what obligation, and what it can bear. The briefs that follow put that record to work.

Election returns. The result: who won, by how much, certified by the county boards and state officers whose job is to sign it. It is the most authoritative data in this book, because somebody signed their name to it, and it is still not a dataset. No federal agency collects American election results. Every national file you have ever seen was assembled privately from fifty-one separate state records (fifty states and the District of Columbia), and several briefs here are about what happens at the seams. The returns chapter says what a certified return can and cannot show, and its companion says which rung of the ladder answers which question: state, county, precinct, ballot.

Rosters of officeholders. A roster answers a different question than a return does: who held the seat, not who won the election, and the two records disagree more often than you would guess. The roster chapter measures that gap, and the cluster's briefs use member lists running back two centuries.

Voter files. The machinery behind the result: who was registered, in which party, and whether they voted. The voter-file chapter shows the one thing no other instrument can do — check whether a survey is lying about turnout. The catch is access. The file is public by law in every state and downloadable in almost none of them; you may be asked to sign something, pay something, or appear in person. *Public* and *obtainable* are different words, and the gap between them is where much of what we do not know about American elections lives.

Roll calls and courts. The clean records: every congressional vote since 1789, and every Supreme Court vote beside it. The roll-call chapter also shows what the record omits: no note of what most votes were about. The briefs build ideology scores on top of it. What those scores are, and whether to trust a manufactured number, is Section IV's opening question.

Campaign finance. The self-reported records: filings to the FEC and the lobbying registry, the only source in this book written by the people it is about. Disclosure is not audit, so this cluster is simultaneously the most detailed data in the book and the least independently checked.

Media and attention. Records kept by platforms while serving pages: pageview logs and follower counts. They are administrative too: the ledger of a server, not a measurement of opinion. The briefs test how far a country looking something up can stand in for a country caring.

Records of ordinary people. The closing cluster, and the one this page's own tables come from: traffic stops, jury strikes, redlining maps, and names. Everyone earlier in the section put themselves forward — ran for office, raised money, registered to lobby. Nobody in these last files volunteered. The person in a stop record was driving to work, and the people graded on a 1937 map were not consulted and mostly never learned that a federal agency had assigned their neighbourhood a letter.

04 The denominator that cannot exist

The rest of this page works the numerator problem on the section's hardest case, the traffic stops, because the problem is easiest to see where it cannot be escaped. The comparison everybody reaches for is stops per resident. Here it is, computed three ways, from the same 905,070 San Francisco stops and three different population counts:

Denominator	Black stops per person	White stops per person	Disparity ratio
residents	3.34	1.09	3.07x
adults 18+	3.86	1.20	3.23x
drives to work	19.88	4.15	4.79x

The disparity — stops per Black person divided by stops per White person — runs from **3.07x** to **4.79x**, a spread of 1.56x, depending on nothing but which published denominator you divide by. All three are Census Bureau products, and all three are defensible in a sentence. Residents is the obvious choice. Adults 18 and over is more defensible, since children do not drive. People who drive to work is the most defensible of all, and it is still wrong, because it misses everyone driving for any other reason.

And none of them is the number you actually want, which is **who was on that road, at that hour, driving in a way that could attract a stop**. No agency collects it. There is no survey of it. It is not a hard number to find; it is a number that does not exist.

This is the problem in its purest form. The file is real, the stops are real, and the disparity is real. And the most natural sentence you could write about it is a sentence about a denominator you chose: "Black drivers are stopped at 3.07 times the rate of White drivers."

05 The move that needs no denominator

There is a way out, and it is worth seeing because it generalises to the whole section.

Ask a different question. Not *who gets stopped*, which needs a population, but *what happened to the people the officer already decided to search*. Both of those are events the institution recorded. The population never enters.

Race	Searched	Search rate	Contraband found	Hit rate
asian/pacific islander	2,861	1.8%	1,033	36.1%
white	11,707	3.1%	2,834	24.2%
other	3,746	3.5%	761	20.3%
hispanic	11,445	9.9%	1,165	10.2%
black	23,622	15.5%	2,183	9.2%

Black drivers are searched at **15.5%** against **3.1%** for White drivers — a rate over stops, and stops are in the file, so that comparison is sound. Then the turn, the **hit rate**, the share of searches that found anything: **searches of White drivers turn up contraband 24.2% of the time, and searches of Black drivers 9.2%.**

A search is a decision made under a standard of suspicion. If the same standard were being applied, the searches would succeed at about the same rate. Searching a group far more often and finding contraband far less often is evidence about the threshold, and **it required no denominator at all** — only two columns the institution had already written down.

06 The denominator that is already in the file

The jury records have the same problem and a luckier solution. A prosecutor may strike a juror for almost any reason except race, and the reason is rarely written down. So the numerator, jurors struck, is recorded, and the thing you would want to divide by is not. What would a race-neutral strike rate even look like?

The file answers that itself, because the state is not the only party striking jurors:

Comparison	Rate
State strikes Black jurors	49.8%
State strikes White jurors	11.2%
Defense strikes White jurors	47.3%
Defense strikes Black jurors	15.0%

The state — the prosecution — struck **49.8%** of eligible Black jurors and **11.2%** of eligible White ones. On its own that is a ratio without a baseline. But the defense, drawing from the very same panels in the very same trials, struck **47.3%** of eligible White jurors and **15.0%** of Black ones — very nearly the mirror image. **The control group — the comparison that shows what the same choice looks like made by the other side — was in the file the whole time.** Two parties, the same pools, opposite incentives; together they are close to an experiment the court ran without intending to.

07 The denominator that is a choice of unit

The redlining maps look like the easy case, and they are the subtlest. Here the population *is* available — the census counts everyone in a graded area. The question is not what to divide by but **what to compare with what**.

Quantity	Value
Cities	33
Pooled A-graded, % Black	16.7
Pooled D-graded, % Black	25.3
Pooled gap, points	8.6
Median within-city gap, points	11.4
Cities where the gap reverses	6

The maps graded neighbourhoods from A, the safest to lend in, to D, the riskiest. Pool every graded area in the country and D-graded areas come out **8.6** points more Black than A-graded ones. Hold the city fixed, compute the gap inside each of the 33 cities, and the median — the middle city, half above it and half below — is **11.4** points. Pooling does not exaggerate here. **It understates**, because a national pool compares Cleveland's D areas against Berkeley's A areas, and those cities differ for a hundred reasons that have nothing to do with a 1930s grade. The comparison the maps invite is *within* a city, because that is the level at which the grade was applied.

The pooled figure also conceals the most interesting thing in the file: **6 of the 33 cities reverse**, with A-graded areas more Black than D-graded ones. A single national number cannot report a reversal. It can only average it away.

08 What this data cannot tell you

It cannot tell you what anyone believed. A record holds acts: a ballot cast, a dollar filed, a stop made. If your question is about opinion or motive, you are back in Section II, asking.

It cannot tell you that any of these institutions discriminated. Every finding above is about rates and thresholds. A hit-rate gap is evidence about the standard applied to a search; it is not a finding about any officer, and no row in any of these files records an intention.

It cannot supply the missing denominators. The driving population does not exist as data, and nothing in this section conjures it. What the section offers instead is the alternative: change the question to one the file can answer, or say plainly which denominator you chose and why.

And much of it cannot be checked against a truth. A self-reported filing has no audit behind it, and there is no register of stops-that-should-have-happened. Where a check does exist, the chapters make the most of it: a voter file against a survey, a recount against a count.

09 What you should have learned

- **An administrative record is the event, not an estimate of it.** Nobody sampled and nobody projected: a ballot was cast, a form was filed, a stop was made. That is the advantage this section has over both counting and asking, and it is why these files can answer questions no survey can reach.
- **The purpose the record was kept for decides what it can answer.** Every column does the job the system needed, and no other. A file that is excellent for running an election is poor for describing an electorate, and neither fact is a defect.
- **The numerator usually arrives without its denominator.** Stops, strikes and filings are counted; the population at risk of being stopped, struck or filed about is often counted by nobody. Divide the same San Francisco stops by three defensible denominators and the disparity ratio moves by a factor of 1.56.
- **Sometimes the comparison the file already contains is the better question.** Search rates need a driving population that does not exist; hit rates do not, because the denominator — searches actually conducted — is in the file. 9.2% of searches of Black

drivers found contraband against 24.2% for white drivers, and that comparison needs nothing from outside.

- **Pooling can hide the finding, not just blur it.** The national redlining gap of 8.6 points is smaller than the median within-city gap of 11.4, and a single national number cannot report that 6 of 33 cities reverse.

10 How to read this section

The clusters are independent, and inside each one the data-type chapter comes first: read it before the briefs that lean on it. If you only have time for two chapters, read the roll-call chapter for the clean case and the campaign-finance chapter for the self-reported one. The contrast between a record that is the event and a record that is a claim about the event is most of this section's argument.

11 Sources

Every figure on this page is read from tables three of the section's own briefs built, so nothing here was collected afresh, and re-running any of those briefs re-runs this argument:

- `policing-by_race.csv`, `acs_denominators.csv`. Stop, search and contraband counts, with three American Community Survey denominators for the same population.
- `jury-selection-strikes.csv`, `trials.csv`. Peremptory strikes by side and race, over 211 trials.
- `redlining-cities.csv`. A- and D-graded population and racial composition, city by city.

The three briefs carry the full account of what each archive is and what it cost to get. The archives themselves are the Stanford Open Policing Project (<https://openpolicing.stanford.edu/data/>), the jury data APM Reports released with *In the Dark* season 2 (<https://github.com/APM-Reports/jury-data>) and *Mapping Inequality* at the Digital Scholarship Lab, University of Richmond (<https://dsl.richmond.edu/panorama/redlining/>). The other records this page describes are named, dated and cited in the chapters that lead their clusters: returns, rosters, voter files, roll calls, filings, pageviews.

Every figure above is recomputed from those tables and checked against the values in `checks.csv`, printed here in full.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`control.csv`, `denominators.csv`, `hit_rates.csv`, `records.csv`, `unit.csv`

Check	Value
Records described	3
Chapters read, none of them downloaded	3

Check	Value
Stops in the policing file	905,070
Published denominators available for the same stops	3
Smallest disparity ratio, and its denominator	3.07x (residents)
Largest disparity ratio, and its denominator	4.79x (drives to work)
Spread between them, as a multiple of the smallest	1.56x
Hit rate, White drivers searched, %	24.2
Hit rate, Black drivers searched, %	9.2
Jurors the state found eligible, Black	1,811
Jurors the state found eligible, White	3,318
Trials in the jury file	211
Strike rates recomputed from struck/eligible, max disagreement with printed, pts	0.01
Cities in the redlining file	33
Pooled A-to-D gap, points	8.6
Median within-city A-to-D gap, points	11.4
Cities in which the gap reverses	6