

Peremptory Strikes and Jury Composition

Peremptory strikes in one Mississippi district, and a pattern no single trial can show

84-355 Democracy's Data

August 2026

In a courtroom in Mississippi's Fifth Circuit District in 1992, a prosecutor stood up and removed six people from a jury panel. Of the people sitting in that box, 15 were his to remove if he wanted them gone: 3 Black and 12 white. He struck all 3 of the Black jurors and 3 of the 12 white ones.

He gave no reason for any of the six, and no reason was asked of him. That is what a **peremptory challenge** is. Each side in an American trial gets a fixed number of them, and the whole point of the device is that it does not require justification. The trial is the first row of the file this chapter is built from, and the file records it in four numbers:

Trial	Black jurors eligible	Struck by the state	White jurors eligible	Struck by the state
1992-2061-Robert Triplett	3	3	12	3

Two words in that table are doing more work than they look. **"Eligible" means the side could actually have struck that person.** A juror already removed for cause, or never reached before the panel filled, was available to nobody. Counting such people in the denominator, the bottom number each rate is divided by, would understate every rate in this chapter. The source records this juror by juror, and these counts respect it.

And **race is missing for a third of the jurors in the wider file.** APM Reports coded race from court records and, where those were silent, from other sources; 4,752 of 14,873 prospective jurors remain "Unknown" and are excluded here. That hole does not close later on. It should be carried through everything that follows.

01 One prohibited reason, in a system built on not requiring reasons

For most of American history the peremptory strike was absolute. Since *Batson v. Kentucky* (1986) there is exactly one forbidden reason: **race**. If the defense makes a *prima facie* showing of racial motive, the prosecutor must supply a race-neutral explanation for the strike, and the judge decides whether to believe it.

That is the entire regime, and its weakness is visible in the sentence that states it. **A rule against one reason, inside a procedure designed not to ask for reasons, can only be enforced by inference.** Nobody writes down "because she is Black." What exists instead is a

record of who was removed. Whether that record can carry the inference is the question this chapter is about.

It is worth being clear why the answer matters beyond the courtroom. The jury is the only branch of American government staffed by lottery. It exists so that the state's most severe power, taking a person's liberty or life, is exercised by ordinary people rather than by professionals. Its legitimacy comes entirely from being drawn out of the community rather than chosen.

A jury assembled by removing one group is not a lottery. It is a jury selected by a lawyer, returning a verdict rendered by people the community did not supply. The harm runs three ways at once. To the defendant tried by a filtered panel. To the excluded juror denied a civic duty on the basis of race. And to the community whose judgment was quietly replaced with somebody's preference.

02 The record the courts never kept, and the newsroom that did

Enforcement by inference requires a record, and no court keeps one. A clerk logs that a juror was excused; nothing in the file the clerk keeps says what race that juror was, because no rule requires it to. The evidence in this chapter therefore exists only because a newsroom built it. APM Reports, the investigative unit of American Public Media, decided during the second season of *In the Dark* that the question could not be answered without it.

Reporters pulled the paper records of 225 trials out of Mississippi's Fifth Circuit, covering twenty-six years, and coded every prospective juror by hand. Who was in the pool, who could be struck by whom, who was struck, and, where the documents allowed it, what race the juror was. They then published the whole thing openly, as three CSV files, which is why a chapter can rest on it at all.

Two things follow from how it was made, and neither is visible in the files. Nothing here was queried, downloaded from an agency, or scraped off a website, because there was nothing to query. The record did not exist until reporters made it, one trial folder at a time, reading paper and writing down what they read.

Getting the result now takes a second from a public repository, and that second is the entire cost to everyone who came afterward. It is as clean an illustration as this book has of how completely a finished CSV hides what it took to produce.

The other consequence is that it stops. A hand-built record covers what somebody had the reason and the budget to build, and there is no second circuit, no later years, and no process by which either would appear. So when a later section says this district cannot speak for the country, that is less a limitation of the sample than a limitation of what a newsroom can afford to read.

Two properties make it the right evidence rather than merely the available evidence. It is juror-level, so a rate can be computed over the people a given side was actually permitted to remove instead of over everyone in the room. And it records both sides' strikes, which supplies a control group nobody had to construct: the defense drew from the same panels, in the same trials, under the same rules, wanting the opposite thing. A file of state strikes alone would leave every disparity open to the reply that the struck jurors differed somehow. This one answers that reply from inside itself.

The thing it does not contain is the thing the law cares about. Reasons are absent from every row, because the procedure was designed not to require them. That gap is a feature of the institution rather than a failure of the reporting, and no better dataset would close it.

The other limits are real and unfixable in a different way. Race here is a reporter's coding of a document, not anyone's self-report, and the documents were often silent. The unknowns are excluded rather than guessed at, and they are unlikely to be a random third. And this district was read precisely because it produced *Flowers v. Mississippi*, which is the strongest possible reason to expect it to be unrepresentative.

The reduction from that to what this chapter reads is short and lossy on purpose. The juror file is walked once. It keeps only the jurors whose recorded eligibility says the side in question could have struck them, and only the two coded races. The result collapses to 4 rows of strike rates and 211 trial rows, one for every trial with at least ten eligible jurors. 5,033 juror-slots survive into the trial table. Nothing juror-level is stored here, so a reader who wanted to recount would go back to the published file and rerun that script, which does no imputation of any kind.

03 What a hand-built record looks like

APM's `jurors.csv` has 14,874 rows and nine columns, and one row is one prospective juror in one trial. Three of them, in seven of the nine columns, each row wrapped onto a second line to fit the page:

Column	What it holds	Juror 35	Juror 38	Juror 8198
id	the reporters' identifier for this prospective juror	35	38	8198
trial	the case, as docket number and defendant	1993-9826-Terry L. Landingham	1993-9826-Terry L. Landingham	2006-0044-Darreal Forrest
race	the juror's race, as recorded	White	Black	White
gender	the juror's sex, as recorded	Male	Female	Female
race_source	how the reporters knew the race — the column almost nothing else in this book has	Jury strike sheet	Jury strike sheet	Jury strike sheet
struck_by	who removed the juror, or that nobody did	Struck for cause	Struck for cause	Struck by the defense
strike_eligibility	which side was entitled to strike this juror at all	n/a	n/a	Both State and Defense

Look at `race_source`. **The file records how it knows.** Race was read off a jury strike sheet for most jurors, off a voir dire transcript for a couple of hundred, and guessed from other sources for a few dozen. For 4,750 of them the value is `No info`.

Almost nothing in this book carries a column like that. Where a value came from is normally something you have to reconstruct from a build script, or guess at. Here the reporters wrote it into the record beside the value, which is what makes the gaps in this file arguable rather than merely present.

Now look at the last two columns, which are where the chapter's arithmetic actually lives. Neither is the yes-or-no it needs to be.

Column	Distinct values	The values
<code>struck_by</code>	9	"Struck by the state", "Struck by the defense", "Struck for cause", "Struck without notation", "Juror excused/absent", "Unknown", and three more
<code>strike_eligibility</code>	5	"Both State and Defense", "State", "Defense", "Neither", "n/a"

Out of that comes this:

Side	Race	Could be struck	Struck	%
state	Black	1,811	902	49.8
state	White	3,318	372	11.2
defense	Black	846	127	15.0
defense	White	2,812	1,330	47.3

4 rows. Fourteen thousand hand-coded jurors, each with a note saying where the record came from, become four numbers and their rates.

The three decisions that make the rate mean anything

The denominator is eligibility, not attendance. `strike_eligibility` records who *could* have removed each juror, and nearly half the file reads `n/a` — jurors already gone for cause, or never reached. Counting those in the denominator would make both sides look far more restrained than they were, and would do it to the state and the defense unequally. **This is the single decision the whole comparison rests on**, and it is available only because the reporters thought to record eligibility as its own field. Most court records do not.

Nine outcomes were collapsed to one question. `struck_by` distinguishes a juror struck by the state from one struck for cause, one excused, one never noted, and one whose fate is simply unknown. The rate in the table counts only `Struck by the state` or `Struck by the defense` as a strike; `Struck without notation`, 571 jurors, counts as not struck by either side. That is the conservative choice and it is still a choice, and it pushes both rates down.

Two races were kept and 4,756 jurors were dropped. Nearly all of them are the No info cases, and they are 31.9% of the file. They are unlikely to be a random third, since a juror whose race the paperwork never recorded is not a random juror. Four more were dropped because they were coded Asian or Latino: too few to compute anything about, and worth naming rather than folding silently into a two-category world. **A binary appears in the output that was not in the input.**

04 The first comparison anybody makes

The file behind the 1992 trial holds 211 of them, every trial with at least ten eligible jurors prosecuted by a single district attorney's office over twenty-six years.

Quantity	Value
Trials in this file	211
Prospective jurors eligible to be struck	5,033
Struck by the state	1,205
Prosecutor's office	Mississippi's Fifth Circuit District
Period	26 years

That narrowness is deliberate. A national average would tell you something about American prosecution in general; this tells you about a decision-maker, the office led by District Attorney Doug Evans, across a career.

Pool every strike that office made and sort by the juror's race:

Race	Eligible to be struck	Struck by the state	% struck
Black	1,811	902	49.8
White	3,318	372	11.2

The state struck 49.8% of the eligible Black jurors and 11.2% of the eligible white ones, a ratio of 4.4 to one.

05 The reply, which is not frivolous

It is tempting to stop there, and stopping there would be a mistake, because there is an answer to that table and it is a serious one:

Struck jurors differ in ways that have nothing to do with race. Some had relatives who had been prosecuted. Some expressed doubts about the death penalty. Some knew a witness. If those characteristics are unevenly distributed across the community — and in a county with this history they certainly are — then a race-neutral striking rule will produce a racial disparity all by itself.

The table above cannot defeat that argument, because it has no control group. It compares Black jurors to white jurors and finds them treated differently, which is exactly what the prosecutor's story predicts. What would settle the matter is another party striking from the *same pool*, in the *same trials*, under the *same rules*, for *different reasons*.

06 The other side of the same room

There is one, and it has been sitting in the file all along. **The defense also strikes.** Same venire, same courtroom, same day, same human beings, opposite incentives.

The logic is worth working out before looking. If Black jurors really were disproportionately strike-worthy for reasons any lawyer would recognize, the defense should strike them too, or at least not strike them dramatically *less*. Suppose instead the two sides turn out to mirror each other. Then whatever shared trait explains the state's behaviour would have to make a juror desirable to the defense exactly as fast as it makes them undesirable to the state.

Striking side	Juror's race	Eligible	Struck	% struck
state	Black	1,811	902	49.8
state	White	3,318	372	11.2
defense	Black	846	127	15.0
defense	White	2,812	1,330	47.3




Figure 1. The two sides strike the same panels in opposite directions. Share of eligible jurors struck, by the striking side and the juror's race, across 211 trials. The state struck 49.8% of eligible Black jurors and 11.2% of eligible white ones; the defense struck 47.3% of eligible white jurors and 15.0% of eligible Black ones. The two tall bars differ by 2.5 points and the two short ones by 3.8.

The pattern reverses almost exactly, and that is worth more than the disparity that opened this section. Whatever made a Black juror strike-worthy to the state made a white juror strike-worthy to the defense at nearly the same rate. The legitimate-characteristics explanation now has to name a trait that is unevenly distributed by race and valued in precisely opposite directions by the two parties. No such trait has been proposed, in this litigation or in the research literature it belongs to.

Which ought to settle it. It does not, and the reason is not statistical.

07 The law asks the question one trial at a time

Batson is litigated in front of a judge who has seen one venire. Aggregate patterns across a career are not before that court and, in the ordinary case, have never been assembled by anybody. The question the judge is actually asked is whether *this* set of strikes, in *this* room, on *this* morning, is explicable without race.

So put the 1992 trial back on its own terms and answer that question honestly. The prosecution had 15 eligible jurors, 3 Black and 12 white, and used 6 strikes. If those 6 had been chosen without regard to race, how often would all 3 Black jurors be among them?

That is an urn problem: draw 6 from 15 without replacement and ask for the chance that the draw swallows all 3 of the 3.

the venire: 15 eligible jurors

what the prosecution actually struck




Figure 2. One trial, and every outcome race-neutral striking could have produced. The 15 eligible jurors of the 1992 trial, and the 6 strikes the prosecution used. The spread shows how many Black jurors those strikes would have caught, had they been drawn without regard to race. The observed outcome, all 3 of 3, is the least likely of the 4 possibilities, occurring in 4.4% of draws.

Quantity	Value
Eligible jurors	15
Black among them	3
Strikes used	6

Quantity	Value
Black jurors struck	3
Chance of a split this lopsided or worse	4.4%

About one chance in 23. Suggestive, and nowhere near enough. A judge would be right to observe that unlikely things happen in single trials all the time, and that 211 trials passed through this district. A 4.4% probability establishes nothing whatever about one prosecutor's state of mind on one morning in 1992.

A single trial almost never carries the evidence. That is not a defect in this data. It is a description of how the rule works.

08 Twenty-six years of single trials

Compute that same probability for every trial in the file and count how many clear the conventional threshold.

Quantity	Value
Trials examined	211
Trials where the split is significant at $p < 0.05$	99
Expected by chance, on the textbook shortcut of 0.05×211	10.6

99 of 211 trials clear it, against the 10.6 that five percent of 211 gives you: about 9 times as many as chance should produce. It is a large number and it is the wrong number, because the benchmark it is measured against is not true of these tests.

Before you read on, commit to a number. That benchmark is about to be replaced by an honest one, computed from each trial's own margins rather than assumed. You have the numbers going in: of 211 trials, 99 came in significant, where the shortcut expects 10.6. **How many significant trials does an honest benchmark say chance alone would produce out of 211?** Write the number down, and write down whether the gap it leaves is wider or narrower than the 9-fold one the shortcut gives.

Multiplying 211 by 0.05 assumes each trial's p-value is equally likely to land anywhere between zero and one. **That assumption fails here, and it fails in the direction that works against the finding.** These are hypergeometric tests on small integer counts. A trial with four eligible Black jurors and three strikes can produce only a handful of distinct p-values, and none of them is 0.04. A test whose next available value below 0.11 is 0.03 simply never fires at 0.05.

Tests like these, on small whole-number counts, are *conservative*, rejecting less often than their nominal rate says they will.

The honest benchmark therefore has to be simulated rather than assumed. Redraw every trial's strikes from its own margins, keeping its own number of eligible Black jurors and

its own number of strikes, 2,000 times, and count how often the resulting courthouse produces a significant trial.

Quantity	Value
Trials significant at $p < 0.05$	99
Expected by chance: the textbook shortcut	10.6
Expected by chance: simulated from each trial's own margins	3.5
Trials significant at $p < 0.01$	51
Trials where every eligible Black juror was struck	15

Most readers write down a number a little under 10.6 and a gap a little narrower than 9-fold, on the reasoning that a benchmark computed carefully ought to be harder to clear. Chance produces 3.5 such trials, not 10.6. The observed count is 28 times the simulated null rather than 9 times the shortcut. **The correction runs in the direction people almost never expect: being careful about the null makes the disparity roughly three times larger, not smaller.** And 15 of the 211 trials removed every single eligible Black juror.

Nobody should have to take a single threshold on trust, and 0.05 is an arbitrary line in any case. Figure 3 does the count at every threshold from 0.005 to 0.2: what actually happened in this district, against what 2,000 simulated courthouses produced with race doing no work at all.




Figure 3. The result does not depend on choosing 0.05. Trials called significant at every threshold from 0.005 to 0.2. What happened in this district is drawn against the mean and the middle 95% of 2,000 simulated courthouses. In each of those, a trial keeps its

own eligible jurors and its own number of strikes, but race does no work. At the conventional threshold the counts are 99 against 3.5. **The red line never enters the gray band at any threshold on the axis.**

Figure 4 drops the thresholds altogether and puts every trial on the same plane: how many Black jurors the state could have struck, against how many it did.




Figure 4. Every trial, jittered so that overlapping ones stay visible. Black jurors the state could have struck, against Black jurors it struck, one dot per trial, colored by the probability of that split under race-neutral striking. Points on the dashed diagonal are the 15 trials in which every eligible Black juror was removed. 51 of the 211 trials sit below $p = 0.01$ on their own, without any help from the other 210.

09 The trial at the far corner

One dot in Figure 4 is further from the rest than any argument about thresholds can reach.

Trial	Black eligible	Black struck	White eligible	White struck	P
2003-0071-Curtis Flowers III	15	15	21	0	1.8e-10
1998-0011C-D-Krishun Williams	17	14	29	3	1.5e-06
2003-0071-Curtis Flowers IV	16	11	20	0	7.3e-06

The most extreme trial in the file is **the third trial of Curtis Flowers. Fifteen Black jurors were eligible and fifteen were struck. Twenty-one white jurors were eligible and none was.** The probability of that split under race-neutral striking is $1.8e-10$, roughly one in five billion.

Flowers was tried **six times for the same crime by the same prosecutor.** In *Flowers v. Mississippi* (2019) the Supreme Court threw out his conviction, and the majority rested substantially on the pattern of whom that office struck across those six trials. The charges were later dropped, and Flowers went home after twenty-three years.

Notice what it took. Not a statistical argument in a single courtroom, which Figure 2 shows would have been laughed out. Not the district-wide mirror in Figure 1, which no judge had ever been shown. It took one defendant tried so many times by one office that the overall pattern and the individual case finally became the same object. **The rule is enforced at exactly the level where the evidence is weakest,** and *Flowers* is the rare case where the two levels collapsed into one.

10 What can be said, and by whom

What this file supports is that in this district, over twenty-six years, the state's peremptory strikes were not race-neutral. The conclusion does not rest on anybody's judgment about which jurors *should* have been struck, which is the assumption that sinks most disparity findings. It rests on the defense's mirror-image behavior, which supplies the comparison from inside the same rooms.

Almost none of that is available to the person it is supposed to protect. The pattern is overwhelming in aggregate and close to invisible one trial at a time, and the 4.4% probability computed for the 1992 trial is roughly what an ordinary *Batson* motion has to work with. It loses.

11 The argument about all of this

The dispute this evidence enters is not really about the numbers. Nobody seriously contests that prosecutors strike Black jurors at higher rates; studies in Philadelphia, North Carolina and Louisiana found the same thing. The dispute is about what follows from it.

One position, Justice Marshall's, stated in a concurrence to *Batson* itself, is that the peremptory challenge cannot be saved. Requiring reasons, he argued, only teaches lawyers to produce them, and the only remedy is to abolish peremptories entirely. Several states have since restricted them, and **Arizona abolished them outright, in all jury trials, civil and criminal, effective 1 January 2022, the first state to do so.** The Arizona Supreme Court's stated reason was that *Batson* had failed: strikes were still being used to discriminate, by both sides.

The other position is that peremptories protect defendants too, and Figure 1 is evidence for it as much as against it. The defense in these same trials struck 47.3% of eligible white jurors. That is a defendant's lawyer using the tool exactly as it was meant to be used, in a district where the tool was the only protection on offer. Abolish it for the state and you abolish it for the defense.

There is also a methodological objection worth taking seriously, and it is the missing third. Race is unknown for 4,752 of 14,873 jurors. If missingness were related both to race

and to being struck, the rates in Figure 1 could shift. The mirror structure makes that a stretch, since the same missing jurors are missing from both sides of it, but “a stretch” is not “impossible.”

12 What this evidence will bear

The strength of this file is that it is not a summary. It was assembled trial by trial from court documents by journalists at APM Reports. It contains its own control group, which is what makes the central comparison possible at all. And it follows a single decision-maker over twenty-six years, rather than averaging across a profession.

Its limits are equally concrete, and each one bounds a different claim. The missing third of jurors bounds the precision of every rate quoted here. The absence of any strike *reasons* means the file records what was done and never what was said. So it can show that the state's behaviour is not explained by race-neutral striking, without ever showing what explanation the prosecutor offered or the judge accepted.

One district bounds how far any of this carries. Mississippi's Fifth Circuit is not America, and this office was extreme enough to reach the Supreme Court, which is a reason to expect it to be unusual.

The probability model bounds the meaning of every p-value above. It assumes that under a race-neutral rule any set of strikes was equally likely. That is the assumption being tested, not a claim about how juries really work.

And nothing here touches outcomes at all. This is a record of who was excluded, not of what verdicts followed, so whether these juries convicted differently is not a question this data can answer.

Three things are left genuinely open, and more of this table will not close any of them. Whether peremptory challenges should exist at all is a question the data bears on and cannot decide. What a court ought to do with aggregate evidence is unsettled in a deeper way: *Flowers* used a history spanning six trials, and almost every defendant has one. And whether the pattern has changed since this file ends is simply unknown, because nobody has read the next twenty-six years of court records the way APM Reports read these.

The prosecutor in 1992 was not required to give a reason, and did not give one. What the record shows is that across 211 trials, the reasons that were never given sorted the panel by race 28 times more often than chance. The lawyer sitting at the other table, drawing from the same panel on the same morning, sorted it the opposite way at nearly the same rate. Neither of them had to explain that, and for most of those twenty-six years nobody asked.

13 Sources

Data

- APM Reports, *In the Dark* season 2, jury data. <https://github.com/APM-Reports/jury-data-jurors.csv>, [trials.csv](https://github.com/APM-Reports/jury-data-trials.csv) and [voir_dire_answers.csv](https://github.com/APM-Reports/jury-data-voir_dire_answers.csv), coded from the court records of 225 trials in Mississippi's Fifth Circuit District, reduced here to the 211 trials with at least ten eligible jurors and the four strike rates in Figure 1; every number

above is computed from the files it wrote, at the moment this document was rendered.

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`strikes.csv`, `trials.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `strikes.csv` has `eligible` and `struck` by side and race. Compute the rate for each combination. How large is the gap, and which side produces it?
- `trials.csv` is one row per trial. Count the trials where every eligible Black juror was struck. Would any single one of those trials prove anything?
- Compute the strike rate per trial and look at the spread. Is this a pattern across a district, or a few prosecutors?
- `trials.csv` names the defendant in each row. Pick one trial and work out what the numbers in that row would have looked like to the judge at the time.

Cases and literature

- *Batson v. Kentucky*, 476 U.S. 79 (1986), including Marshall, J., concurring.
- *Flowers v. Mississippi*, 588 U.S. ____ (2019).

Further reading

- Grosso, C. M. and O'Brien, B. (2012). "A Stubborn Legacy: The Overwhelming Importance of Race in Jury Selection in 173 Post-Batson North Carolina Capital Trials." *Iowa Law Review* 97: 1531-1559.

REBUILD THIS DATA YOURSELF, WITH AN AI ASSISTANT

I want to rebuild a public dataset from its original source, without any starting files. Please do the following and show your work.

THE SOURCE. APM Reports, the investigative newsroom behind the podcast "In the Dark", season 2: jury-selection records from 225 trials in Mississippi's Fifth Circuit Court District over 26 years, read out of court files by reporters. The data sit in the GitHub repository

`APM-Reports/jury-data`; fetch the raw files

<https://raw.githubusercontent.com/APM-Reports/jury-data/master/jurors.csv>

<https://raw.githubusercontent.com/APM-Reports/jury-data/master/trials.csv>

No account or key is needed. In the juror file each row is one

prospective juror in one trial – 14,874 of them – with the juror's race, who was allowed to strike them (`strike_eligibility`), and who actually struck them (`struck_by`). The trial file is one row per court proceeding, 305 in all.

WHAT TO BUILD.

1. Strike rates by race, separately for the prosecution and the defense. Count a juror in a side's denominator only if that side could have struck them – a juror removed for cause, or never reached, was available to nobody. Use only jurors coded Black or White: race is unknown for 4,752 jurors, a third of the file, and they have to be left out.

2. A per-trial table: for every trial with at least ten eligible jurors of known race, the Black and white jurors the state could strike, and how many of each it struck.

CHECK YOUR WORK. The copy this book used was fetched in August 2026.

If your rebuild matches it, you should find:

- the state struck 902 of 1,811 eligible Black jurors and 372 of 3,318 eligible white jurors – about 50% against 11%
- the defense ran the mirror image: 127 of 846 Black jurors struck, 1,330 of 2,812 white
- 211 trials clear the ten-juror bar

The repository is the archive of a finished investigation and rarely changes; if your counts differ, check whether the newsroom revised the files before assuming you made an error.