

Traffic Stops and the Benchmark Problem

Traffic stops in one city, and a disparity that changes sign when you change the denominator

84-355 Democracy's Data

August 2026

Over ten years the San Francisco police recorded 905,070 traffic stops. Hispanic drivers were stopped 0.86 times as often as white drivers, which makes them the least-stopped large group in the city. Hispanic drivers were also stopped 1.01 times as often as white drivers, which is exactly even.

Both figures come from the same stops. They differ only in what the stops were divided by: the first by the number of Hispanic residents of San Francisco, the second by the number of Hispanic San Franciscans who drive to work. Neither denominator is dishonest. Neither is right, either, and the choice between them decides the sign of the finding.

That is the difficulty at the center of every argument about policing data, and it is not the difficulty most people expect. The problem is not that records are missing. **The problem is that the comparison is ill-defined**, and the useful part of this chapter is what remains sayable once that is admitted.

01 Nobody was required to write any of this down

There is no national record of traffic stops, because no law creates one.

What exists is whatever each police department chose to keep, in whatever format its records system happened to use. It was released only where a state disclosure law obliged it. The Stanford Open Policing Project assembled what there is the slow way, by public-records requests, one agency at a time. Then it did the harder half. It forced dozens of mutually incompatible extracts into one layout, with the same column names and the same coding, so that a method could be carried from one department to another.

This chapter uses one department out of that collection, San Francisco, from the project's 2020-04-01 release. That is 905,070 stops between 2007 and 2016, one row each, in a file far too large to travel inside a chapter. That standardization is the project's real contribution, and it is also the thing to keep in view. Every column below has been through somebody's translation.

The stops answer only half the question, so the other half comes from an entirely different kind of source. The population figures are American Community Survey counts for San Francisco County, taken from the two five-year releases that tile the stop window exactly. They were pulled through the Census API on 2026-08-10. That is 22 calls in all, each logged with the size and SHA-256 of what came back, and the API key stripped out of the recorded address.

Set those two paragraphs beside each other and the asymmetry is the point. The denominator came from an agency whose job is publishing numbers, in an afternoon, over a documented interface, and anyone can re-run those calls tomorrow and check the answers against the hashes recorded here. The numerator, the top number of every rate here, came from police departments that publish nothing, one letter at a time, and it cannot be re-run at all. It is a release, not a feed. It is dated, the stops end in 2016, and a version of this chapter about last year would have to begin by filing the requests again and waiting.

The cost of getting data this way is that the data has an expiry date somebody else controls. None of that is visible in the two small CSVs sitting here, which look exactly like any other download.

Both are the right evidence for this chapter, and neither is the evidence anyone would have designed. The outcome test needs one thing no published summary can supply: whether a search happened and what it turned up, carried on the same record. Plenty of agencies record one and not the other. The file that would answer the question outright – what the drivers who were *not* searched had in their cars – cannot exist, because nobody looked, which is what “not searched” means. The census side is imperfect in a way worth naming precisely.

The residents table splits the population cleanly. The adult and commuter tables are “race alone” versions that fold in Hispanic respondents of those races, nudging two of the groups’ counts up. And the commuter table covers workers who *live* in San Francisco, so it counts a resident driving out to San Mateo and misses the one driving in from Oakland. Where no honest match existed at all – the “other” category, past the residents rung – the row was left out rather than approximated.

Cleaning here mostly meant throwing things away on purpose. The source file carries a row per stop; it is streamed once and reduced to 5 rows by race and a small year-by-race table, so nothing individual-level is stored here at all. Its boolean columns are text, and only the literal string TRUE is read as true, which means a blank counts as a no rather than as a mystery. Nothing was imputed, meaning no missing value was filled in by guessing, and nothing was reweighted. What a reader would have to redo is the download and that one reduction; everything after it is arithmetic on five rows.

02 A traffic stop is a decision, and the file records only the outcome

A traffic stop is the most common encounter between an American citizen and the state’s coercive power. Roughly one adult in ten is stopped in a given year. What happens next is decided by one person, at the roadside, in seconds, under a legal standard the law leaves deliberately vague. A warning, a citation, a search, an arrest.

Each row of the Stanford Open Policing Project’s San Francisco file is **one stop**, with roughly these fields:

Field	What it holds
date, time	when the stop happened
location	district or intersection

Field	What it holds
subject_race	the officer's perception, entered on the form
subject_sex, subject_age	as recorded by the officer
reason_for_stop	the violation cited
search_conducted	TRUE / FALSE
search_basis	consent, probable cause, incident to arrest
contraband_found	TRUE / FALSE, recorded only if a search happened
outcome	warning, citation, arrest

Two entries in that list decide everything that follows.

subject_race is the officer's perception, not the driver's self-report. It is a box somebody ticked at the roadside. That is a limitation, and it means the variable is not the census's variable and the two must never be silently mixed. But for a question about *the officer's behavior* it is arguably the correct variable, because what matters is what the officer believed he was looking at.

contraband_found exists only where search_conducted is true. There is no counterfactual anywhere in this file. You never learn what a driver who was not searched had in the car, and that absence turns out to be the whole problem.

Quantity	Value
Stops in the file	905,070
Years covered	2007 to 2016
Stops that led to a search	53,381
Searches that found contraband	7,976
Race categories recorded	5

Note how narrow the interesting part is. **Searches are rare**, 5.9% of stops, and everything this chapter can establish happens inside that thin slice.

03 Nine hundred thousand rows, and the two that show the work

That field list was a description. This is the file. Below are the header line and two stops from `ca_san_francisco_2020_04_01.csv`, the Stanford release, as they come out of the zip. The lines are long, and are folded at their commas to fit the page. `location`, `lat` and `lng` are replaced with `[redacted]`, because a street corner and a timestamp are a record about a person, and nothing here needs them. Everything else is verbatim.

Column as it arrives	What it holds	Stop 1	Stop 2
raw_row_number	the row's number in the department's own export	869921	870048
date	date of the stop	2014-08-01	2014-08-01
time	time of the stop	00:01:00	14:34:00
location	where it happened — redacted here	[redacted]	[redacted]
lat	latitude — redacted here	[redacted]	[redacted]
lng	longitude — redacted here	[redacted]	[redacted]
district	police district	NA	NA
subject_age	the driver's age	NA	NA
subject_race	the driver's race, as standardised by the project	asian/pacific islander	white
subject_sex	the driver's sex	female	male
type	vehicular or pedestrian	vehicular	vehicular
arrest_made	whether an arrest followed	FALSE	FALSE
citation_issued	whether a citation was issued	FALSE	TRUE
warning_issued	whether a warning was issued	TRUE	FALSE
outcome	the outcome, collapsed to one word	warning	citation
contraband_found	whether contraband was found	NA	TRUE
search_conducted	whether any search happened	FALSE	TRUE
search_vehicle	whether the vehicle was searched	FALSE	TRUE
search_basis	the stated legal basis for the search	NA	other
reason_for_stop	why the stop was made	Mechanical or Non-Moving Violation (V.C.)	Moving Violation
raw_search_vehicle_description	the department's own words for the search	No Search	Search without Consent, Positive Result
raw_result_of_contact_description	the department's own words for the outcome	Warning	Citation

Read the last two columns first, because they are the whole lesson. They are prefixed `raw_`, and they hold the San Francisco Police Department's own words in the department's own vocabulary — phrases nobody outside that department would guess. Beside them sit the standardized columns this chapter actually reads. Here is the same two stops, with each department phrase set against what the release made of it:

The department wrote	The release records
No Search	search_conducted = FALSE, contraband_found = NA
Warning	outcome = warning

The department wrote	The release records
Search without Consent, Positive Result Citation	search_conducted = TRUE, contraband_found = TRUE outcome = citation

Somebody read the left column and wrote the right one. The Stanford project keeps the original string beside its translation precisely so that the translation can be audited, and this chapter's entire finding rests on it being right. It is a judgment call made by people who are not in this chapter, on material this chapter never sees.

Three more things in those 22 fields decide what can be asked. `contraband_found` is NA on the first stop and TRUE on the second, and it is not missing data: it is NA because nobody searched, so there was nothing to find. That single pattern is why the counterfactual does not exist. `district` and `subject_age` are NA on both — columns the common schema defines and this department did not fill, which is what standardization looks like from the inside. And every boolean is the text TRUE or FALSE, alongside the text NA, so something has to decide what a blank means. Here only the literal TRUE counts as true, which turns every blank and every NA into a no rather than a mystery.

Now the clean file, which is small enough to print entire. This is `by_race.csv` in full, exactly as the build writes it:

Race	Stops	Searched	Contraband found	Arrests
white	372318	11707	2834	3970
asian/pacific islander	157684	2861	1033	1221
black	152196	23622	2183	3540
hispanic	116014	11445	1165	2284
other	106858	3746	761	910

905,070 stops became 5 rows. Every column of the raw file was thrown away except `subject_race`, `search_conducted`, `contraband_found`, `arrest_made` and the year inside `date`. That reduction was deliberate rather than incidental. **No individual-level record is stored here at all**, so the folder cannot be re-identified, re-purposed, or quietly used to answer a question it was not collected for. What that costs is everything conditional. No question about time of day, or district, or the interaction of race and sex, can be asked of these five rows, and getting it back means re-running the download.

The grouping key is `subject_race`, whose 5 values are the officer's categories rather than the census's. `asian/pacific islander` is one box on a form, and several dozen lines of the census's race table. The chapter never merges the two without saying so.

A row with a blank in that column is dropped by the build rather than filed under `other`. "The officer did not record it" and "the officer recorded other" are different facts, and only one of them is about the driver. In this release the rule costs nothing, because the 5 categories account for every one of the 905,070 stops. But a rule that happens to catch no rows today is still a decision, and the next release is not obliged to cooperate.

04 The population at risk of being stopped was never counted

Here is the table anybody produces on opening this file with no theory at all.

Race (officer-perceived)	Stops	% of all stops
white	372,318	41.1
asian/pacific islander	157,684	17.4
black	152,196	16.8
hispanic	116,014	12.8
other	106,858	11.8

Black drivers were 16.8% of stops. San Francisco's Black population over this period was roughly 5 to 6% of residents. It is very tempting to set those two numbers side by side, and it is a mistake, because **the residential population is not the population at risk of being stopped.**

To compare a stop count against a population you would have to know who was *driving*, where, and when. San Francisco is a commuter city: a large share of the cars on its streets on a weekday belong to people who live somewhere else. Driving is not distributed like residence. It varies with car ownership, with work schedules, and with which arterial roads happen to run through which neighborhoods.

And a perfect count of drivers still would not be enough, because you would also need to know **how many drivers of each group were doing something that warranted a stop.** Nobody records that. Nobody could.

You would need	Does it exist?
Who was driving	No
Where and when they were driving	No
Who committed a violation	No
Who was carrying contraband	Only for those searched

This is not a funding problem that a better survey would fix. **There is no census of who was on the road**, the denominator cannot be constructed, and the share in the table above therefore proves nothing on its own. Most public arguments about policing stall exactly here, with each side asserting a denominator neither can produce.

Figure 1 shows what is at stake in the choice. The same 905,070 stops are divided by three *different* real denominators from the American Community Survey. Each group is plotted as a ratio to white drivers, so a value of one means "stopped at the same rate as white drivers."

Before you look at the figure, put a number on its fourth column. The first three divide these stops by people, and you have already seen what two of them do to Hispanic drivers: 0.86 against residents, 1.01 against people who drive to work — under-stopped, then dead level. The fourth column asks something else entirely. It abandons

population and takes the stop as given: of the drivers an officer had already pulled over, what share were searched, again as a ratio to white drivers. Commit to a figure for Hispanic drivers before you read on. Below one, like the first column? At one, like the third? Above it — and if above, how far above?




Figure 1. The same stops, four denominators, and a disparity that changes sign. Stop rate relative to white drivers on a log scale.

The first three columns divide the identical 905,070 stops by three real American Community Survey counts for San Francisco County. The fourth drops population denominators entirely and asks instead how often a stop led to a search. Against residents Hispanic drivers read 0.86, against adults 0.98, against people who drive to work 1.01, and on searches 3.14. Black drivers move from 3.07 to 4.79 across the three, and read 4.94 on searches. The gray band is the denominator nobody has.

Read the Hispanic line first, because it is the one that breaks the argument most people want to make. Against residents Hispanic drivers look *under-stopped*, at 0.86 times the white rate. Against adults the figure is 0.98. Against people who drive to work it is 1.01, dead level. Three defensible denominators, three different stories, **and the sign of the disparity flips between the first and the third.**

Black drivers move too, from 3.07 to 4.79, in the other direction. Which conclusion you reach has been decided by a choice the data cannot adjudicate, and the column that would adjudicate it sits in the gray band, unbuilt.

The fourth column is where most guesses fail. Having watched Hispanic drivers come in under one and then level, readers tend to keep them somewhere near one. On searches per stop they read 3.14 — not near one, and not on the side of it the first three columns pointed to.

There is a further reason to distrust those first three columns even on their own terms. The ACS counts are self-reported race; the stop file records an officer's perception. They are different variables, and every ratio joining them is a rough benchmark rather than a measurement.

So much for that. Notice the fourth column, which is where the rest of this chapter goes.

05 Take the stop as given

Change the question. Instead of asking who gets stopped, treat the stop as having happened and ask **what happened next**. Among the people an officer had already pulled over, how often did a search follow?

Race	Stops	Searches	% of stops searched
black	152,196	23,622	15.52
hispanic	116,014	11,445	9.87
other	106,858	3,746	3.51
white	372,318	11,707	3.14
asian/pacific islander	157,684	2,861	1.81

Black drivers were searched in 15.52% of stops and white drivers in 3.14%, a ratio of 4.9 to one. Hispanic drivers, the group that looked under-stopped or level on every population denominator available, were searched 3.14 times as often as white drivers.

This is a better comparison than the share of stops. Both groups have already been picked out by the same officers under the same conditions, and no count of who was driving is needed to make it. It is still not enough, and the objection is the standard one: **perhaps the stopped populations genuinely differed**. If officers were on average encountering more grounds for suspicion among drivers of one group than another, a higher search rate would follow from correct policing rather than from a lower bar.

That objection cannot be answered with search rates, because answering it would require knowing how suspicious each driver actually was, which is precisely what nobody observes.

06 Condition once more

In 2001 three economists, **Knowles, Persico and Todd**, proposed a way out that needs no denominator at all. Condition one step further: take the search as given, and ask what the officer found.

An officer searches when suspicion crosses some internal threshold. If the same threshold is applied to everyone, the searches it authorises should succeed at roughly equal rates across groups. A lower threshold means more searches on weaker evidence, and weaker evidence means a lower success rate. **So unequal hit rates imply unequal thresholds.**

The comparison is now between searches, not between populations. Nothing has to be known about who was driving, who was stopped, or how suspicious anyone was. Figure 2 draws the price of that, which is that each conditioning step throws away nearly everything.




Figure 2. Each conditioning step buys a valid comparison by discarding almost all the data. 905,070 stops, of which 53,381 produced a search (5.9% of stops) and 7,976 produced contraband (0.88% of stops). The dashed box on the left is everyone who was on the road and was never counted; it has no width because it has no number. The final band is not a rendering fault, and neither renderer inflates it: 7,976 finds really is 0.88% of 905,070 stops.

It is worth predicting the answer before reading it. Given a search rate roughly 5 times higher for Black drivers, should the share of searches that turn up contraband be higher, lower, or about the same? Each answer has a serious defender, and each implies a different account of what the department was doing.

07 What the searches found

Race	Searches	Contraband found	% of searches finding contraband
black	23,622	2,183	9.2
hispanic	11,445	1,165	10.2
other	3,746	761	20.3
white	11,707	2,834	24.2
asian/pacific islander	2,861	1,033	36.1

Contraband was found in 9.2% of searches of Black drivers and 24.2% of searches of white drivers. Searched 4.9 times as often, and found with something 2.6 times less often. Under the outcome test that is evidence of a lower evidentiary bar, and it required no denominator to reach.

Two groups could be a coincidence. Figure 3 puts all five on the same plane: search rate along the bottom, hit rate up the side.





Figure 3. The more often a group was searched, the less often the searches found anything. Share of stops leading to a search against share of searches finding contraband, one circle per officer-perceived race group, circle area proportional to stops. The dashed line is the overall hit rate, 14.9%, where equal thresholds would place every group. Black drivers sit at 15.52% searched and 9.2% found; Asian and Pacific Islander drivers at 1.81% searched and 36.1% found. The ordering is monotone across all 5 groups.

That steady ordering is what makes an innocent explanation hard to construct. It would have to account for the Black and Hispanic results at the same time as the reverse pattern among Asian and Pacific Islander drivers. They were searched least often and found with contraband most often, across 905,070 stops.

08 Ten years, and the ordering never flips

Ten years are stacked into every number so far. A pattern visible only in aggregate, produced by one or two unusual years, would be a different finding from one that repeats.



Year	White hit rate (%)	Black hit rate (%)	White higher?
2007	26.2	11.1	yes
2008	24.2	9.9	yes
2009	19.3	8.7	yes
2010	27.6	8.9	yes
2011	21.6	6.0	yes
2012	20.7	5.9	yes
2013	21.9	7.6	yes
2014	30.7	9.4	yes
2015	29.8	13.1	yes
2016	28.2	16.0	yes

Figure 4. The levels move a great deal; the ordering does not move at all. Share of searches finding contraband, by year and officer-perceived race, 2007 to 2016. Every group's hit rate falls through 2012 and rises afterwards, so the level is clearly responding to something outside this file. In all 10 years searches of white drivers found contraband more often than searches of Black drivers.

Whatever produced the gap in Figure 3, it was not one unusual year.

09 What can be said, and what cannot

Within this department, over these ten years, searches of Black and Hispanic drivers were authorized on weaker evidence than searches of white drivers. That is a statement about the bar an officer applied, and it was reached without ever knowing who was driving, which is the whole reason it can be made at all.

What cannot be said is how much of that gap is individual animus, how much is departmental policy, and how much is a consequence of where officers were deployed. All three produce the same aggregate pattern, this data cannot separate them, and the distinction matters enormously for anyone proposing to do something about it.

10 The argument about this method

The part usually left out of accounts of the outcome test is what happened when its inventors ran it. **Knowles, Persico and Todd applied their own test to Maryland and found no evidence of prejudice against Black motorists.** They concluded that the disparity in search rates there fitted officers maximizing arrests, rather than acting on racial animus. The method and the conclusion are separable, and it is worth sitting with the fact that the tool used above was built by people whose own application of it came out the other way.

The most serious objection is technical and unresolved. The outcome test compares **average** hit rates, but the theory it rests on is about the **marginal** search, the weakest one an officer chooses to make. If the underlying distribution of suspiciousness differs between groups, average hit rates can differ even when the threshold is identical. Simoiu and colleagues (2017) formalized this problem, named it infra-marginality, and proposed alternatives. It is a good reason to treat the outcome test as strong evidence rather than as proof, and the hedged phrasing two paragraphs above is hedged for exactly that reason.

The largest study to date is Pierson et al. (2020), covering roughly **100 million stops** nationally. It used two designs at once: the outcome test, and a “veil of darkness” test, which exploits the fact that a driver’s race is harder to see after dark. It found evidence that the bar for searching Black and Hispanic drivers was lower than for white drivers. The literature has not converged, and nothing here should be read as claiming it has.

11 What this evidence will bear

The strength of this file is that it is not a summary. It holds individual records, and it records both the search and its outcome, which is what makes the outcome test possible at all; many jurisdictions record one and not the other. The coding is standardized across jurisdictions by the project that assembled it, so the method travels even where the levels do not.

Its limits bound different claims, and it is worth being specific about which. Race here is an officer’s perception entered on a form. That is the right column for a question about the officer’s behavior, and the wrong one for any comparison against census categories. So Figure 1’s first three columns are weaker evidence than Figure 3, even before their denominators are argued about.

"Contraband found" is itself a judgment. What counts as a find, and whether a find is written down, are decisions made by the same officers whose behavior is being measured, so the hit rate is not an untouched outcome.

One city through 2016 limits the reach of all of it. Nothing here describes San Francisco today, and the levels certainly do not travel. The denominator problem never actually goes away, because the outcome test routes around it rather than solving it. And nothing here is about an individual officer: these are department-level patterns, and a department-level pattern can be produced by a small number of people behaving very differently from their colleagues.

Three questions stay open, and none of them will be closed by this file. Infra-marginality is real and unresolved, and it bounds the confidence of the central claim rather than contradicting it. Whether the pattern generalizes is a question for the national data, not for one city. And what produces the threshold is invisible here: animus, policy, deployment, or the assignment of officers to neighborhoods. That is inconvenient, because those four would call for four different remedies.

What survives all of that is the shape of Figure 1. Ask who gets stopped and the answer depends on a denominator that does not exist. The same Hispanic drivers can be shown as under-stopped, or as exactly level, depending on a choice made before any analysis begins. Ask instead what happened to the people an officer had already pulled over, and the denominator problem disappears. The officer supplied the comparison group. The disparity that survives that move is the one worth arguing about, and it is larger, not smaller, than the one that does not.

12 Sources

Data

- Pierson, E. et al. Stanford Open Policing Project, San Francisco, CA, 2020-04-01 release. <https://openpolicing.stanford.edu/data/> — one row per stop, reduced here to the by-race and by-year counts every figure above is drawn from.
- U.S. Census Bureau, American Community Survey, San Francisco County — the three population denominators in Figure 1. The exact ACS table behind each row is recorded, and `acs_denominators.csv` carries it beside the count. The counts come from the Bureau's own five-year ACS API, which documents every table and geography it will answer for. <https://www.census.gov/data/developers/data-sets/acs-5year.html>

The data itself

Every figure above is drawn from these tables. They are plain CSV, they open in a spreadsheet, and they are yours to take further.

`acs_denominators.csv`, `by_race.csv`, `by_year.csv`

Your turn

None of these is answered above, and every one can be answered from the tables you just opened.

- `by_race.csv` has `stops`, `searched`, `contraband_found` and `arrests`. Compute a search rate and a hit rate for each group. Do the two measures point the same way?

- `acs_denominators.csv` offers several denominators for the same population. Re-compute the stop disparity with each one. How much does the conclusion depend on the choice?
- `by_year.csv` runs the same counts across years. Plot the search rate by race over time. Did anything change, and when?
- A hit rate needs no population denominator at all. Explain why, then say what it still cannot tell you about who was stopped.

Academic literature

- Knowles, J., Persico, N. and Todd, P. (2001). "Racial Bias in Motor Vehicle Searches: Theory and Evidence." *Journal of Political Economy* 109(1): 203-229.
- Pierson, E., Simoiu, C., Overgoor, J. et al. (2020). "A large-scale analysis of racial disparities in police stops across the United States." *Nature Human Behaviour* 4: 736-745.
- Simoiu, C., Corbett-Davies, S. and Goel, S. (2017). "The problem of infra-marginality in outcome tests for discrimination." *Annals of Applied Statistics* 11(3).

REBUILD THIS DATA YOURSELF, WITH AN AI ASSISTANT

I want to rebuild a public dataset from its original sources, without any starting files. Please do the following and show your work.

THE FIRST SOURCE. Stanford Open Policing Project, stops by the San Francisco Police Department, release dated 2020-04-01. One zip, about 25 MB, at

https://stacks.stanford.edu/file/druid:yg821jf8611/yg821jf8611_ca_san_francisco_2020_04_01.csv.zip

No account or key is needed. It unzips to a 177 MB CSV of 905,070

stops, 2007 through 2016 – each row one stop, with the driver's race as the officer perceived it, whether a search was conducted, whether contraband was found, and whether an arrest was made.

THE SECOND SOURCE. U.S. Census Bureau, American Community Survey

5-year estimates, through the Census Data API – population

denominators for San Francisco County (state 06, county 075). Two vintages tile the stop years exactly, 2007-2011 and 2012-2016:

<https://api.census.gov/data/2011/acs/acs5>

<https://api.census.gov/data/2016/acs/acs5>

This one DOES need a key. It is free and instant from

https://api.census.gov/data/key_signup.html – sign up, then attach

the key to every request. Keep the state and county codes as text:

"06" stops being California the moment it becomes 6.

WHAT TO BUILD.

1. One row per race group: stops, searches, contraband found, arrests. Then the two rates the comparison turns on – the share of stops that led to a search, and the share of searches that found contraband.
2. The same counts split by year, to see the pattern over time.
3. Population counts by race from both ACS vintages, averaged: total residents, adults 18 and over, and people who drive to work.

CHECK YOUR WORK. The copies this book used were fetched in August 2026. If your rebuild matches them, you should find:

- 905,070 stops carrying a race code
- 372,318 stops of white drivers and 152,196 of Black drivers
- 11,707 searches of white drivers found contraband 2,834 times;
23,622 searches of Black drivers found it 2,183 times

The release date is part of the URL, so this exact file should not change; if the project posts a newer release, its numbers will differ because the data were extended or revised, not because you made an error.