Key: The Races Nobody Contested

84-355 Democracy's Data | House competition | Instructor copy — do not distribute

Jonathan Cervas

August 11, 2026

All code is given to students. Grade the interpretation.

Structure, as of the rebuild. The brief (`house-competition-brief.Rmd`) is now a chapter rather than a parallel copy of the lab: prose organised by argument, under idea headings, with two numbered figures and no numbered steps. The lab (`house-competition.Rmd`) still runs Steps 1-10 with the code exposed. **Students should read the brief first. Step numbers below refer to the lab**; the chapter is cited by figure number or section title.

Pairs with the redistricting lab. That one measures district partisanship from presidential votes; this one measures what happened in the House races themselves. Same question, two instruments, and they mostly agree — which is worth saying, because they need not have.

Every number below was run against the data files.

The session in one sentence

In the 1950s, 71.3% of Southern House races had one candidate, that is nearly gone, competition fell anyway, and the whole finding lives in a column that is empty.

The findings

```
dec <- function(v) round(tapply(by[[v]], 10 * (by$year %/% 10), mean), 1)
rbind(uncontested_south     = dec("pct_uncontested_south"),
      uncontested_non_south = dec("pct_uncontested_non_south"),
      competitive           = dec("pct_competitive"),
      landslide             = dec("pct_landslide"))
```

```
#>                        1940 1950 1960 1970 1980 1990 2000 2010 2020
#> uncontested_south      60.5 71.3 45.5 35.5 34.4 27.6 27.3 18.8  9.7
#> uncontested_non_south   7.6  5.7  3.4  6.7  8.7  8.1 10.2  8.7  6.7
#> competitive            23.4 21.1 18.6 15.2 12.6 15.4  9.6 14.2 16.4
#> landslide              13.3 10.6 17.2 26.4 31.4 22.1 25.4 21.9 21.6
```

- **South uncontested: 60.5% (1940s) → 71.3% (1950s) → 18.8% (2010s), and 9.7% so far in the 2020s.** Non-South never leaves 3.4–10.2%.
- **Competitive races (within 5 points): 23.4% → 9.6% (2000s) → 14.2% (2010s).** The single worst year is **2004, at 5.1%.**
- **Landslides (20+ points): 13.3% → peak 36.3% in 1986.**
- 2,568 of 17,403 races have no vote share, and where the file records incumbency at all, **2,221 of 2,223 of those are incumbent wins.**

The three things to get said

1. Start with the 14,777 empty rows. Load the file raw and show them. This is the most transferable five minutes in the lab: **a standard research dataset, used for decades, that silently doubles in size if you do not look.** Ask what would have happened had a student computed a mean without checking — every number halved, no error, no warning.

2. The missing column is the finding. There is no uncontested flag. `dv` is a *two-party* share, and one candidate means no two-party share exists. **Make the room reconstruct that inference from the evidence** — of the uncontested races with an incumbency flag the incumbent won 2,221 of 2,223, and missingness is 34.7% South against 7.3% elsewhere — rather than being told. It is the same "absence is the signal" move as the CPS lab's blank pre-1980 rates.

3. Then the substance, which is genuinely striking. In the 1950s South, the Democratic primary *was* the election. That is not a close race or a safe seat — it is no general election at all, in a democracy, for decades. **And competition fell nationally anyway.** One-party dominance did not end; it changed form, from empty ballots to lopsided ones.

The line for the board: a race with one candidate is the least competitive race there is, and it has no margin at all.

Step 6 is the one to slow down on

The denominator question — all races versus contested races — is the third or fourth appearance of "per what?" this term, and here it is unusually clean:

```
cont <- d[!d$uncontested, ]
rbind(all_races      = round(100 * tapply(d$competitive, 10 * (d$year %/% 10), mean), 1),
      contested_only = round(100 * tapply(cont$competitive, 10 * (cont$year %/% 10), mean), 1))
```

```
#>                 1940 1950 1960 1970 1980 1990 2000 2010 2020
#> all_races       23.3 21.1 18.6 15.2 12.6 15.4  9.6 14.2 16.4
#> contested_only 29.3 26.9 21.6 17.6 15.0 17.8 11.3 16.1 17.8
```

Both are defensible; they answer different questions. **Push students to argue for including the uncontested races**, since a race with one candidate is maximally uncompetitive and excluding it flatters the system.

Walking the steps

Steps	What they do	Min	Watch for
1-3	One district-election, the scope, and the three definitions	8	**The NA in `dv` is the finding.** A student who would drop those rows has deleted the answer.
4-5	The flat national trend, and who holds uncontested seats	7	20.2% to 8.7% across 78 years, with single-election swings wider than the whole drift. It looks like nothing happened. Let that land as a disappointment.
6	Split the country in two	12	**The pivot.** Two opposite trends cancelling in a national average.
7-8	Competitive and landslide; then split districts	9	Four measures, four methods, same direction.
9	All five series at once	5	The converging uncontested lines are the image to remember, and the heavy black national average sits between them the whole way: that is Step 6 in one picture.
10	What you can say	5	Elite decision, not voter verdict: of the uncontested races whose incumbency the file records, 99.9% had an incumbent.

Two figures were added to the chapter and are worth planning around. Figure 1, under *A challenger on the ballot is a low bar*, is the material of Steps 7-8: a per-decade distribution of winning margins among contested races, which is the same claim as the 5-point and 20-point cuts without the thresholds; use it to ask what a different cut point would have done. Figure 2, under *Four measures, one direction*, is the chart of Step 9, and it now includes the national uncontested share as a heavy black line, so the "correctly computed mean of two subpopulations" is visible sitting between the two regional lines instead of being asserted underneath the figure.

Step 6 is the pivot. Show the flat national line in Step 4, let the room conclude 'nothing changed', then split by region. The aggregation trap is more memorable when they walked into it.

If you are short on time, protect the pivot and the closing step. The numbers are what make the argument concrete; the sequence is the argument.

Option A — Midterms

Verified output

```
rbind(competitive = round(100 * tapply(d$competitive, d$midterm, mean), 1),
      uncontested = round(100 * tapply(d$uncontested, d$midterm, mean), 1))
```

```
#>                0    1
#> competitive 15.7 15.7
#> uncontested 14.3 15.3
```

Almost nothing. Competitive races: 15.7% in presidential years, 15.7% in midterms. Uncontested: 14.3% against 15.3%.

3/4: Runs it and reports the difference is small.

4/4: Says the null is the interesting result, and says why it is surprising — midterms have lower turnout and a different electorate, so one might expect different margins. **District partisanship swamps the electorate's composition.**

4/4, best: Notes that a gap this small — 0.0 points on competitiveness, 1.0 on uncontested — across 17,403 races is precisely estimated, so this is a real null rather than an underpowered one — and contrasts it with the *midterm penalty* lab, where the president's party reliably loses seats. **Midterms move which party wins without moving how close the races are.** That is a genuinely sharp observation and worth full credit.

When your source and your predecessor disagree

Worth ten minutes, because the answer is not always "trust the elder."

The Clerk parse and Jacobson overlap for 2004–2014 — 2,260 districts, median difference **0.03 points**. Only 13 districts differ by more than a point. Every one was run down, and they did not all resolve the same way:

Districts	Who was wrong	How we knew
SC-06, 2004	**Us**	"Constitution" was missing from the fusion party list, so 4,157 votes never reached the Republican. Adding it matched Jacobson exactly.
NY ×4	**Us**	Fusion lines listing two parties ("Conservative, Libertarian") went unmatched.
LA-01, 2014	**Us**	We took the leading candidate of each party; the Clerk's own Recapitulation *sums* them, and so does Jacobson.

Districts	Who was wrong	How we knew
NE-03 2004, UT-01 2006	**Jacobson**	The candidate blocks *and* the Recapitulation table both give our figures. Two independent presentations inside the official document agree with each other and disagree with him.
CT ×5, 2008	**Jacobson**, arguably	He counts fusion lines in New York and South Carolina and not in Connecticut. CT-01 minus the Working Families line is 71.68; he reports 71.7.
LA ×2, TX-22, TX-23	**Nobody**	Open primaries, December runoffs, and a write-in campaign. No rule reproduces him, and the Clerk mixes two elections in one row. Flagged, not fixed.

The teachable move is the third column. We did not decide who was right by seniority or by which number looked nicer. The Clerk's document happens to print the same results twice — once as candidate blocks, once as a per-state Recapitulation — and where those two agree with each other, they outrank a single derived figure from anywhere else.

Ask the class what they would have done with a 2-point disagreement against a dataset that scholars have used for forty years. The honest answer is that most people quietly adopt the established number, and that this is how errors survive.

The column-name trap, and what to do with it

Worth five minutes of class time, because it happened here.

The Jacobson file carries three similar columns:

column	what it is
`dv`	Democratic share of the two-party **House** vote, this election
`dvp`	the same thing for the **previous** election
`dpres`	Democratic share of the two-party **presidential** vote in the district

`dvp` reads as "d-v-presidential" and is "d-v-previous". **An earlier build of this lab used it as the presidential vote**, computed split districts as `(dv > 50) != (dvp > 50)` — which is a *seat that changed party*, not a split district — and then concluded that the file's own `split` column was unreliable because it disagreed.

The file was right. Three things established that, and they are worth showing:

1. `dvp` correlates with the previous election's `dv` at **r = 0.968**, mean absolute difference **0.72 points**, across 12,226 district-years.

2. Recomputed with `dpres`, the flag agrees with the file's `split` column **100%** of the time; recomputed with `dvp` it agrees only 71.9% of the time.
3. `dpres` for AL-01 in 2012 is **37.70**; The Downballot's independent calculation is **37.70161**.

The lesson is not "read the codebook." There is no codebook — the Stata file ships with empty variable labels. The lesson is that a column whose meaning you have inferred should be checked against something outside the file before you build a finding on it, and that a disagreement with the file's own derived column is evidence about *you* first.

Coverage, before anybody plots pct_split

`pct_split` is `NA` for **1946, 1948, 1950, 1962 and 1966**. Presidential results by congressional district are not published by any government — they are calculated by outside analysts from precinct returns, and nobody did the 1940s. `split_coverage` gives the share of districts each year's figure rests on; anything under 90% is masked.

Students who plot it without checking will draw a line straight through 5 gaps. That is the point of the option.

Option B — Split districts

Verified output

```
round(tapply(d$split_district, 10 * (d$year %/% 10),
             function(x) 100 * mean(x, na.rm = TRUE)), 1)
```

```
#> 1940 1950 1960 1970 1980 1990 2000 2010 2020
#>  NaN 28.2 29.4 36.6 38.6 26.8 16.7  9.9  4.5
```

`NaN` in the 1940s, peaking at 38.6% in the 1980s, then 16.7% in the 2000s, 9.9% in the 2010s and 4.5% in the 2020s.

3/4: Plots the decline and calls it nationalisation.

4/4: Is careful about what "split" means here — whether the district's presidential and House winners were of different parties, which requires both to be recorded. **The 1940s are `NaN`, not zero**: nobody computed presidential results by district for those years, so there is nothing to divide by. From 2016 on the flag is computed from `dv` and `dpres`, which **drops uncontested races automatically**, so the denominator shifts at the source splice as well as with coverage.

4/4, best: Refuses to read the decade path confidently off 8 averages, and notices that the `NaN` is a coverage statement rather than a finding. Also worth credit: asking where the flag comes from, and checking it rather than assuming — recomputed from `dv` and `dpres` it matches the file's own `split` column 100% of the time, whereas the `dvp` mistake described above would have matched only 71.9%.

Option C — The South, defined

Verified output

```r
c(n_south_states = length(unique(d$state[d$south == 1])))
```

```
#> n_south_states
#>             11
```

```r
sort(unique(d$state[d$south == 1]))
```

```
#>  [1]  1  4  9 10 18 24 33 40 42 43 46
```

11 states, coded numerically — the Confederacy.

3/4: Reports that it is 11 states and that the definition is a choice.

4/4: Tests whether the finding depends on it. It does not, and cannot: at 60.5–71.3% uncontested in the mid-century South against 3.4–10.2% elsewhere in any decade, **no plausible reassignment of a border state changes the picture.** A student who checks rather than assumes has done the work.

4/4, best: Notes that the 11-state definition is itself a theoretical claim — that the relevant thing about the South is Confederate history rather than, say, Democratic-primary dominance — and that Oklahoma, Kentucky and West Virginia are excluded by that logic despite similar politics in the period. **The variable encodes an argument.**

Grading

Four-point scale from the syllabus.

- **Never penalise code.** It is your code.
- **The move that earns a 4** is reasoning from the *absence* of data — the empty rows and the missing column — rather than only from the numbers present.
- **A 2** is reporting the competitiveness decline without engaging with the denominator question after Step 6 raised it.
- **Reward anyone who is disturbed by the 1950s South figure of 71.3%.** It is the correct reaction and the reason the lab exists.

Anticipated questions

- *"Why do some columns stop before the data does?"* — The series runs 1946–2024, but it is two sources spliced: Jacobson supplies 15,228 races through 2014, and the Clerk of the House the remaining 2,175 from 2016 on. The Clerk's tables report votes, not careers, so **`incwin` and `dvp` are empty for every race after 2014** — 2,175 rows. Anything about incumbency or the previous election therefore stops at 2014, which is why the incumbency figures above are quoted over the races where the flag exists. The splice itself is sound: the two sources overlap for 2004–2014 and agree to a median of 0.03 points across 2,260 districts.

- *"Were uncontested seats always Democratic in the South?"* — Overwhelmingly in the early period. The file records the incumbent won but not the party in the columns we kept, so **do not assert it from this data** — say it is historically established elsewhere.
- *"Is 5 points the right definition of competitive?"* — It is a convention. Anyone can change it in one line, and they should be invited to.
- *"Does redistricting explain the decline?"* — Partly, and the redistricting lab is where that argument gets made properly. Note that the decline runs across decades with very different redistricting regimes.

Connections

- **Redistricting** — the same question from presidential votes; compare the "empty middle" finding directly.
- **The midterm penalty** — Option A's null sits against that lab's strongest result.
- **CPS turnout** — missing values that mean something, second appearance.
- **EAVS, precincts, redlining** — "per what?", again.